

Becoming the Trainer: Attacking ML Training Infrastructure



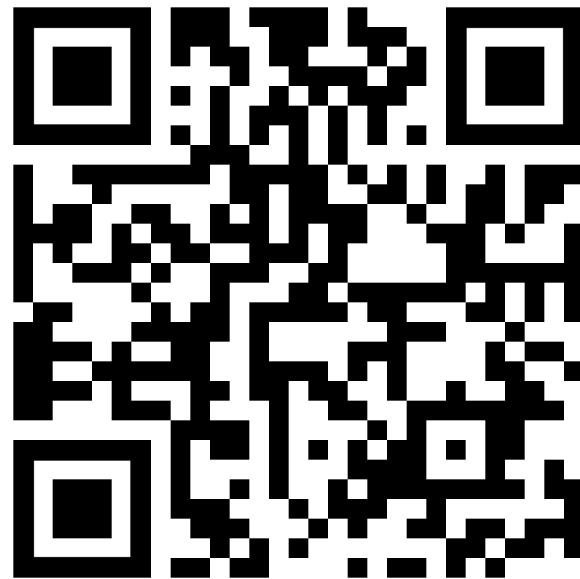
Brett Hawkins (@h4wkst3r)
Adversary Services, IBM X-Force Red



Blog Post and Tool

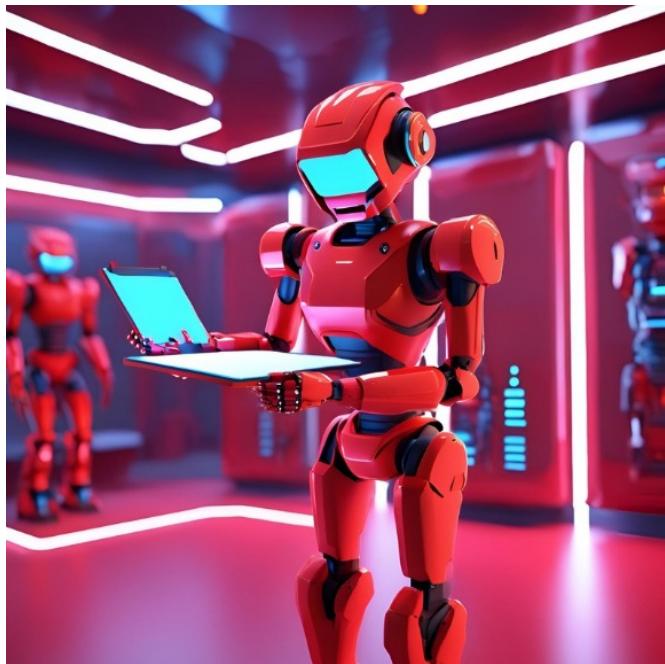


Blog Post



MLOKit

Agenda



1. Introduction
2. Background
3. Attacking ML Training Environments
 - Attack Scenarios
 - Demos
4. Protecting ML Training Environments
5. Conclusion
6. Q&A

Introduction



Who am I – Brett Hawkins

<https://h4wkst3r.github.io>



Current Role

Team Lead,
Adversary Services
IBM X-Force Red



Conference Speaker

Black Hat (US & EU),
BlueHat,
ShmooCon,
DerbyCon, Wild
West Hackin' Fest,
BSides, Hackers
Teaching Hackers



Open-Source Tool Author

SharPersist,
InvisibilityCloak,
SCMKit, ADOKit,
MLOKit

Research Drivers



Threat actors targeting AI/ML environments



Lack of research on attacking and defending ML infrastructure



Adoption of ML technologies by enterprises



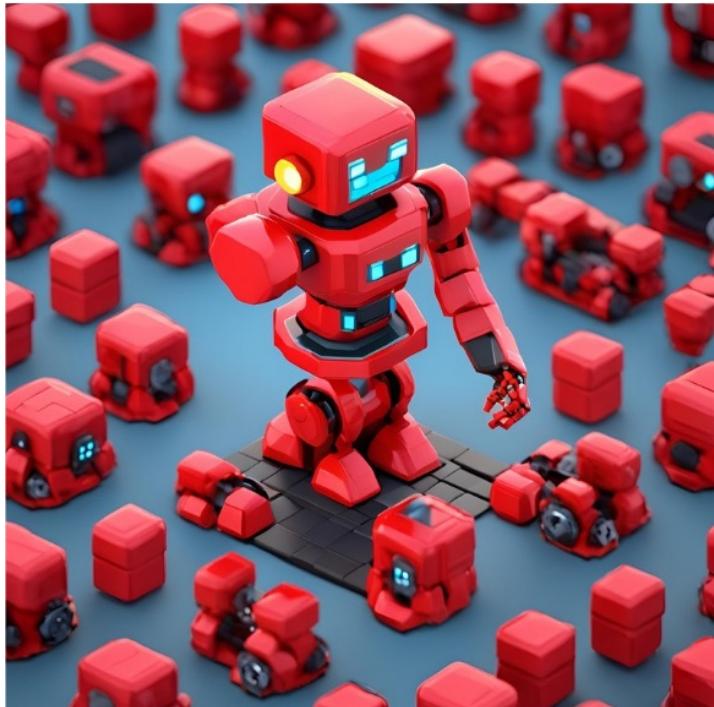
Lack of tooling to simulate attacks against platforms used in ML training envs

Threat Actor Motivation

The screenshot shows a blog post on the Oligo website. The title is "ShadowRay: First Known Attack Campaign Targeting AI Workloads Actively Exploited In The Wild". The authors are Avi Lumelsky, Guy Kaplan, and Gal Elbaz, posted on March 26, 2024. The post has 17 min read time and includes social sharing icons for LinkedIn, X, Facebook, and Email. A subtitle at the bottom states: "Thousands of publicly exposed Ray servers compromised as a result of Shadow Vulnerability".

Theft of models and weights, backdooring models for initial access or persistence, expanding access via lateral movement and privesc, sensitive data theft or deploying ransomware, model modification/poisoning for misclassification, degradation, fraud or ml-based detection evasion.

Attendee Takeaways



How to steal
models from
model registries



How to poison
models within ML
training platforms



How to defend
key components
of ML training
infrastructure



How to get code
execution via
attacks on ML
training
infrastructure

What is new in this research?

```
[*] INFO: Listing Model Artifact Location Info:  
Account Name: testworkspace5178193999  
Datastore Type: AzureBlob  
Container Name: azureml  
Path: ExperimentRun/dcid.AutoML_91114fd1-6657-4bf0-b51d-6f868e2c2033_42/outputs/mlflow-m  
[*] INFO: Getting associated datastore for model artifacts:  


| Account Name            | Container Name   D |
|-------------------------|--------------------|
| testworkspace5178193999 | azureml            |

  
[*] INFO: Uploading model artifacts  
[*] INFO: Uploading: conda.yaml  
[*] INFO: Uploading: MLmodel  
[*] INFO: Uploading: model.pkl  
[*] INFO: Uploading: python_env.yaml  
[*] INFO: Uploading: requirements.txt  
[+] SUCCESS: Model has been poisoned with model artifacts specified in source directory
```



Advanced attacks
against ML training
environments



New detection
rules (Azure ML
and SageMaker)



MLOKit tool
updates – NEW
supported
platforms and
attacks

My Perspective



I am

Offensive
Cybersecurity
Specialist

I am not

Data Scientist
AI/ML Engineer

Cloud Engineer

Detection Engineer

DevOps Engineer

Software Engineer

Background



Prior Work

Links to prior work are provided in appendix slides

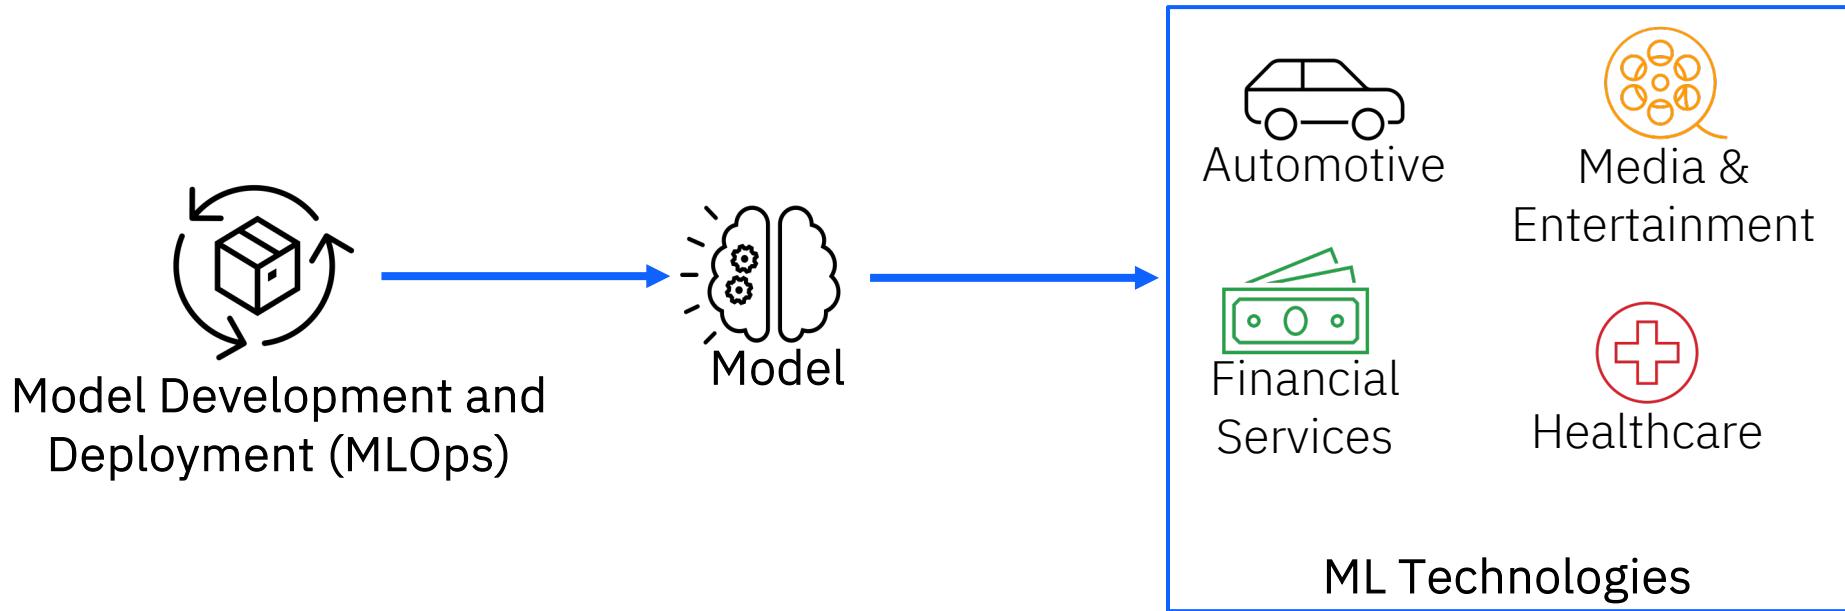
Chris Thompson ([@retBandit](#)) & I – ShmooCon 2025

Disrupting the Model: Abusing MLOps Platforms to Compromise ML Models and Enterprise Data Lakes

Or Azarzar ([@azarzaror](#)) – Blog Post (2021)

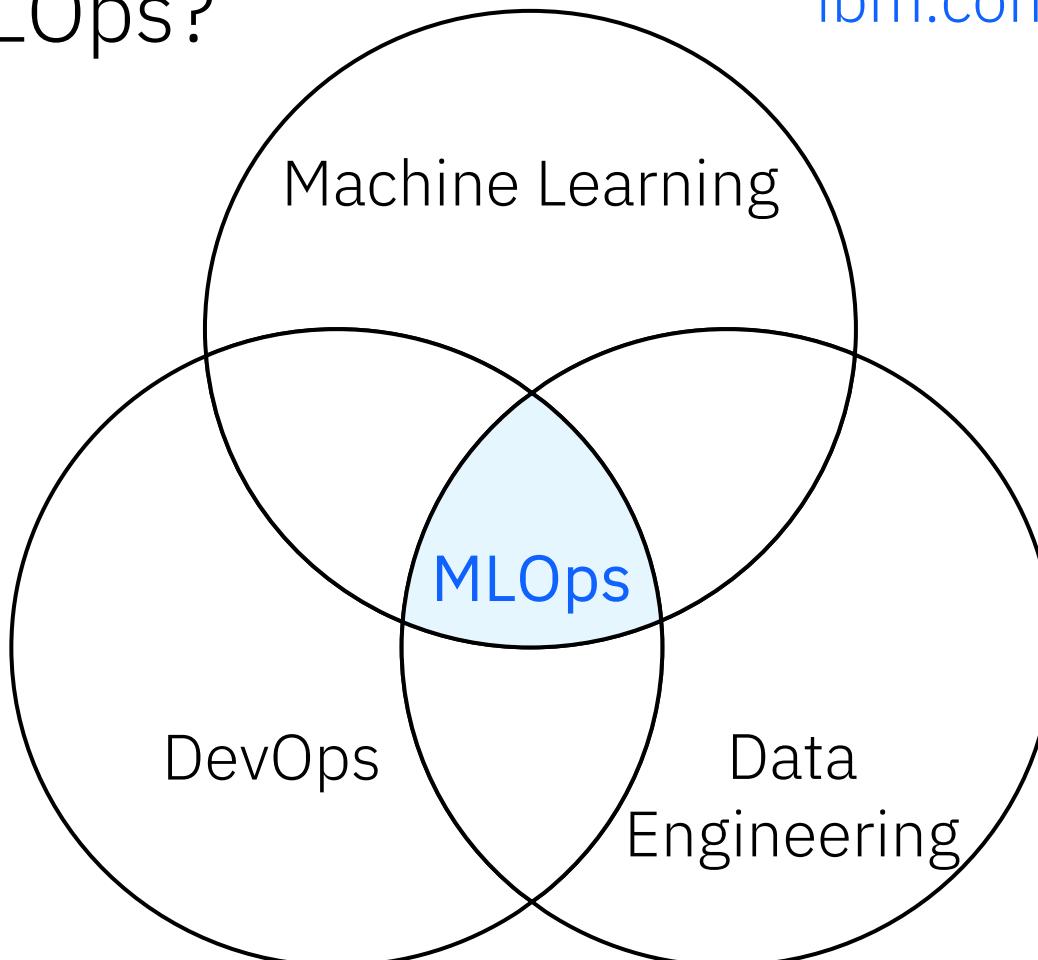
Protect Your Environment When Working with Amazon SageMaker

ML Technology Use Cases



What is MLOps?

ibm.com/topics/mlops



Popular MLOps Platforms



vertex.ai



Amazon SageMaker



DataRobot



Azure Machine Learning



databricks

W&B



Commercial

mlflow™

METAFLOW



TensorFlow



Kubeflow

semetic

Flyte

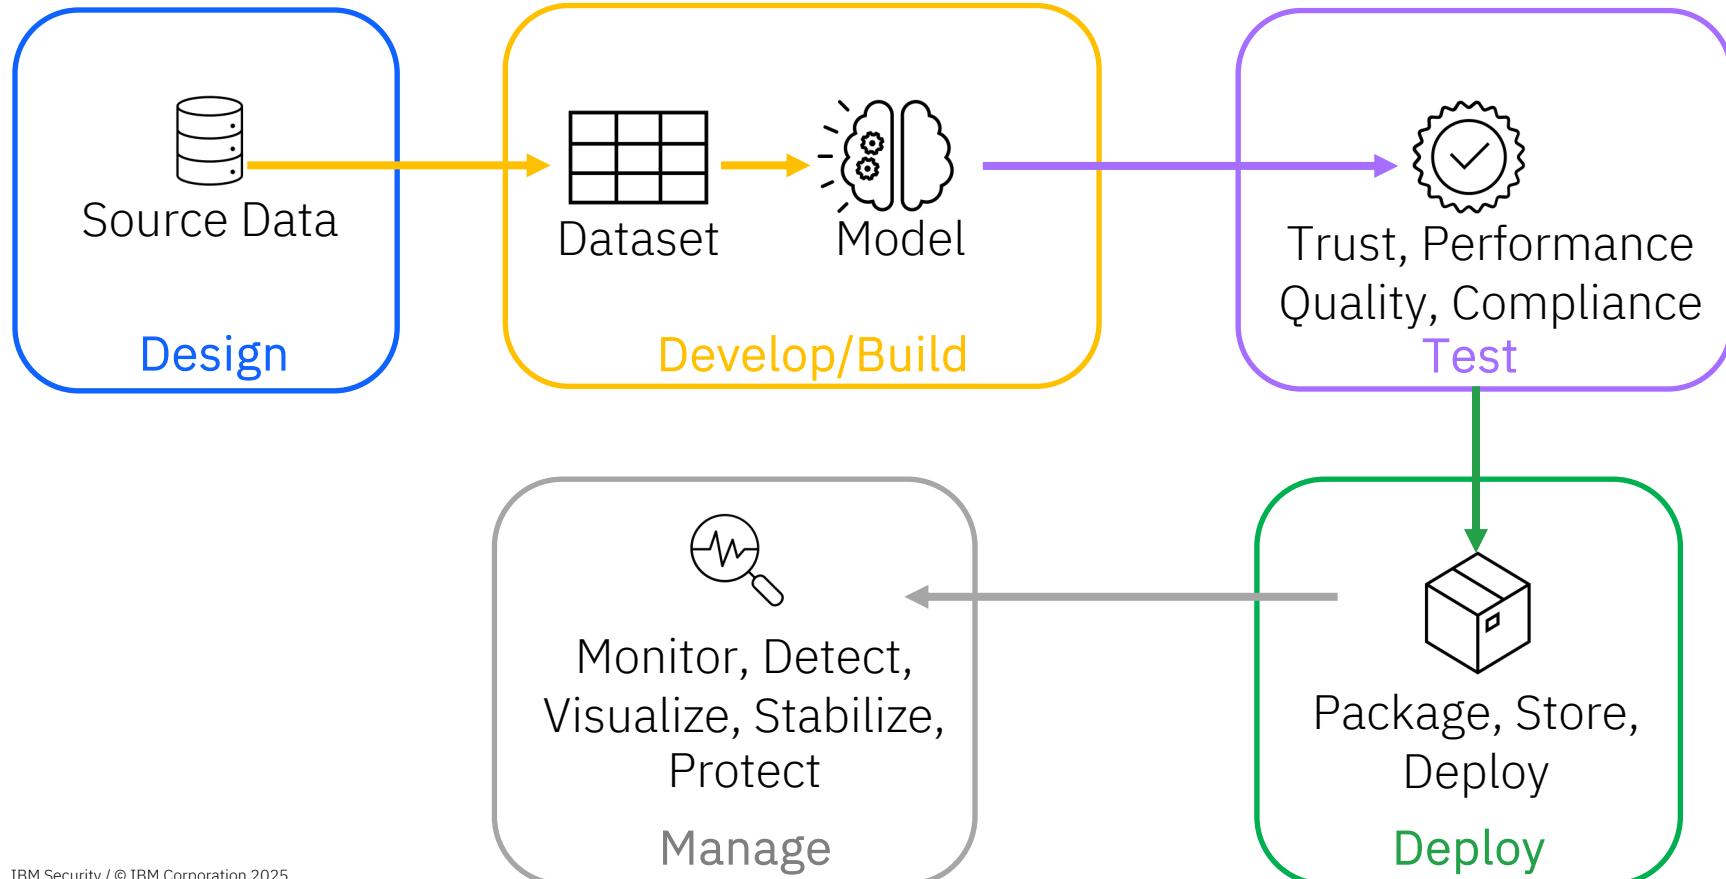
Apache Airflow

CORE

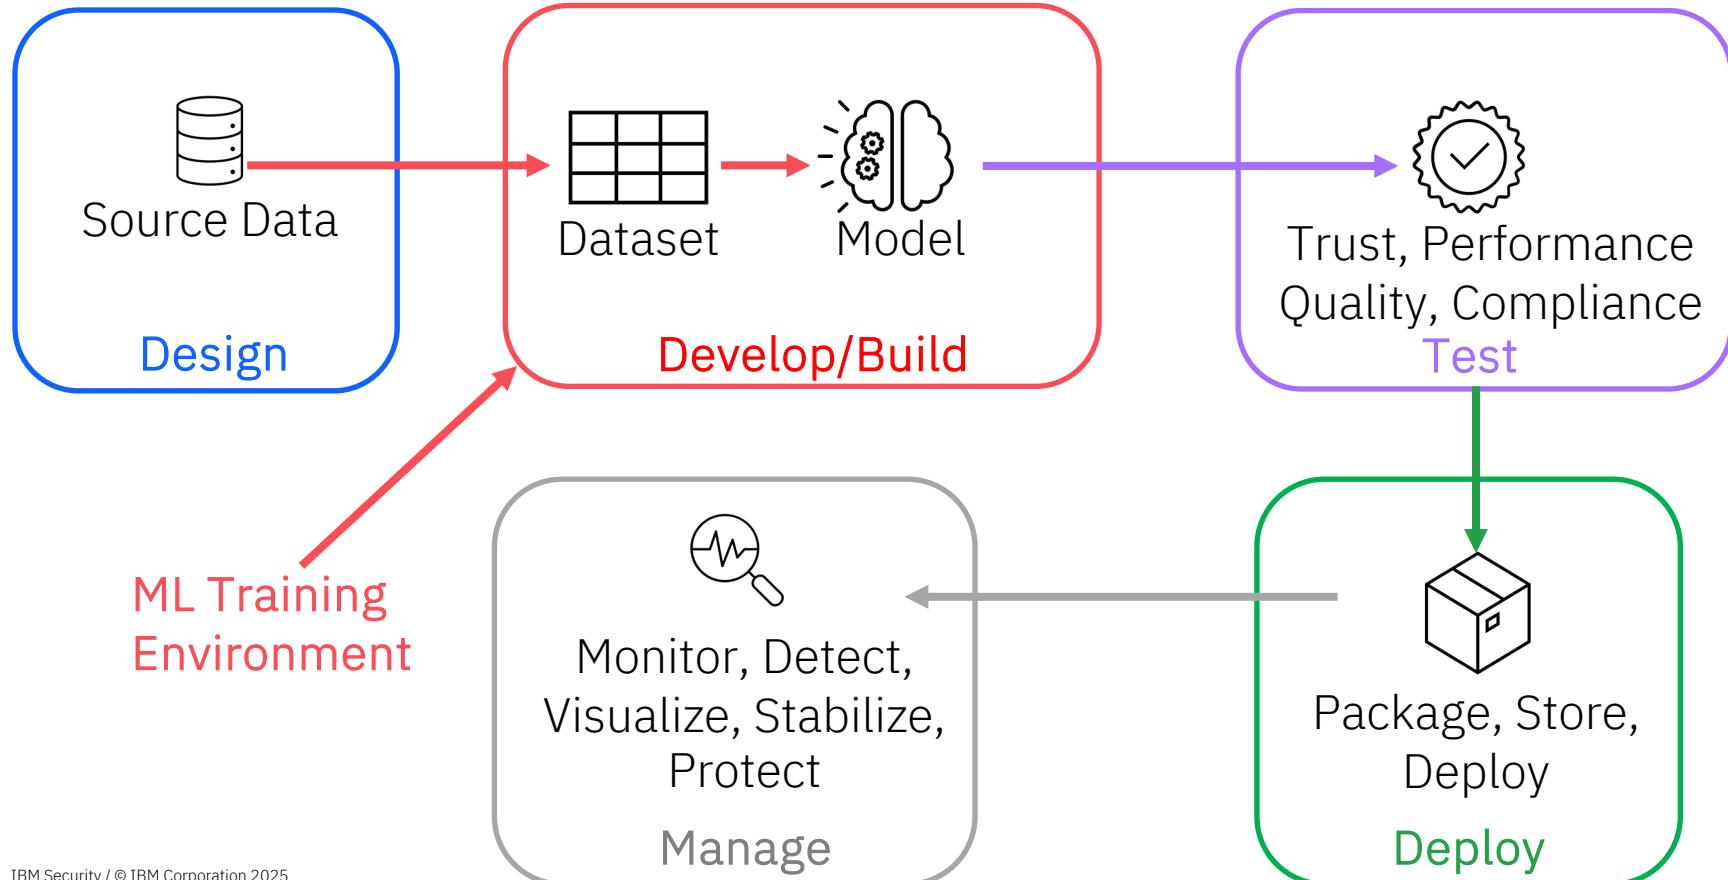
MLReef

Open-Source

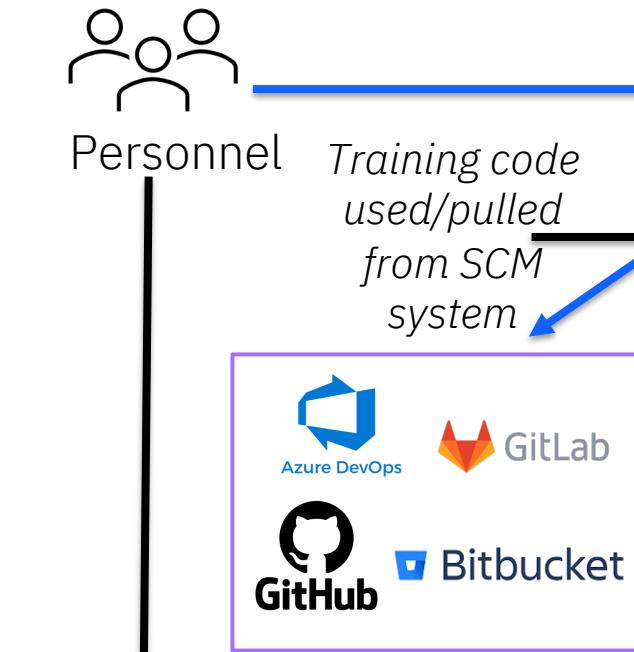
MLOps Lifecycle



MLOps Lifecycle - ML Training Environment

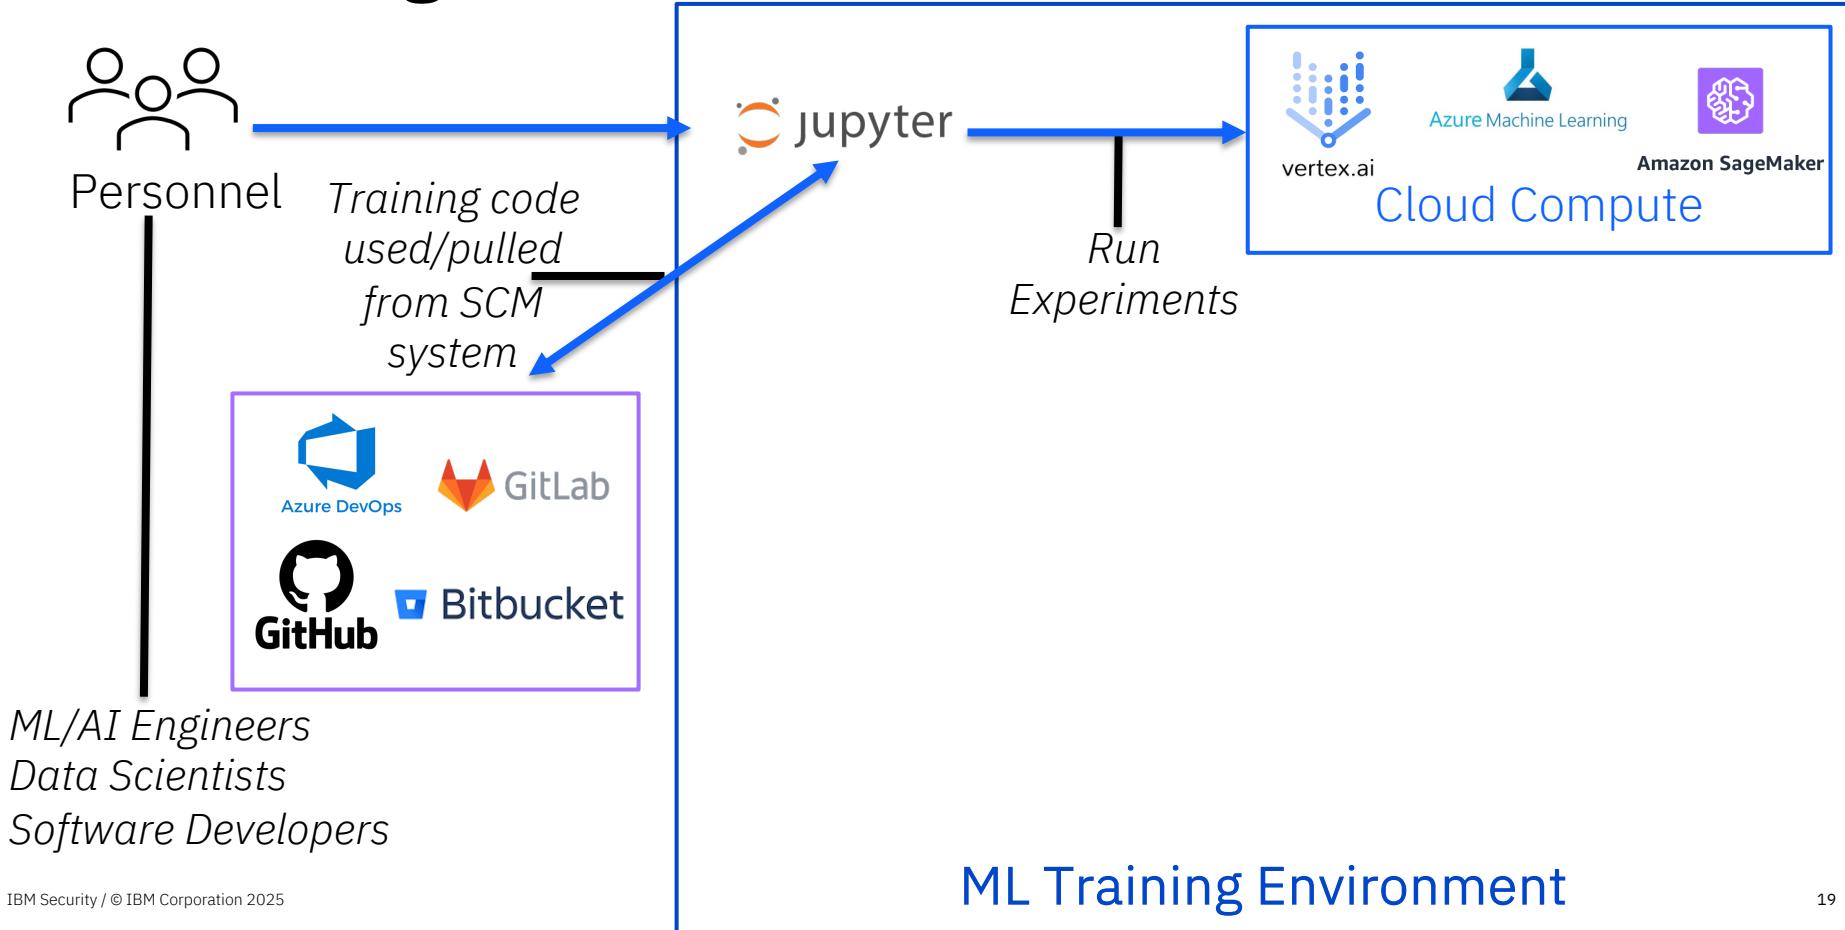


ML Training Environment Infrastructure

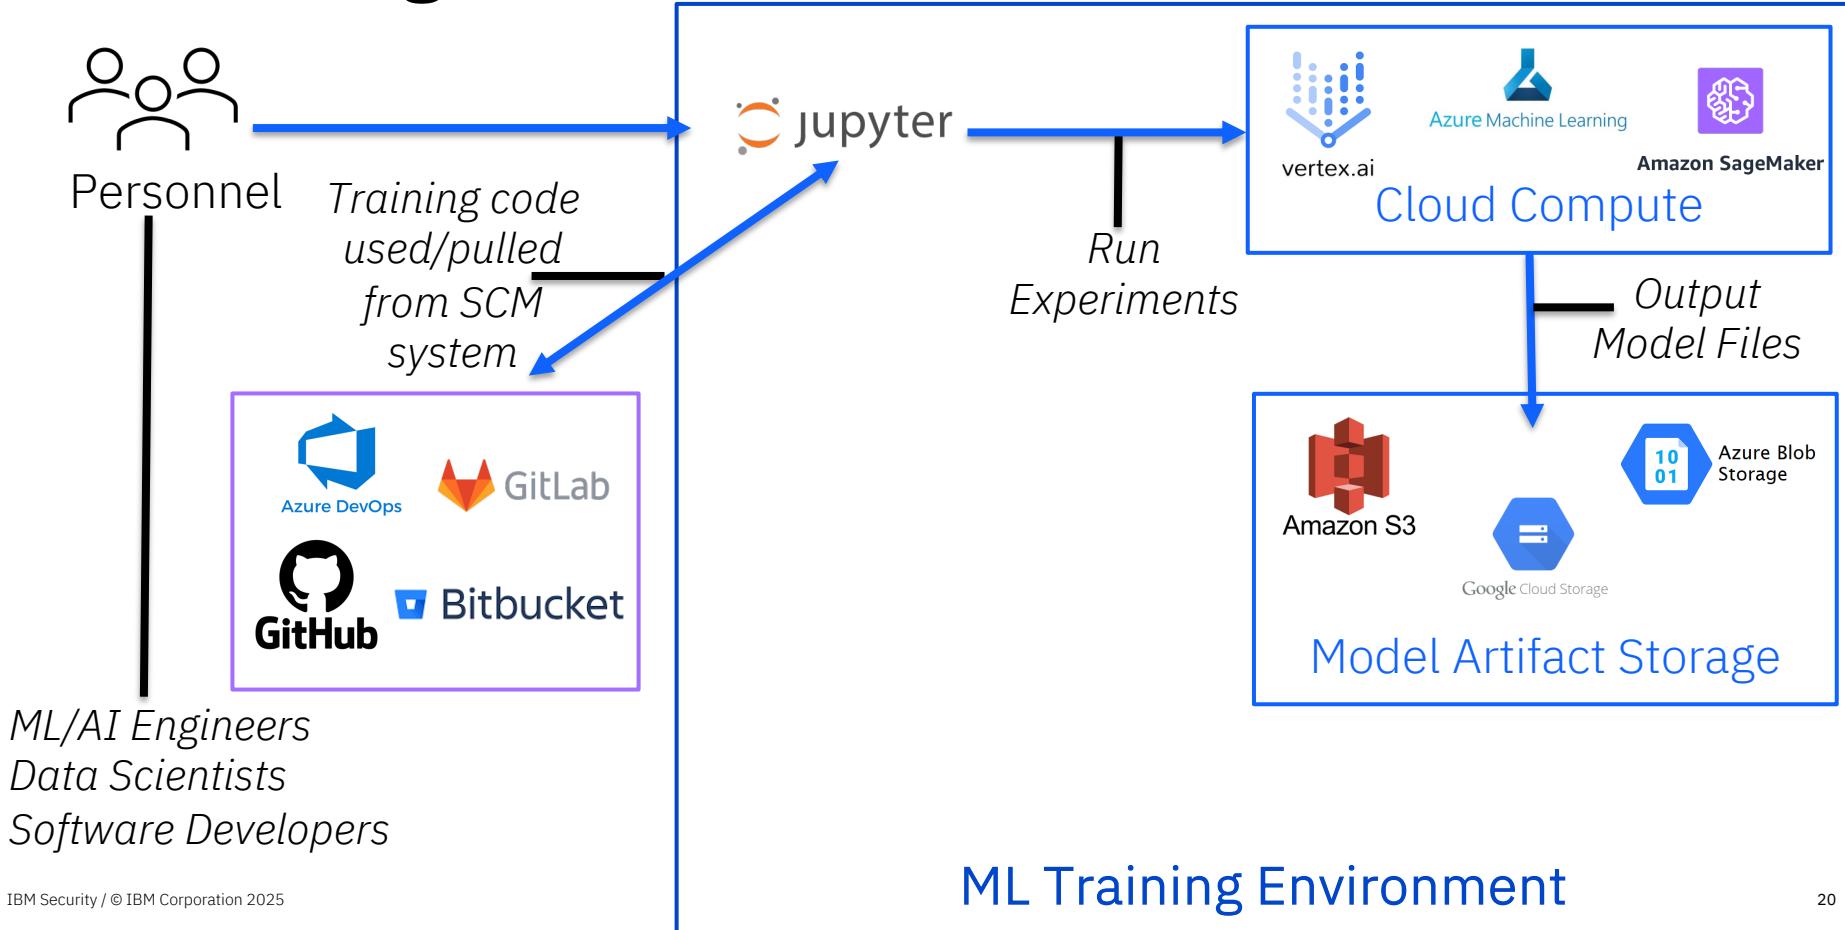


ML Training Environment

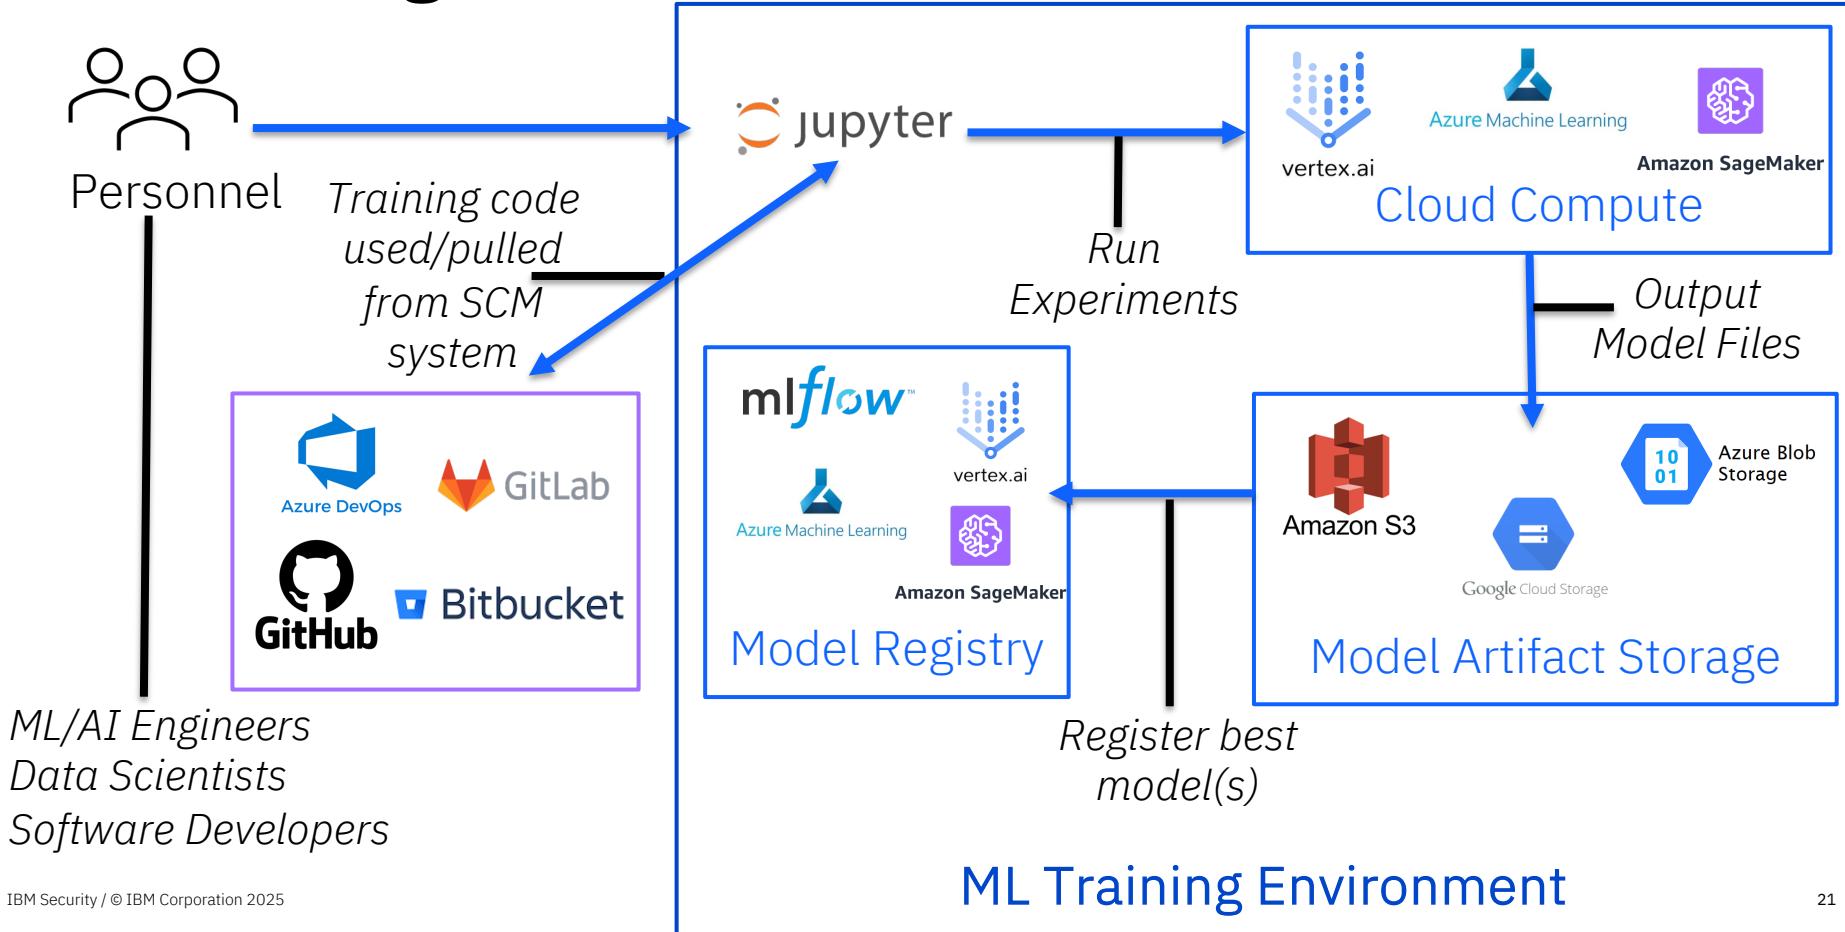
ML Training Environment Infrastructure



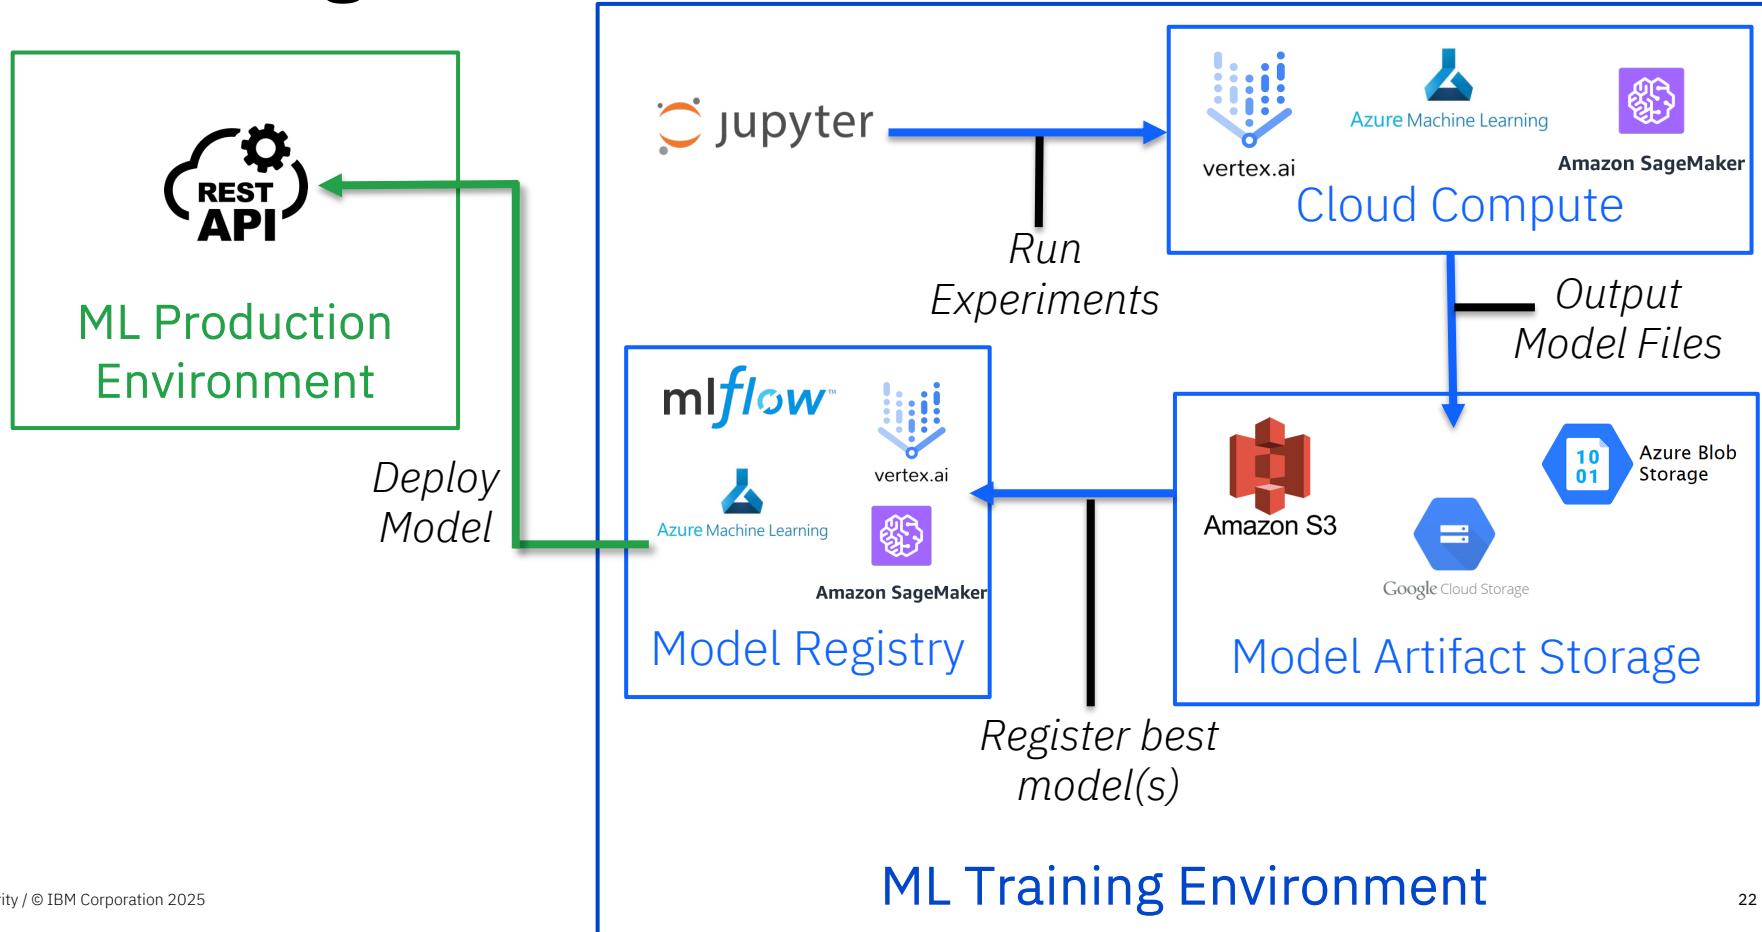
ML Training Environment Infrastructure



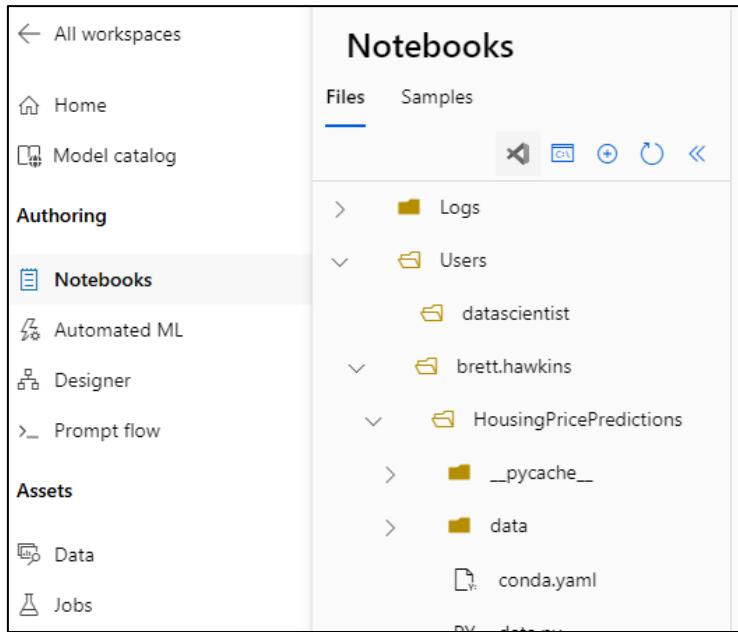
ML Training Environment Infrastructure



ML Training Environment Infrastructure



ML Training Infrastructure Components



The screenshot shows a left sidebar with navigation links: All workspaces, Home, Model catalog, Authoring (selected), Notebooks (selected), Automated ML, Designer, Prompt flow, Assets, Data, and Jobs. The main area is titled 'Notebooks' with tabs for 'Files' (selected) and 'Samples'. It displays a file tree with the following structure:

- > Logs
- < Users
 - > datascientist
 - < brett.hawkins
 - > HousingPricePredictions
 - > __pycache__
 - > data
 - cond.yaml



Notebook Env

Contains ML training code



Model Registry
Track and version models



Cloud Compute
Infrastructure that performs ML training from ML training code

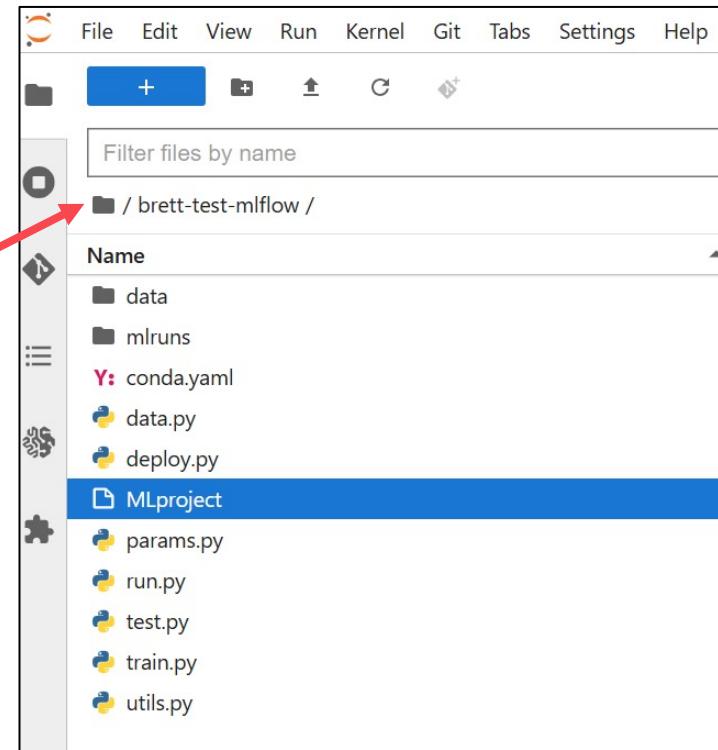


Model Artifact Storage
ML training artifact outputs (model weights, model files, etc.)

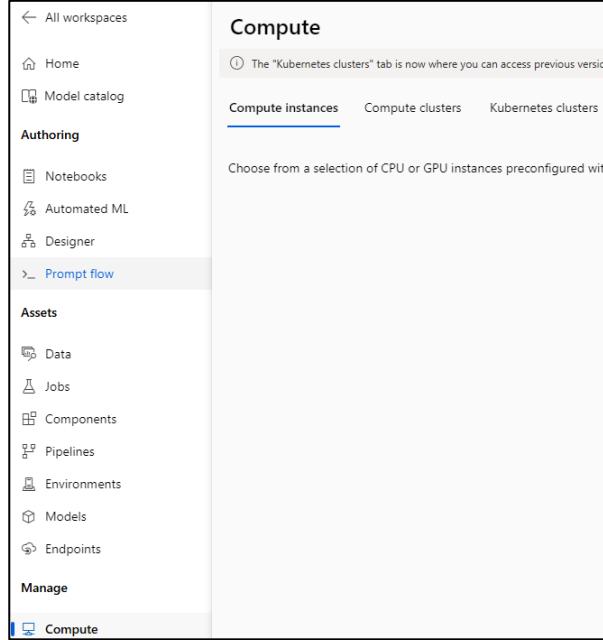
Notebook Env- SageMaker

The screenshot shows the AWS SageMaker AI dashboard. On the left, there's a sidebar with navigation links like Notebooks, Admin configurations, SageMaker AI dashboard, and JumpStart. The main panel displays details for a notebook instance named "instance/brett-test-mlflow-notebook". It includes sections for Lifecycle configuration, Status (InService), Creation time (Dec 04, 2024 19:43 UTC), Last updated (Dec 05, 2024 15:08 UTC), Volume Size (5GB EBS), Platform identifier (Amazon Linux 2, Jupyter Lab 3 (notebook-al2-v2)), and Minimum IMDS Version (2). Below this is a "Git repositories" section with a table:

| Name | Repository URL |
|-------------------------------|---|
| brett-test-mlflow - (Default) | https://git-codecommit.us-east-2.amazonaws.com/v1/repos/brett-test-mlf |



ML Training Infrastructure Components



The screenshot shows the 'Compute' tab in the ML Workbench interface. The left sidebar includes links for 'All workspaces', 'Home', 'Model catalog', 'Notebooks', 'Automated ML', 'Designer', 'Prompt flow' (which is selected), 'Assets' (with sub-links for Data, Jobs, Components, Pipelines, Environments, Models, and Endpoints), 'Manage', and 'Compute' (which is also selected). The main content area is titled 'Compute' and contains a message: 'The "Kubernetes clusters" tab is now where you can access previous version'. Below this are three tabs: 'Compute instances' (selected), 'Compute clusters', and 'Kubernetes clusters'. A note below says 'Choose from a selection of CPU or GPU instances preconfigured with'.



Notebook Env
Contains ML
training code



Model Registry
Track and version
models



Cloud Compute
Infrastructure that
performs ML
training from ML
training code



Model Artifact
Storage
ML training
artifact outputs
(model weights,
model files, etc.)

ML Training Infrastructure Components

mlflow 2.18.0 Experiments Models

Experiments

Displaying Runs from 5 Experiments

Runs Evaluation Experimental Traces Experimental

Search Experiments

- Default
- ElasticNet
- Some_Experiment_2
- some-random-experiment
- Some_Experiment_3

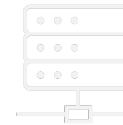
| Run Name | Created | Dataset |
|-----------------------|----------------|--------------------------------------|
| silent-kite-450 | 10 seconds ago | dataset (ea8b3b58) Training Dataset. |
| enthused-goat-111 | 22 seconds ago | dataset (ea8b3b58) Training Dataset. |
| efficient-pig-620 | 34 seconds ago | dataset (ea8b3b58) Training Dataset. |
| secretive-wolf-621 | 48 seconds ago | dataset (ea8b3b58) Training Dataset. |
| bursting-steed-895 | 1 minute ago | dataset (ea8b3b58) Training Dataset. |
| traveling-perch-716 | 1 minute ago | dataset (ea8b3b58) Training Dataset. |
| unequalled-donkey-781 | 1 minute ago | dataset (ea8b3b58) Training Dataset. |



Notebook Env
Contains ML training code



Model Registry
Track and version models



Cloud Compute
Infrastructure that performs ML training from ML training code



Model Artifact Storage
ML training artifact outputs (model weights, model files, etc.)

Model Registry and Artifact Storage - MLFlow

Points to Experiment Run

The image displays two screenshots of the MLflow interface. The left screenshot shows the 'Registered Models' page with a registered model named 'output-reg-model'. It includes details like 'Version 1', 'Registered At: 2024-12-05 07:44:12', and 'Source Run: rumbling-ray-97'. The right screenshot shows the 'Artifacts' tab for an experiment run named 'rumbling-ray-97' under the 'ElasticNet' experiment. It lists various artifacts such as code, MLmodel, conda.yaml, input_example.json, model.pkl, python_env.yaml, requirements.txt, and serving_input_example.json. A red box highlights the 'Source Run' field in the Model Registry and the 'Artifacts' section in the Experiment Run details.

mlflow 2.18.0 Experiments Models

Registered Models > output-reg-model >

Version 1

Registered At: 2024-12-05 07:44:12

Source Run: rumbling-ray-97

Stage (deprecated): None ⓘ

> Description Edit

Model Registry

mlflow 2.18.0 Experiments Models

ElasticNet > rumbling-ray-97

Overview Model metrics System metrics Artifacts

ElasticNet

- code
- MLmodel
- conda.yaml
- input_example.json
- model.pkl
- python_env.yaml
- requirements.txt
- serving_input_example.json

ElasticNet

Path: s3://bretttestbucket8878

MLflow Model

The code snippets below demonstrate model registry.

Model Registry and Artifact Storage - MLFlow

Points to Experiment Run

The screenshot shows the MLflow Model Registry interface. A registered model named "output-reg-model" is listed. The details shown include:

- Version 1
- Registered At: 2024-12-05 07:44:12
- Source Run: [rumbling-ray-97](#)
- Stage (deprecated): None
- Description Edit

Model Registry

The screenshot shows the MLflow Experiment Run interface for run "rumbling-ray-97". The "Artifacts" tab is selected, displaying the following information:

- ElasticNet
- Path: s3://brettestbucket8878
- MLflow Model**
- The code snippets below demonstrate how to use the MLflow Model.

The artifacts listed under "ElasticNet" are:

- code
- MLmodel
- conda.yaml
- input_example.json
- model.pkl
- python_env.yaml
- requirements.txt
- serving_input_example.json

Model Artifacts from
Experiment Run

The screenshot shows the Amazon S3 bucket interface for the folder "ElasticNet/". The "Objects" tab is selected, displaying the following contents:

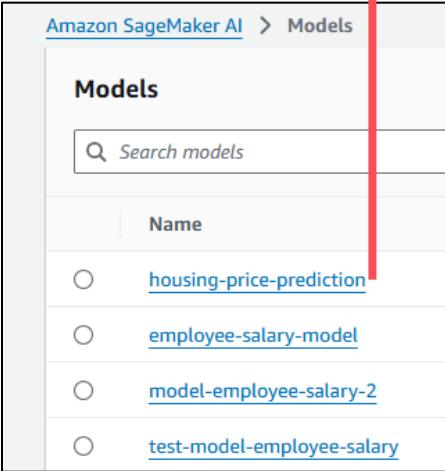
- Objects (8) Info
- Code
- MLmodel
- conda.yaml
- input_example.json
- model.pkl
- python_env.yaml
- requirements.txt
- serving_input_example.json

Properties tab is also visible.

Artifact
Storage (S3)

Model Registry and Artifact Storage - SageMaker

*Points to
Training Job*



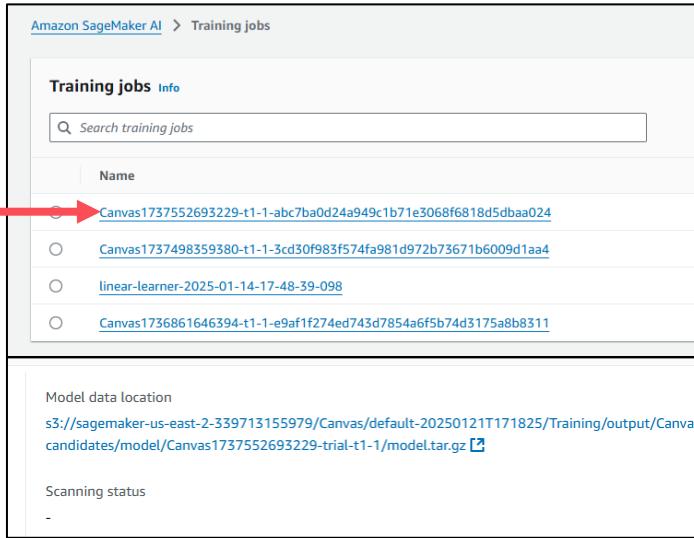
Amazon SageMaker AI > Models

Models

Search models

| Name |
|--|
| housing-price-prediction |
| employee-salary-model |
| model-employee-salary-2 |
| test-model-employee-salary |

Model Registry



Amazon SageMaker AI > Training jobs

Training jobs [Info](#)

Search training jobs

| Name |
|---|
| Canvas1737552693229-t1-1-abc7ba0d24a949c1b71e3068f6818d5dbaa024 |
| Canvas1737498359380-t1-1-3cd30f983f574fa981d972b73671b6009d1aa4 |
| linear-learner-2025-01-14-17-48-39-098 |
| Canvas1736861646394-t1-1-e9af1f274ed743d7854a6f5b74d3175a8b8311 |

Model data location
<s3://sagemaker-us-east-2-339713155979/Canvas/default-20250121T171825/Training/output/Canvascandidates/model/Canvas1737552693229-trial-t1-1/model.tar.gz>

Scanning status
-

Model Registry and Artifact Storage - SageMaker

*Points to
Training Job*

Amazon SageMaker AI > Models

Models

Search models

| Name |
|--|
| housing-price-prediction |
| employee-salary-model |
| model-employee-salary-2 |
| test-model-employee-salary |

Model Registry

Amazon SageMaker AI > Training jobs

Training jobs [Info](#)

Search training jobs

| Name |
|---|
| Canvas1737552693229-t1-1-abc7ba0d24a949c1b71e3068f6818d5dbaa024 |
| Canvas1737498359380-t1-1-3cd30f983f574fa981d972b73671b6009d1aa4 |
| linear-learner-2025-01-14-17-48-39-098 |
| Canvas1736861646394-t1-1-e9af1f274ed743d7854a6f5b74d3175a8b8311 |

*Artifact
Storage (S3)*

Canvas1737552693229-trial-t1-1/

Objects [Properties](#)

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can upload, download, and manage objects.

Find objects by prefix

| Name | Type |
|------------------------------|------|
| model.tar.gz | gz |

*Model Artifacts from
Training Job Run*

Model Registry and Artifact Storage – Azure ML

Points to Training Job



| Name | Version | Type | Job (Run ID) |
|-----------------------|---------|--------|---|
| salary-model | 1 | MLFLOW | helpful_leg_n2slbwmcq_40 |
| heart-failure-model | 1 | MLFLOW | kind_guava_06ftffx2s2_38 |
| taxifare-output-model | 1 | MLFLOW | AutoML_91114fd1-6657-4bf0 |



helpful_leg_n2slbwmcqq Completed

Overview Data guardrails Models + child jobs Outputs + logs

Refresh Edit and submit (preview) Register model

outputs

- _automl_internal

featurization

- data

pipeline

- conda.yaml
- MLmodel
- model.pkl
- python_env.yaml
- requirements.txt

Model Registry

Model Registry and Artifact Storage – Azure ML

| Model List | | | | |
|-----------------------|---------|--------|---|--|
| Name | Version | Type | Job (Run ID) | |
| salary-model | 1 | MLFLOW | helpful_leg_n2slbwmcq_40 | |
| heart-failure-model | 1 | MLFLOW | kind_guava_06ftffx2s2_38 | |
| taxifare-output-model | 1 | MLFLOW | AutoML_91114fd1-6657-4bf0 | |

Model Registry

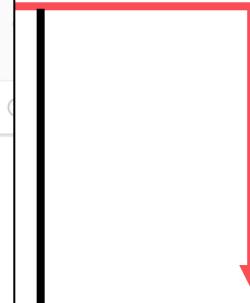
Points to
Training Job

helpful_leg_n2slbwmcq Completed

Overview Data guardrails Models + child jobs Outputs + logs

Refresh Edit and submit (preview) Register model

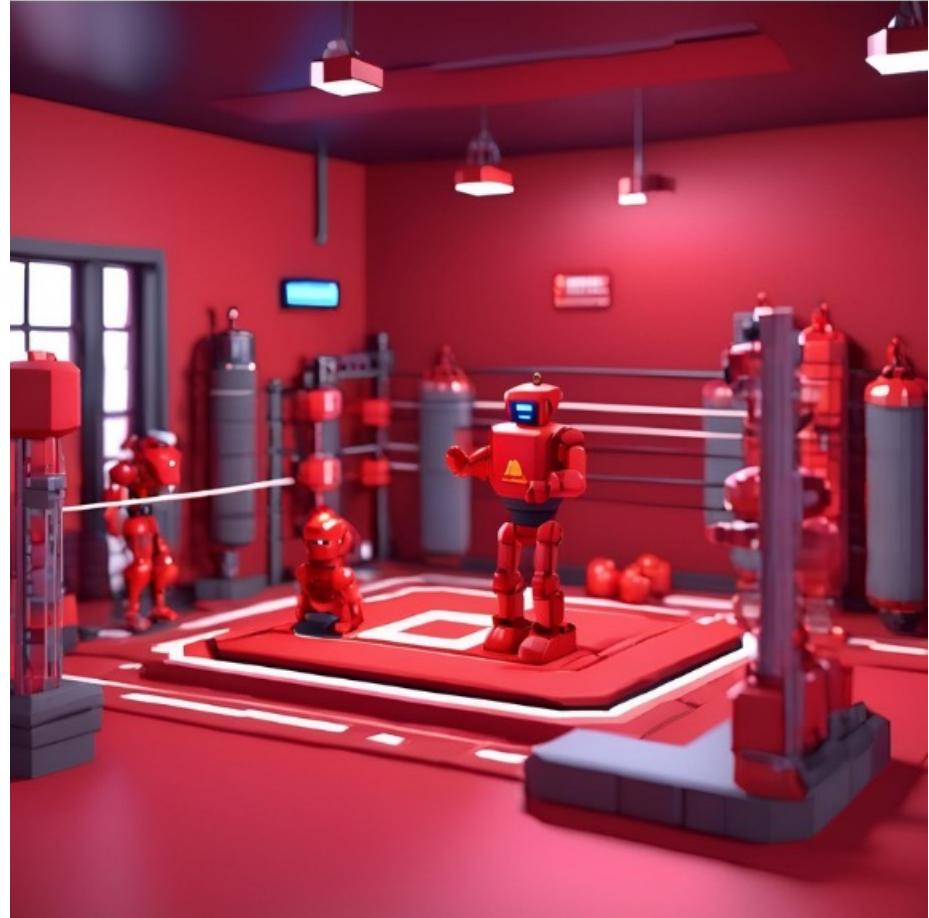
outputs _automl_internal featurization data pipeline conda.yaml MLmodel model.pkl python_env.yaml requirements.txt



Azure Blob
Storage

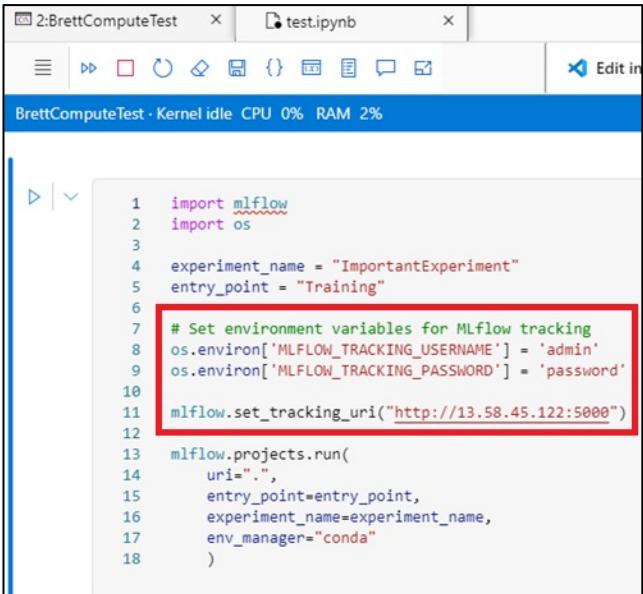
Model Artifacts
from Training
Job Run

Attacking ML Training Environments



Key Components – Attacker Perspective

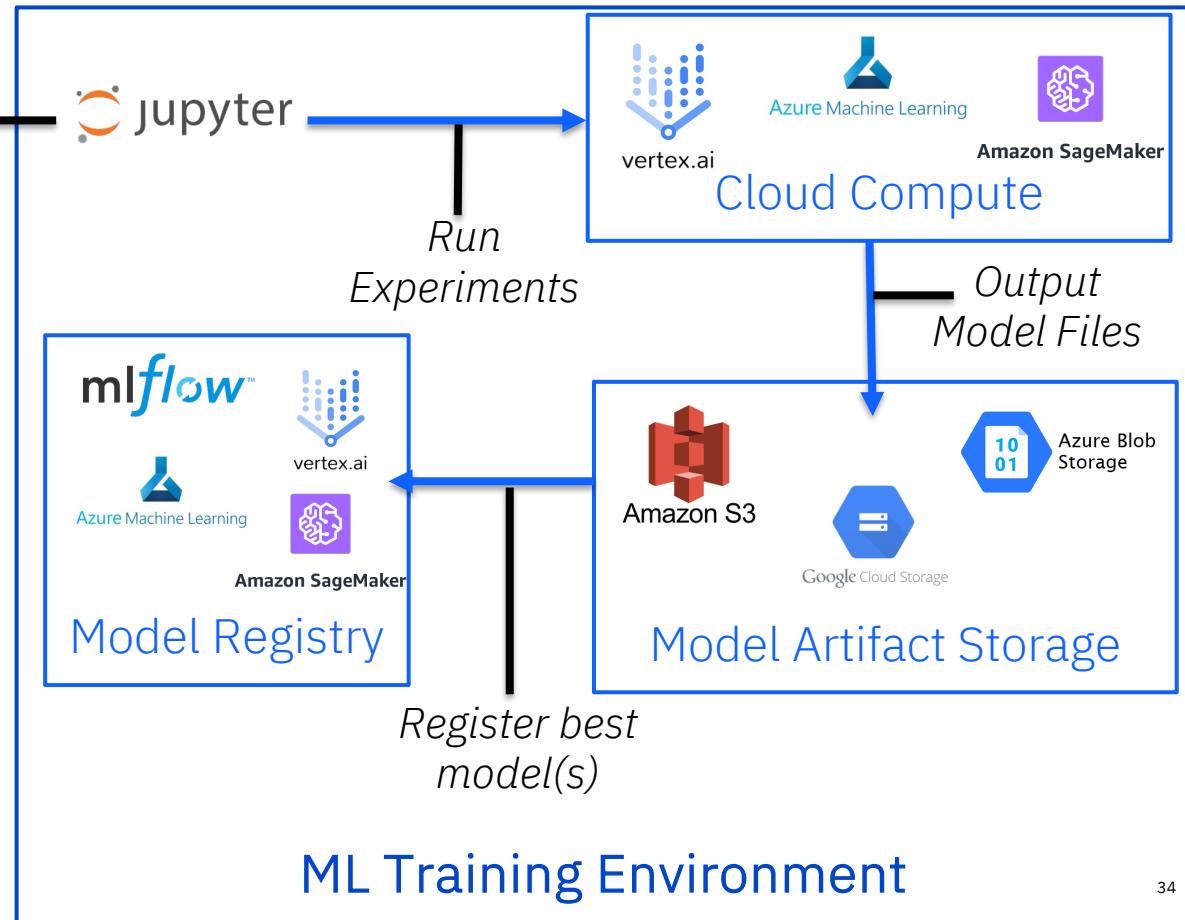
*Credentials and
info on other
infrastructure*



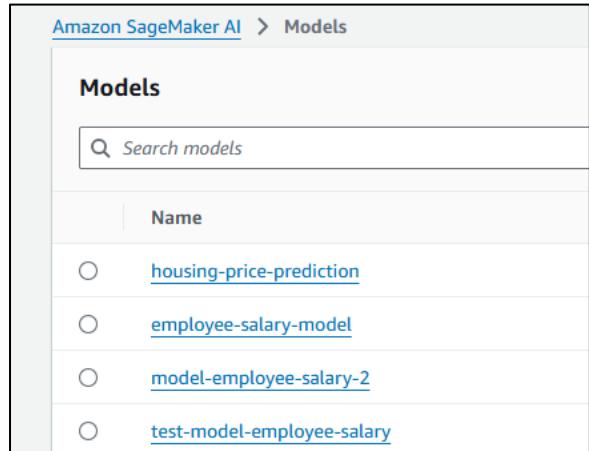
A screenshot of a Jupyter Notebook interface titled "BrettComputeTest". The notebook contains a single cell with the following Python code:

```
1 import mlflow
2 import os
3
4 experiment_name = "ImportantExperiment"
5 entry_point = "Training"
6
7 # Set environment variables for MLflow tracking
8 os.environ['MLFLOW_TRACKING_USERNAME'] = 'admin'
9 os.environ['MLFLOW_TRACKING_PASSWORD'] = 'password'
10
11 mlflow.set_tracking_uri("http://13.58.45.122:5000")
12
13 mlflow.projects.run(
14     uri=".",
15     entry_point=entry_point,
16     experiment_name=experiment_name,
17     env_manager="conda"
18 )
```

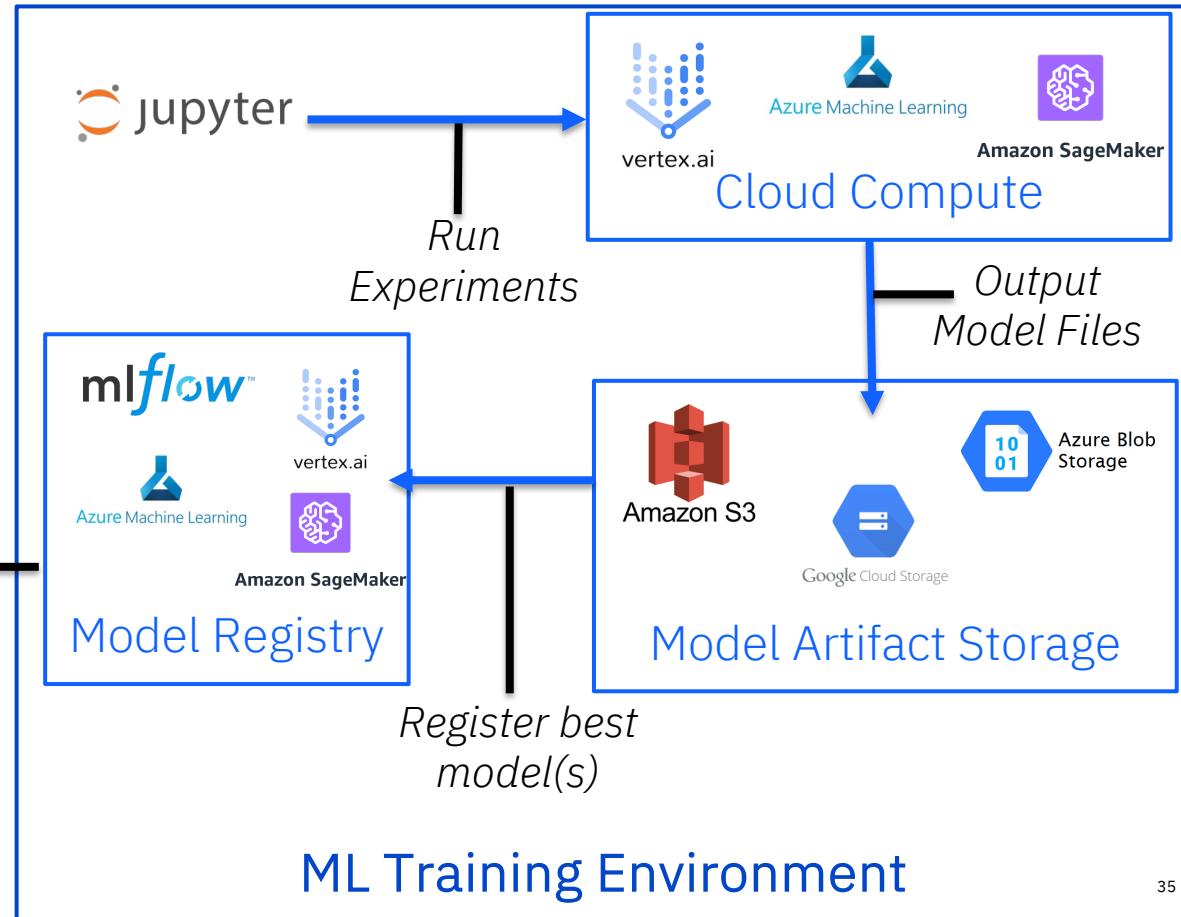
The line `mlflow.set_tracking_uri("http://13.58.45.122:5000")` is highlighted with a red box.



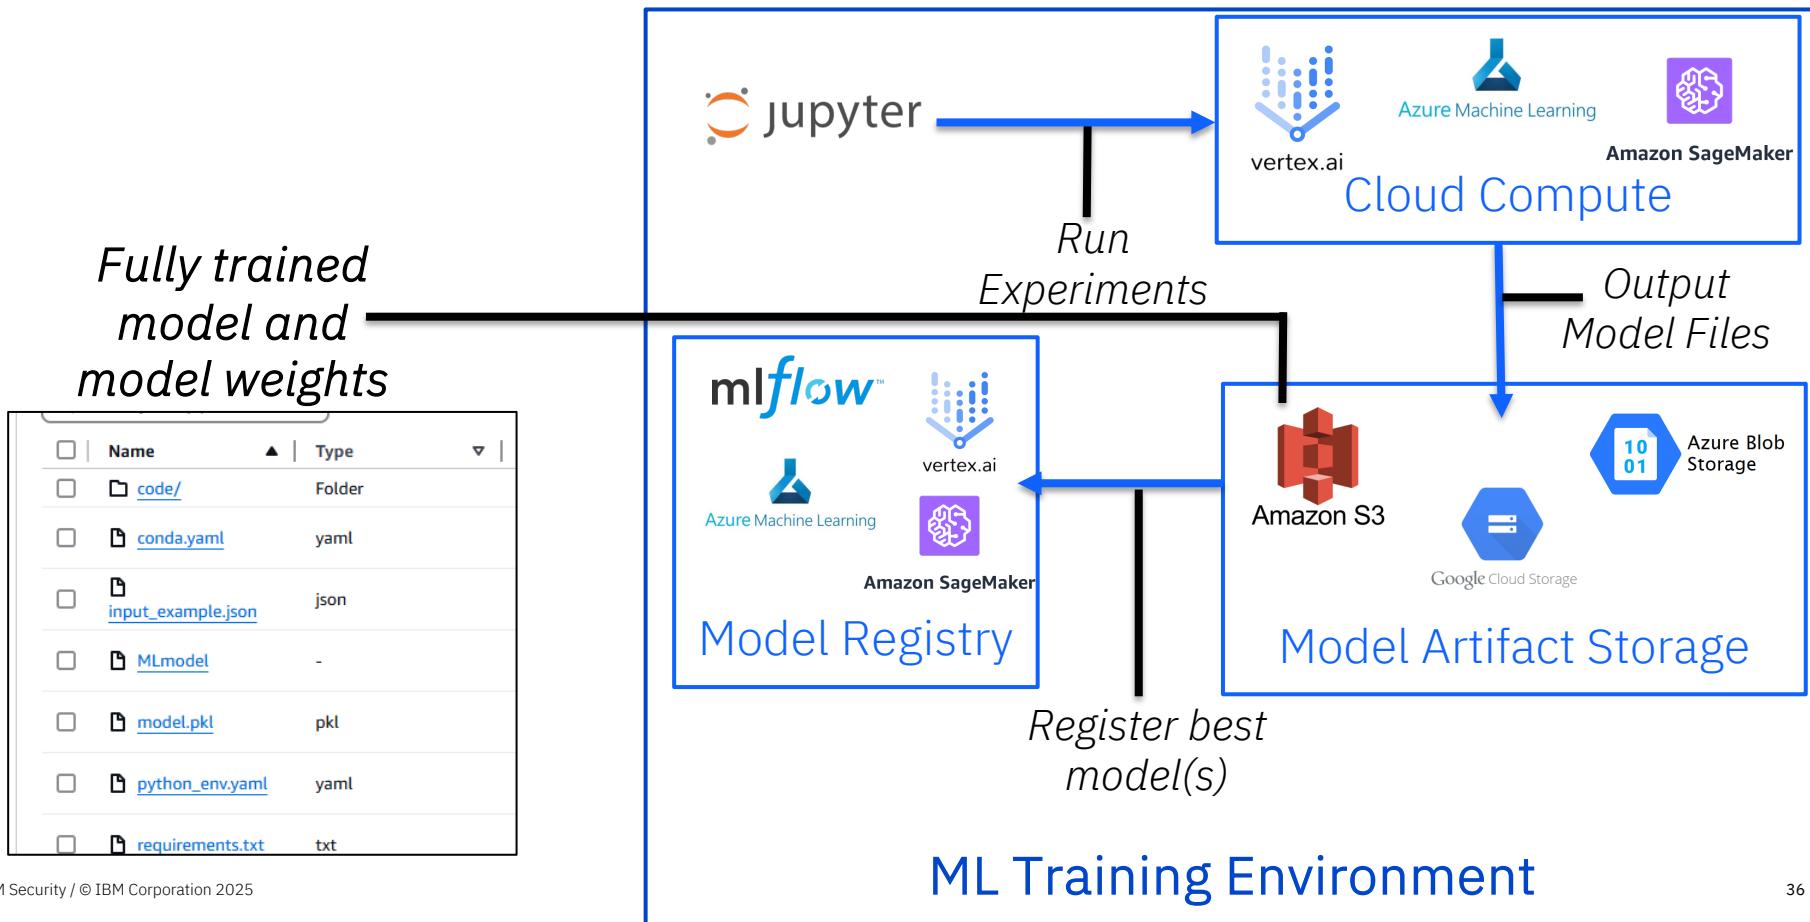
Key Components – Attacker Perspective



*Useful for model
reconnaissance*



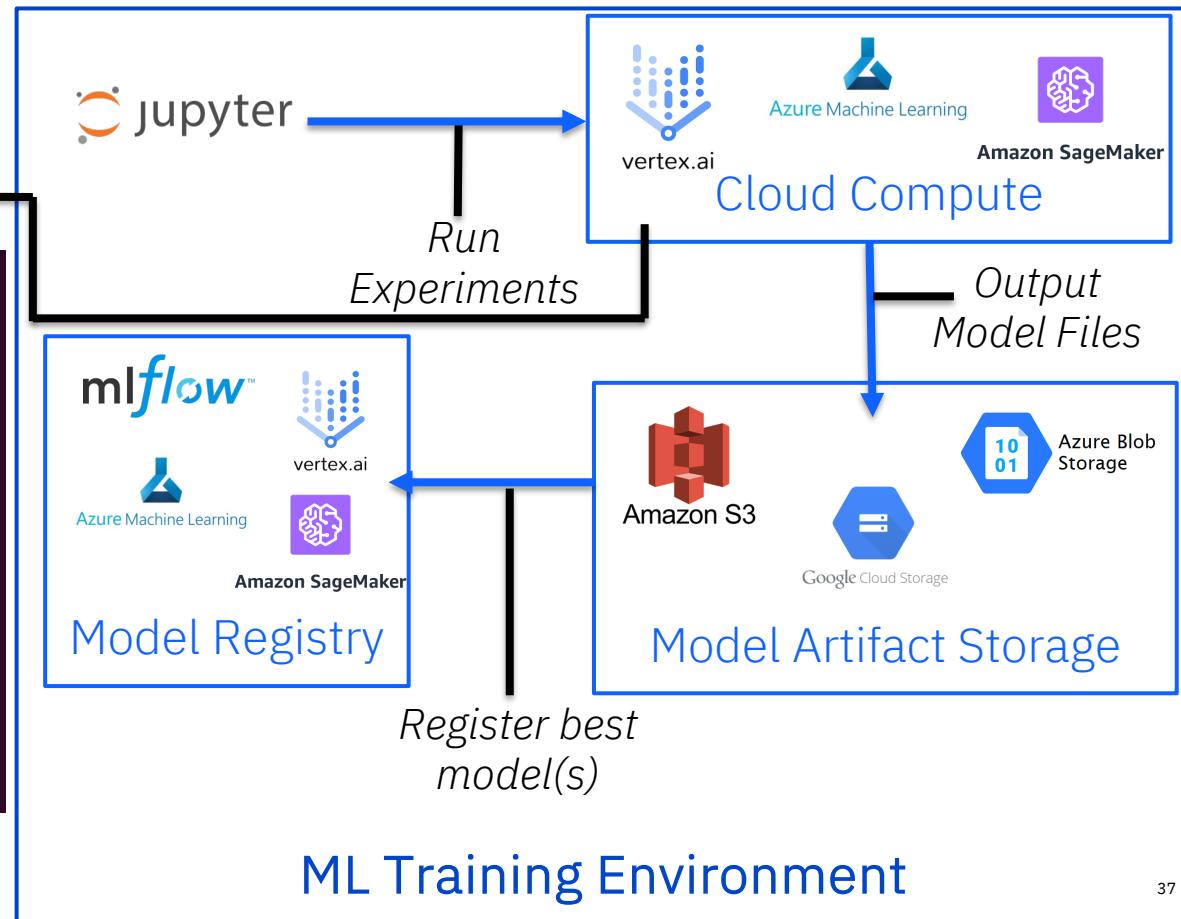
Key Components – Attacker Perspective



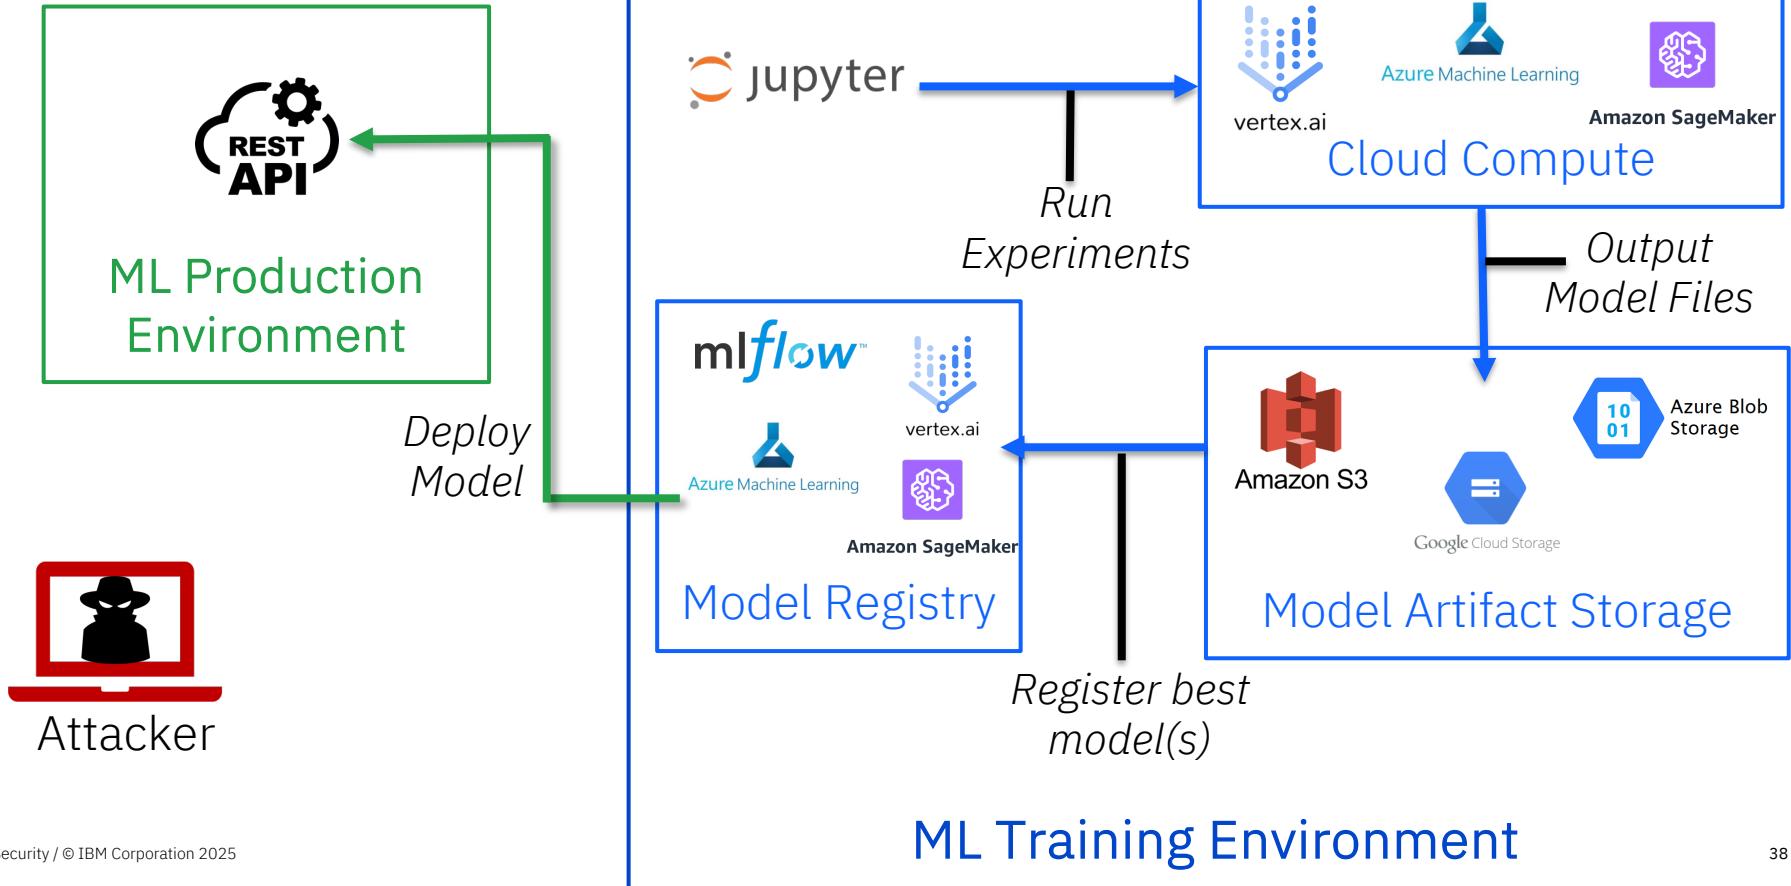
Key Components – Attacker Perspective

Sensitive environment variables and ML training code

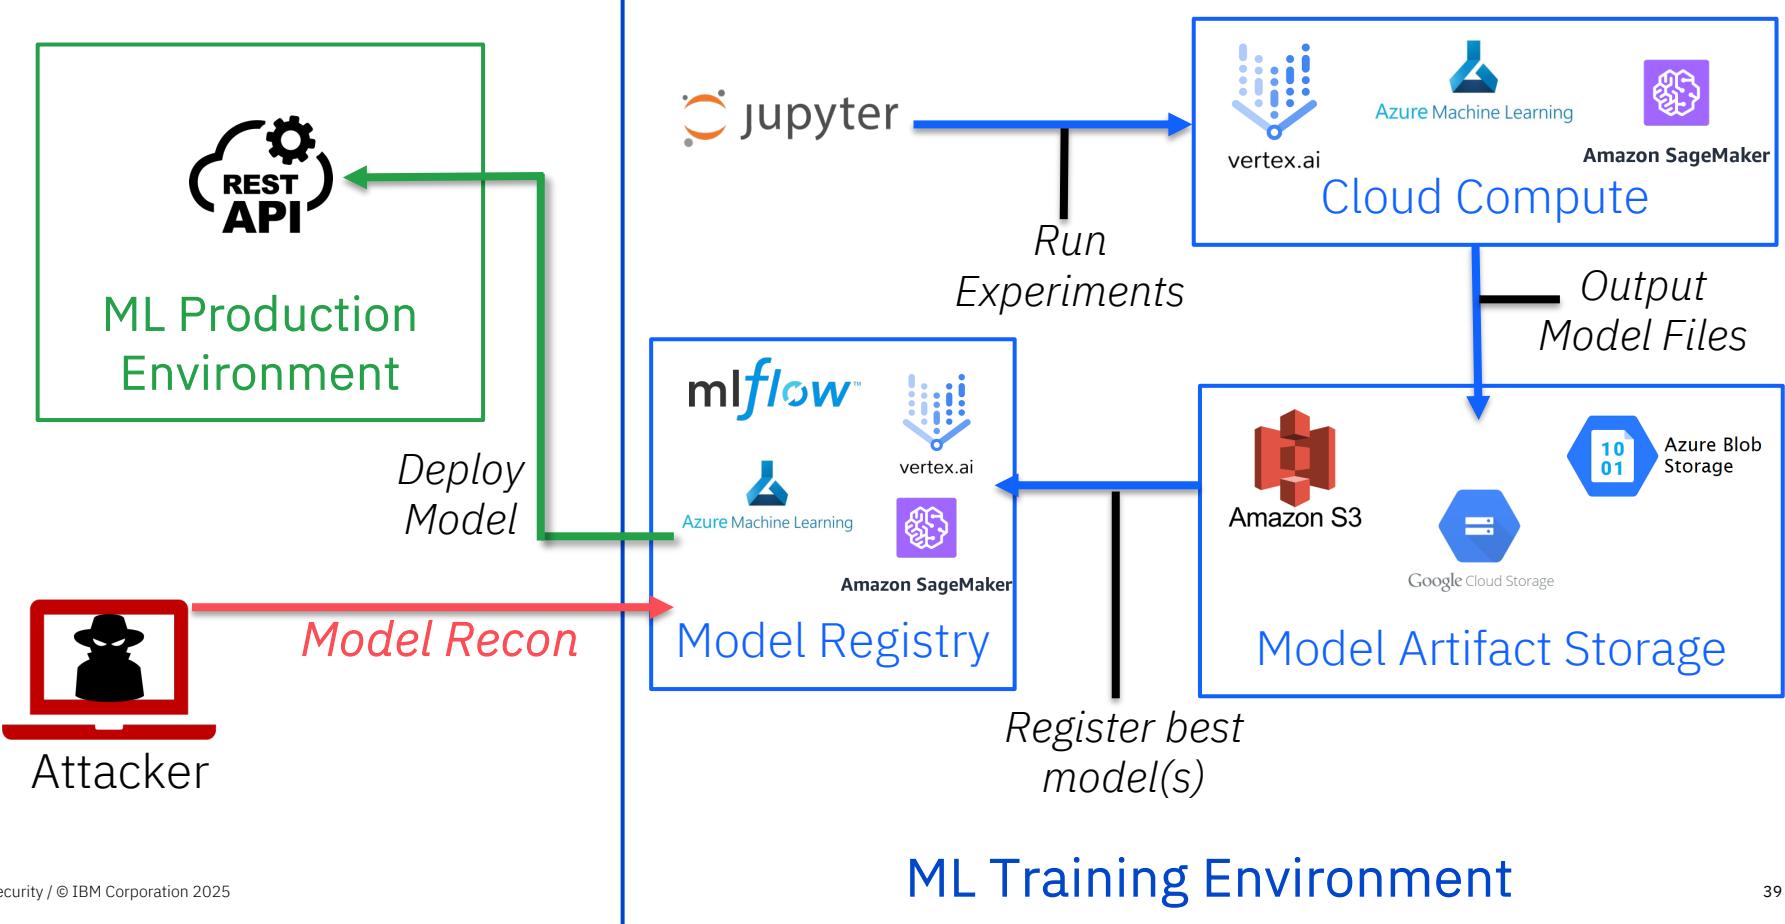
```
(base) [ec2-user@ip-172-16-11-6 brett-test-mlflow]$ ls -la  
total 64  
drwxrwxr-x 7 ec2-user ec2-user 4096 Dec  5 15:35 .  
drwxr-xr-x 6 ec2-user ec2-user 4096 Dec  5 15:08 ..  
-rw-rw-r-- 1 ec2-user ec2-user 1292 Dec  4 19:48 conda.yaml  
drwxrwxr-x 2 ec2-user ec2-user 4096 Dec  4 19:48 data  
-rw-rw-r-- 1 ec2-user ec2-user 1485 Dec  4 19:48 data.py  
-rw-rw-r-- 1 ec2-user ec2-user 879 Dec  4 19:48 deploy.py  
drwxrwxr-x 8 ec2-user ec2-user 4096 Dec  5 15:35 .git  
drwxrwxr-x 2 ec2-user ec2-user 4096 Dec  5 14:30 .ipynb_checkpoints  
-rw-rw-r-- 1 ec2-user ec2-user 131 Dec  5 15:35 MLproject  
drwxrwxr-x 4 ec2-user ec2-user 4096 Dec  5 15:15 mlruns  
-rw-rw-r-- 1 ec2-user ec2-user 986 Dec  4 19:48 params.py  
drwxrwxr-x 2 ec2-user ec2-user 4096 Dec  4 19:58 __pycache__  
-rw-rw-r-- 1 ec2-user ec2-user 270 Dec  5 14:30 run.py  
-rw-rw-r-- 1 ec2-user ec2-user 520 Dec  4 19:48 test.py  
-rw-rw-r-- 1 ec2-user ec2-user 1342 Dec  4 19:48 train.py  
-rw-rw-r-- 1 ec2-user ec2-user 409 Dec  4 19:48 utils.py  
(base) [ec2-user@ip-172-16-11-6 brett-test-mlflow]$ ls -la data  
total 904  
drwxrwxr-x 2 ec2-user ec2-user 4096 Dec  4 19:48 .  
drwxrwxr-x 7 ec2-user ec2-user 4096 Dec  5 15:35 ..  
-rw-rw-r-- 1 ec2-user ec2-user 452865 Dec  4 19:48 test.csv  
-rw-rw-r-- 1 ec2-user ec2-user 462137 Dec  4 19:48 train.csv  
(base) [ec2-user@ip-172-16-11-6 brett-test-mlflow]$
```



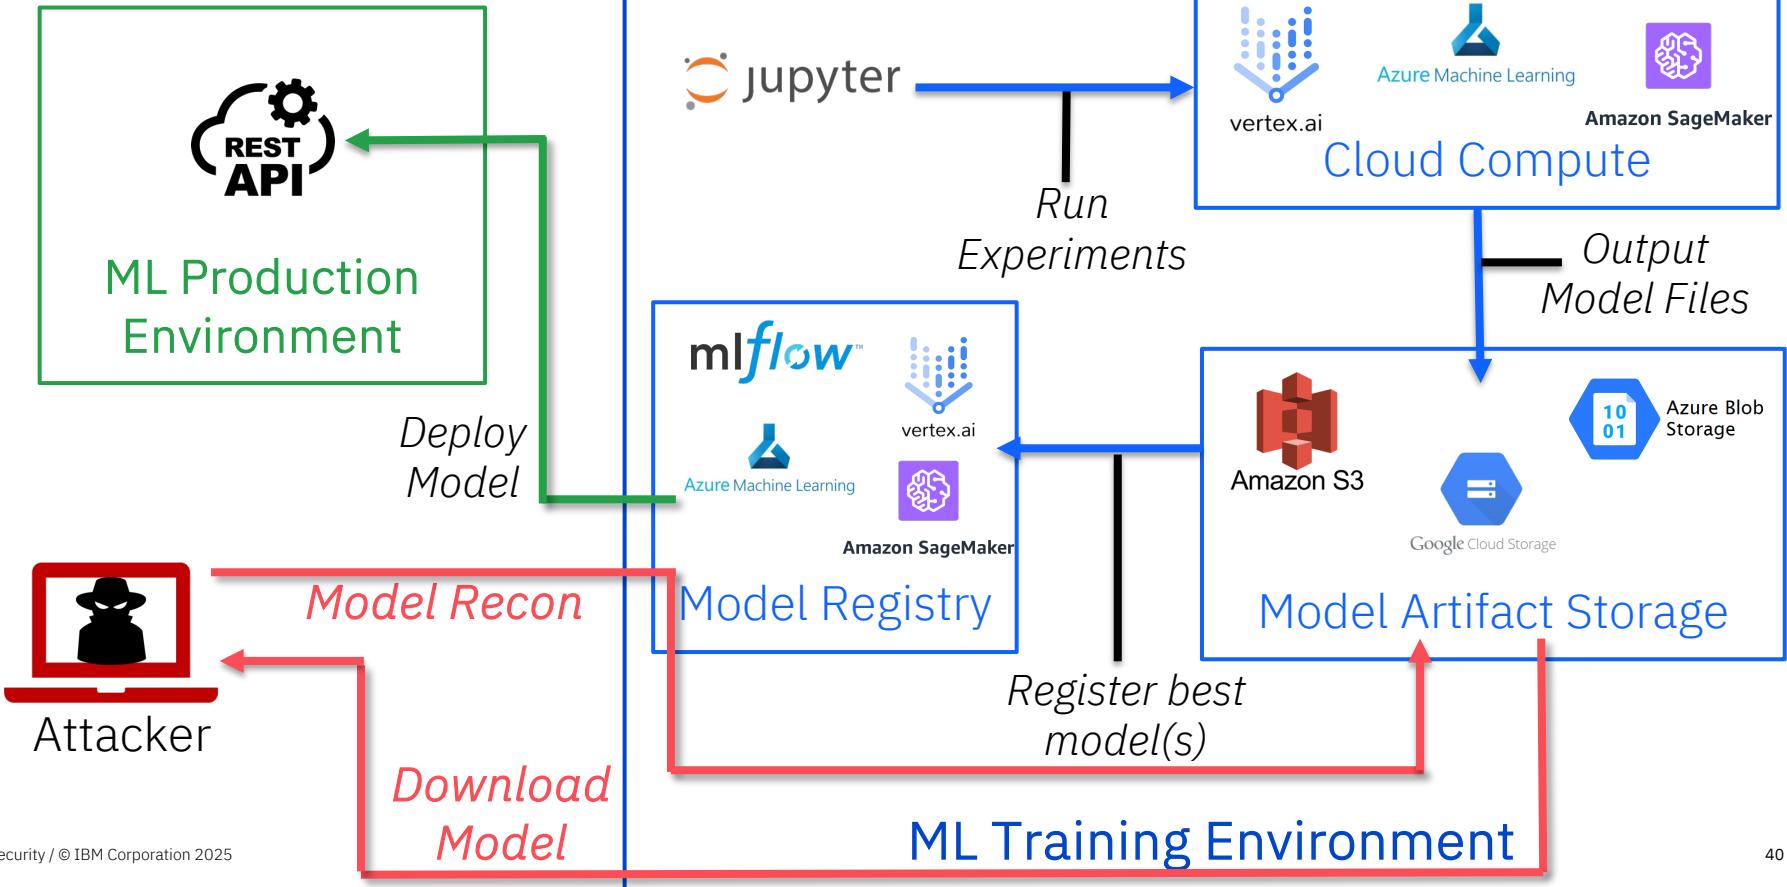
Model Theft



Model Theft



Model Theft



Model Theft - Impact

IP Theft

Model Extraction & Replication

Adversarial Attacks & Evasion

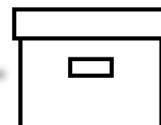
Backdooring

Compromise of System Security

Competitive Intelligence & Strategy Analysis

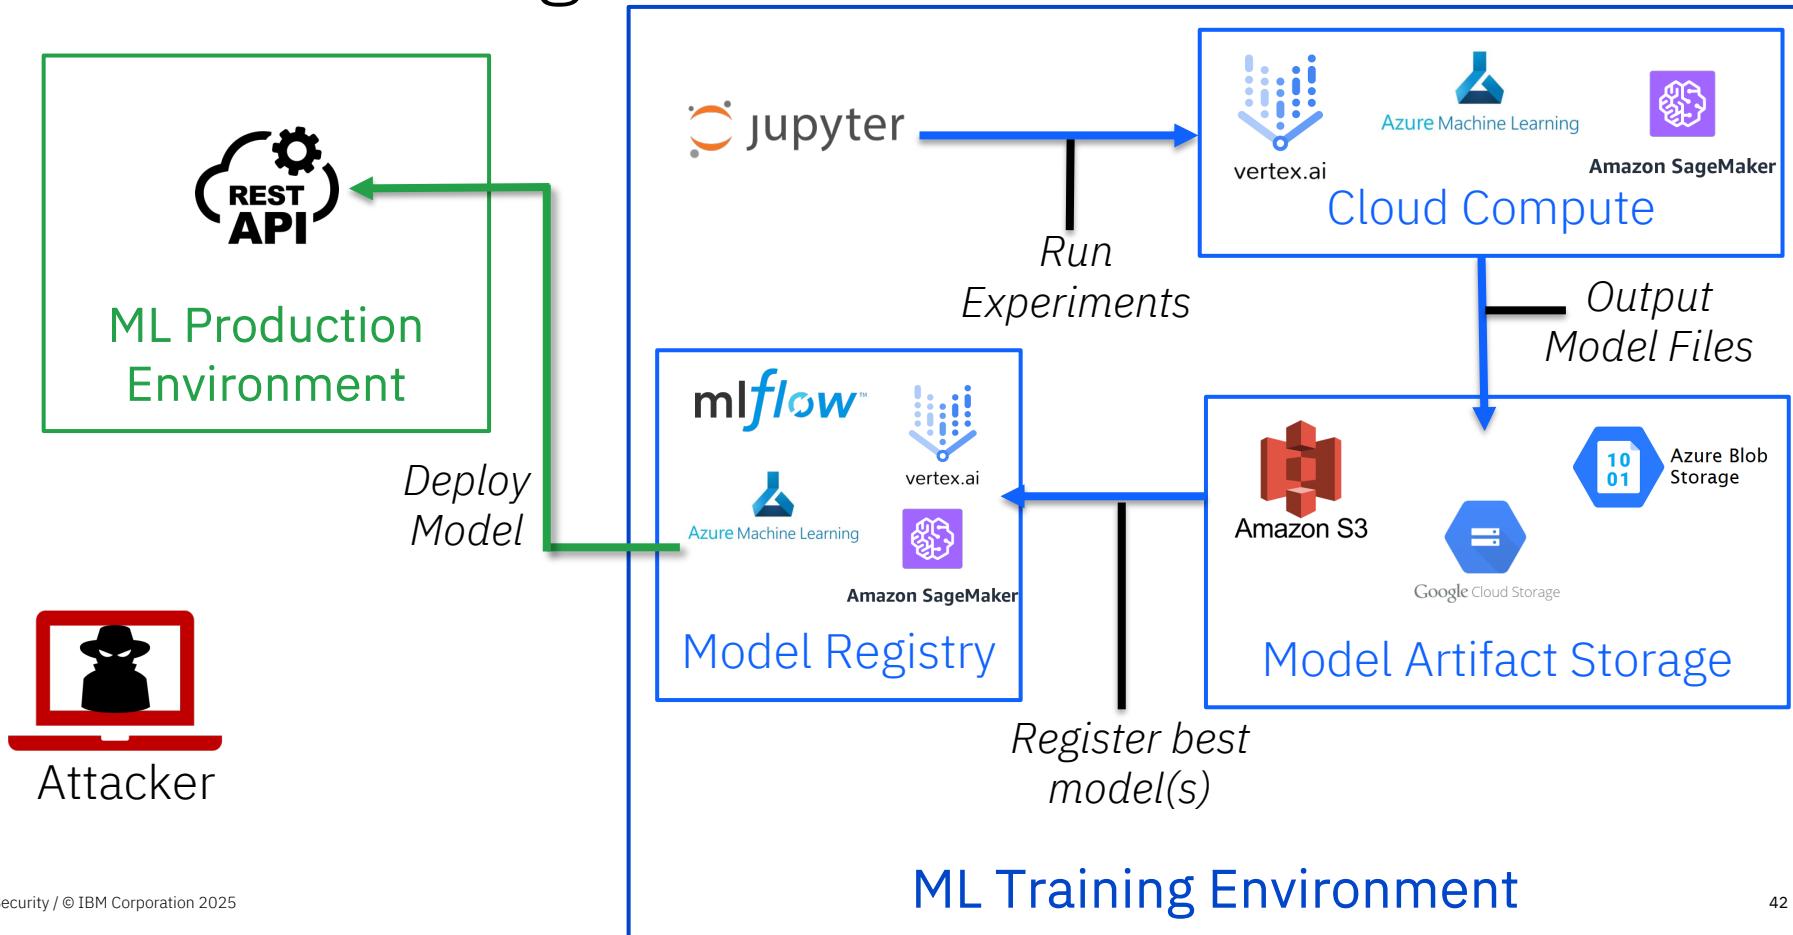


Attacker

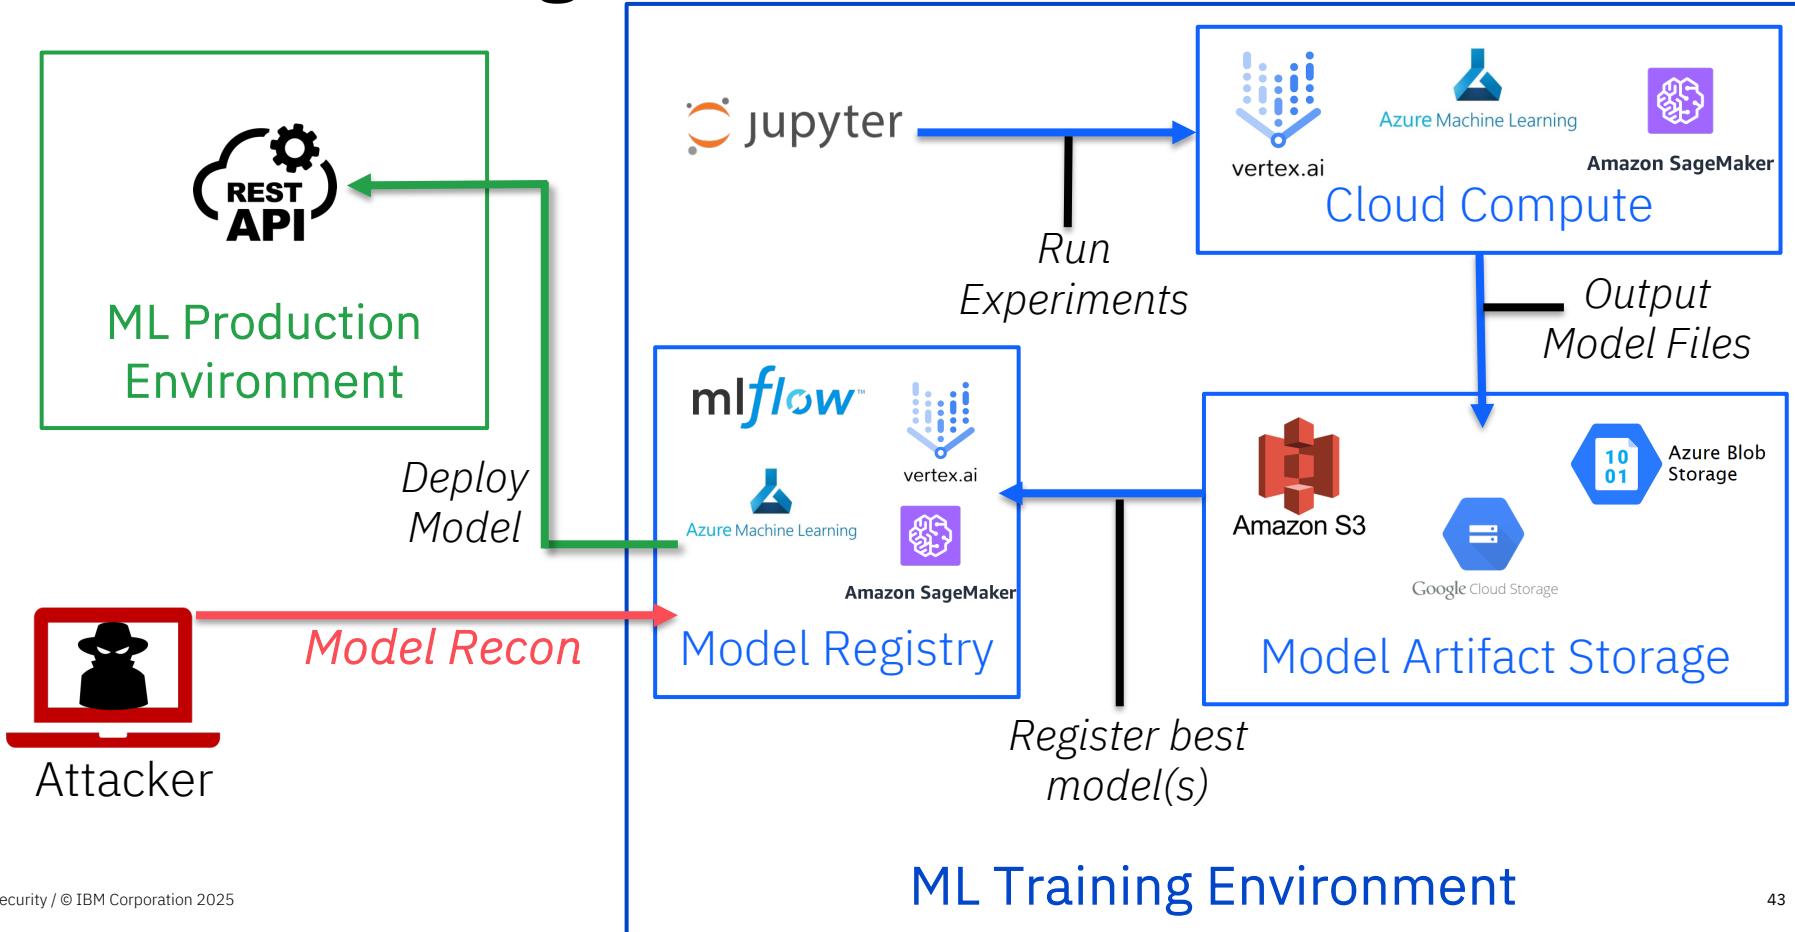


Stolen
Model Artifacts

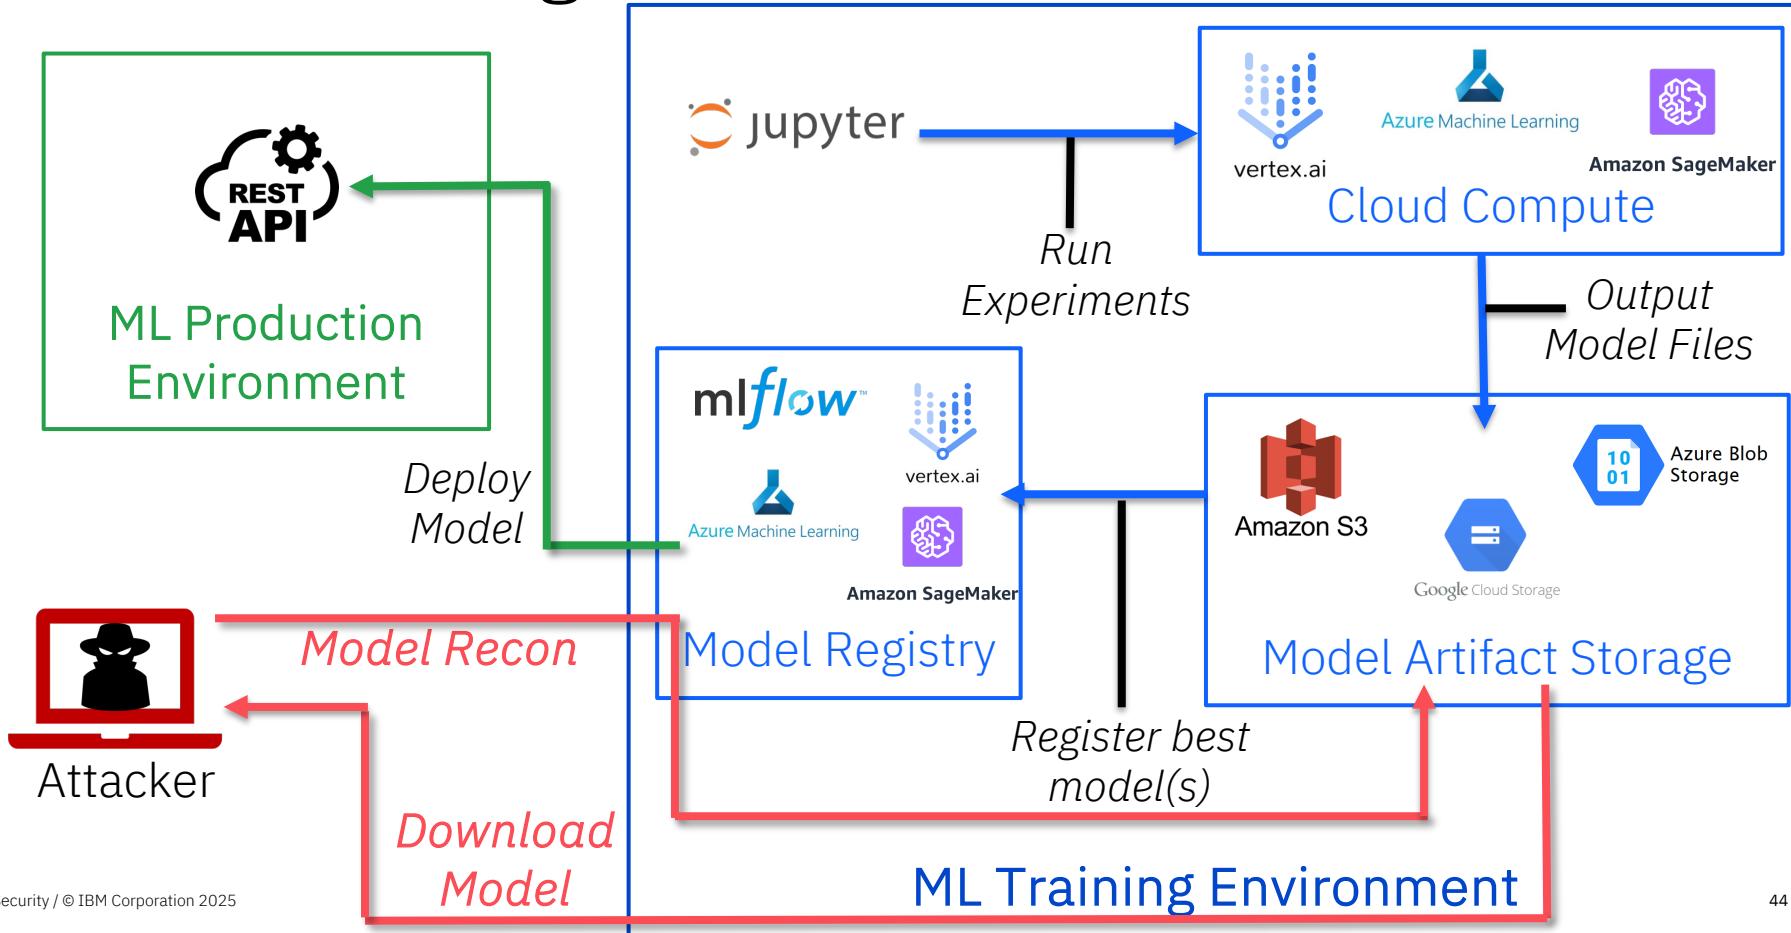
Model Poisoning – Code Execution



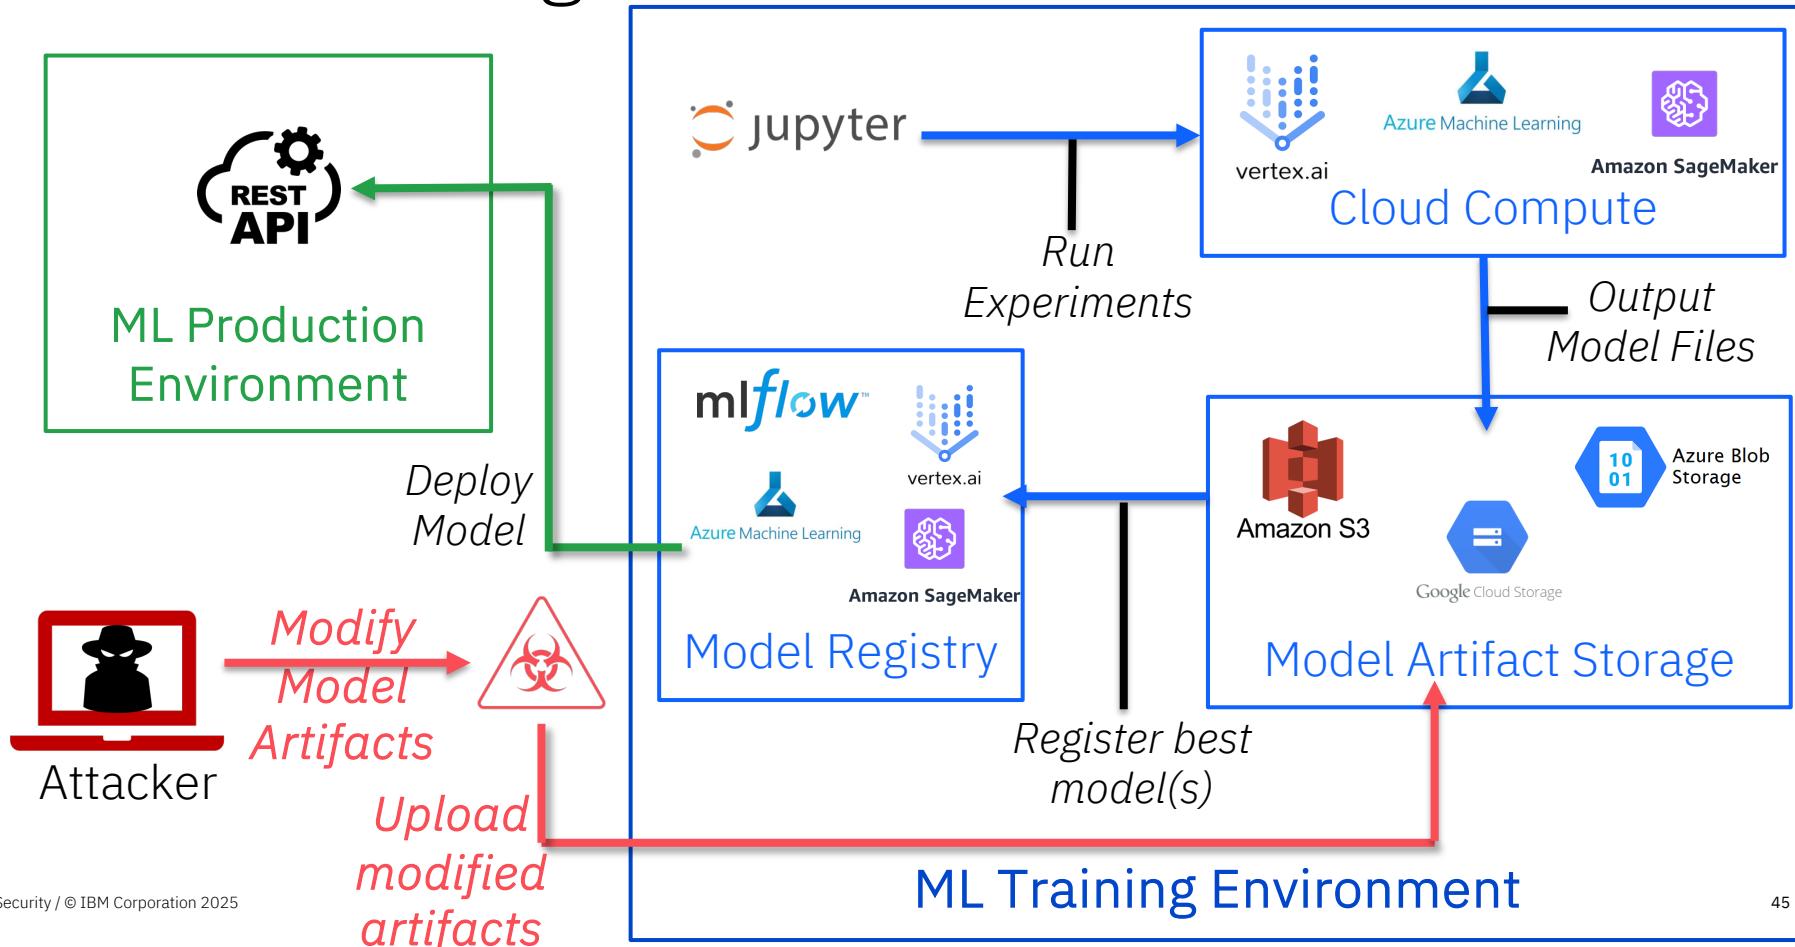
Model Poisoning – Code Execution



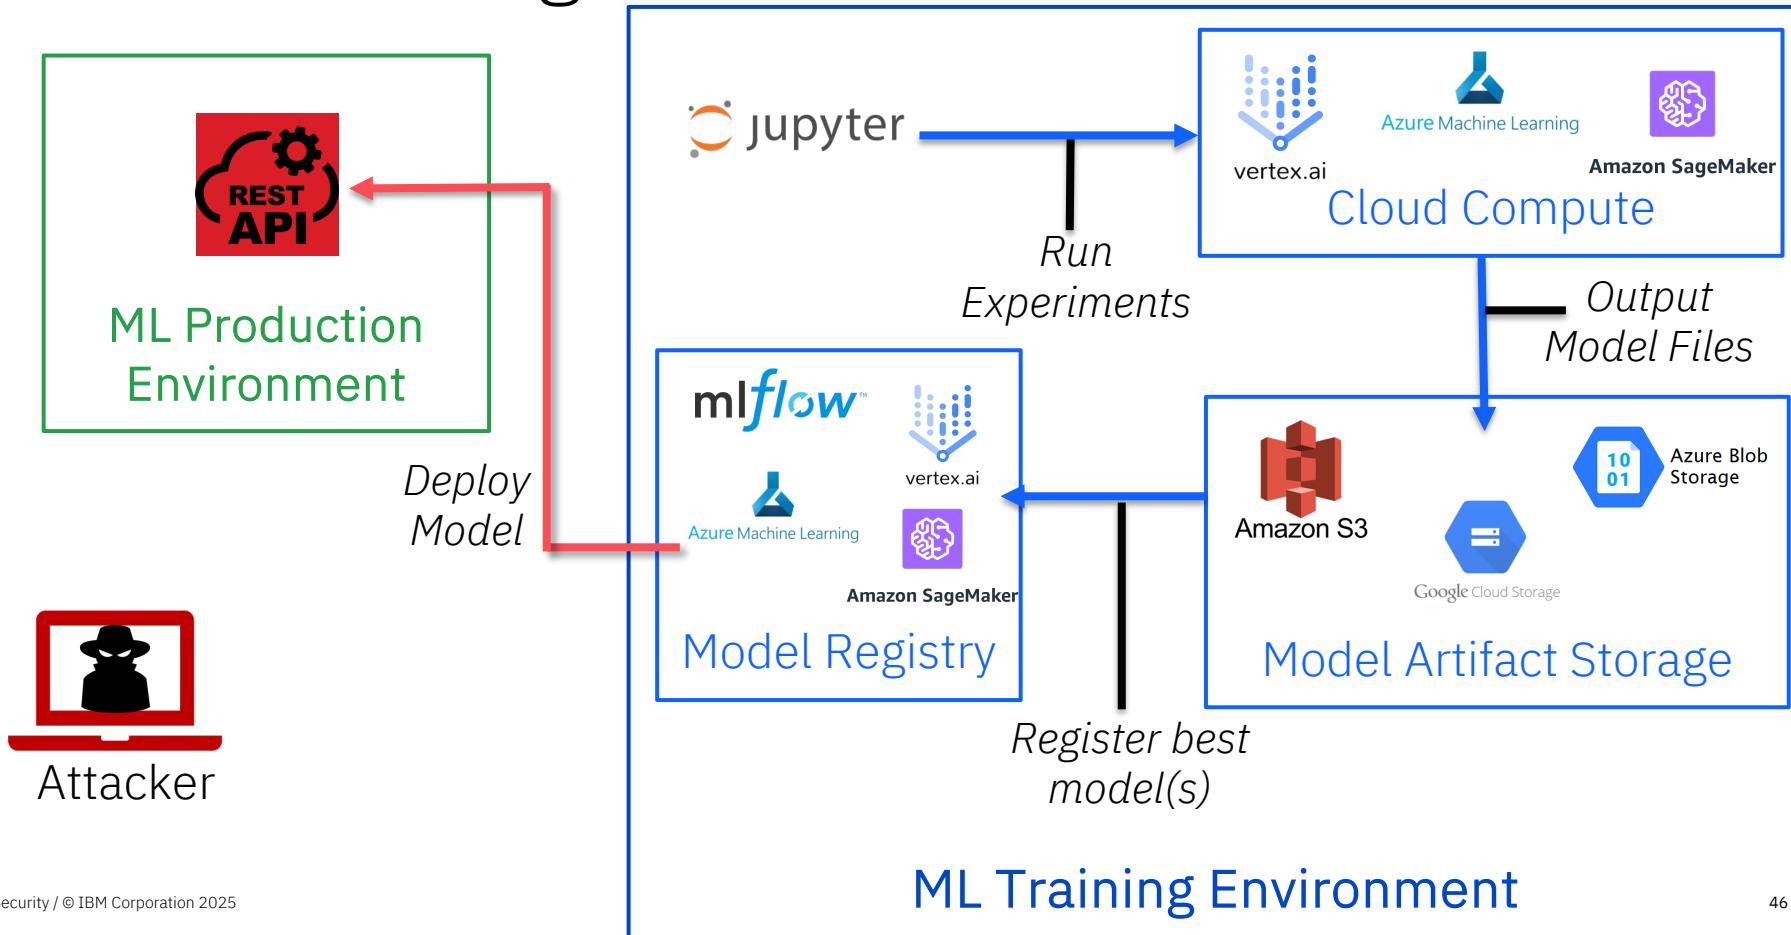
Model Poisoning – Code Execution



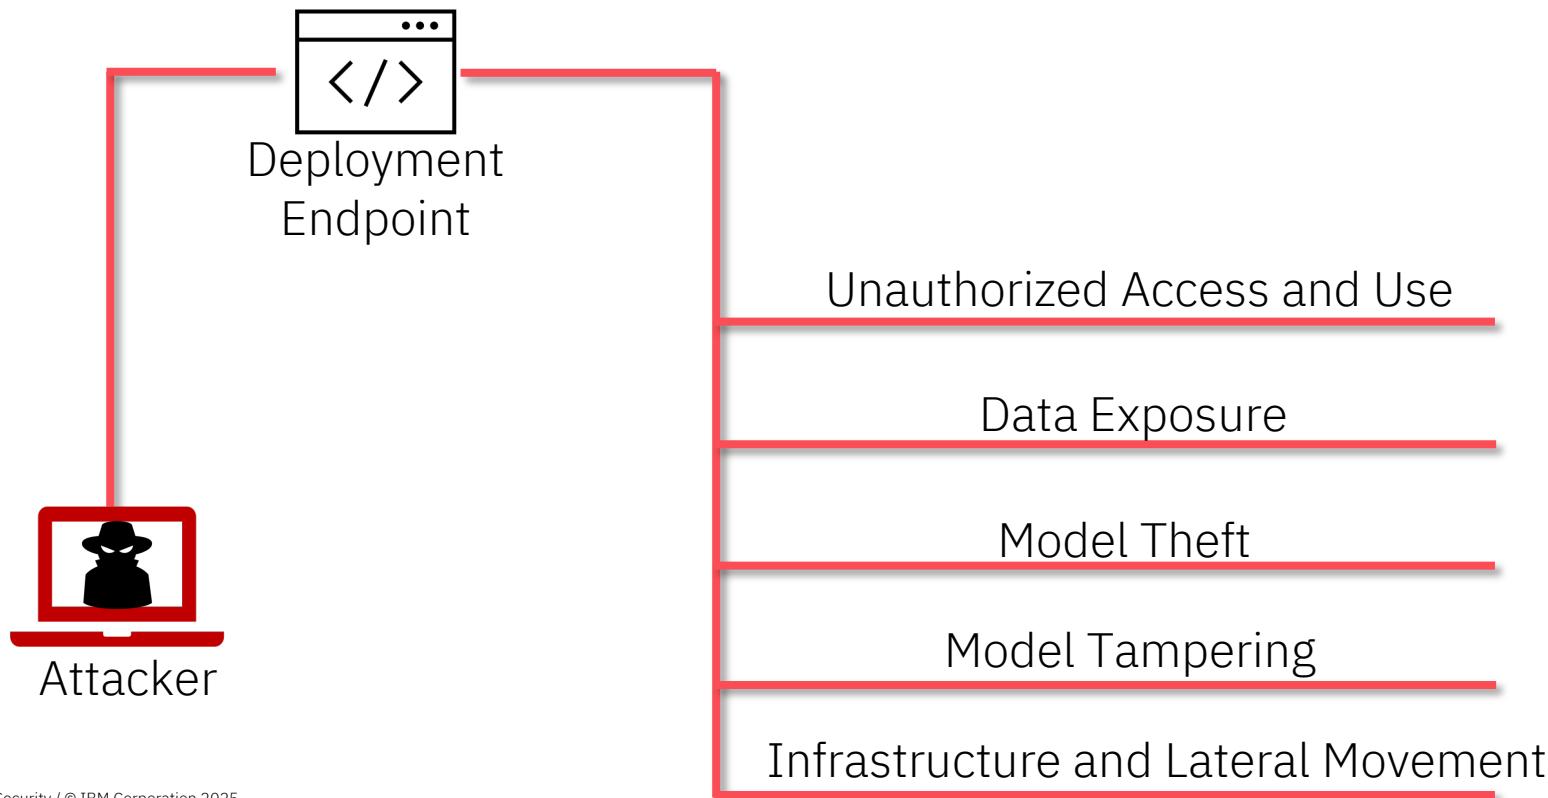
Model Poisoning – Code Execution



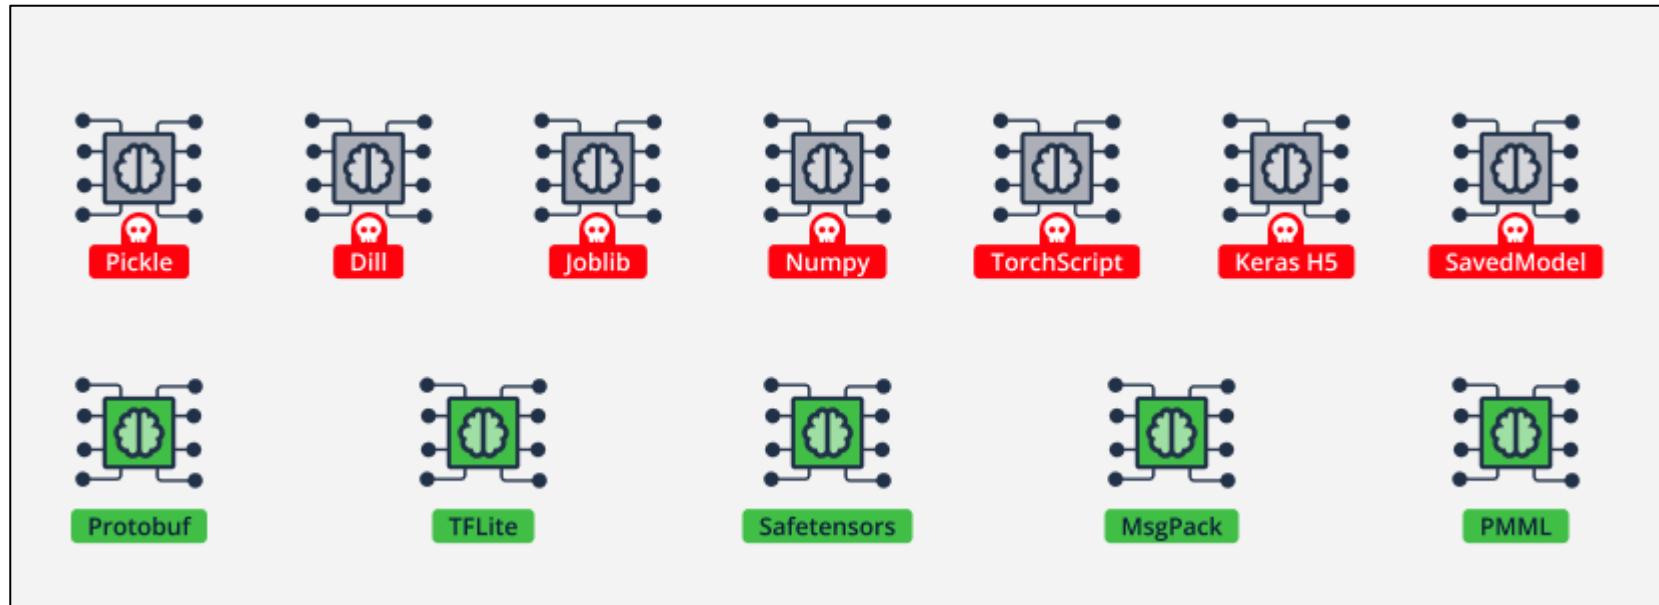
Model Poisoning – Code Execution



Model Poisoning – Code Execution - Impact



Model Formats – Support Code Execution on Load



<https://jfrog.com/blog/from-mlops-to-mloops-exposing-the-attack-surface-of-machine-learning-platforms/>

Creating Malicious Models for Code Execution

- **MaliciousPickles** - <https://github.com/coldwaterq/MaliciousPickles>
- **Charcuterie** - <https://github.com/moohax/Charcuterie>
- **Fickling** - <https://github.com/trailofbits/fickling>
- **HiddenPickle** - <https://github.com/hiddenlayerai/HiddenPickle>

MLOKit

github.com/xforced/MLOKit

```
[*] INFO: Performing download-model module for sagemaker
[*] INFO: Checking credentials provided
[+] SUCCESS: Credentials are valid

      Model Name |          Creation Date
-----|-----
            employee-salary-model | 1/22/2025

[*] INFO: Downloading model artifacts
[*] INFO: Model artifacts location
s3://sagemaker-us-east-2-339713155979/Canvas/default-20250121T171825/Train:1/model.tar.gz

[*] INFO: Checking access to S3 bucket with name: sagemaker-us-east-2-339713155979
[+] SUCCESS: You have access to S3 bucket with name: sagemaker-us-east-2-339713155979

[*] INFO: Listing all files in prefix of: Canvas/default-20250121T171825/Training/output/Canvas1737498359380/sagemaker
[*] INFO: Downloading file at: Canvas/default-20250121T171825/Training/output/Canvas1737498359380/sagemaker/model.tar.gz
[+] SUCCESS: model.tar.gz written to: C:\Demo\MLKit-OiEJQGbz
```



REST API Abuse
Conduct actions programmatically



Authentication
API Key, Access Token, Security Creds, User/Pass

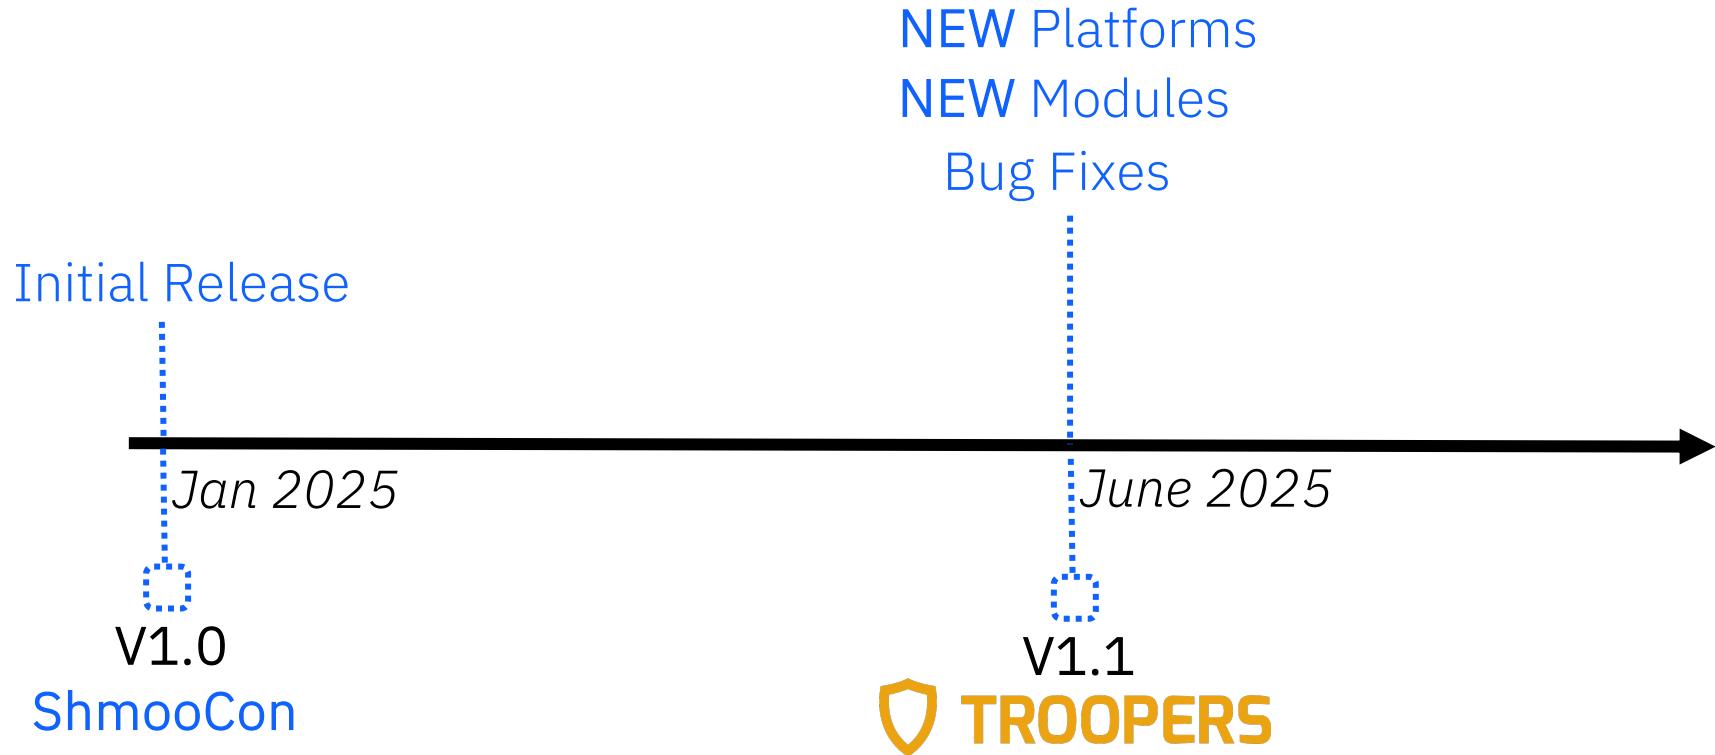


9 Modules
Recon, Training Data Theft, Model Theft, Model Poisoning, Notebook Attacks

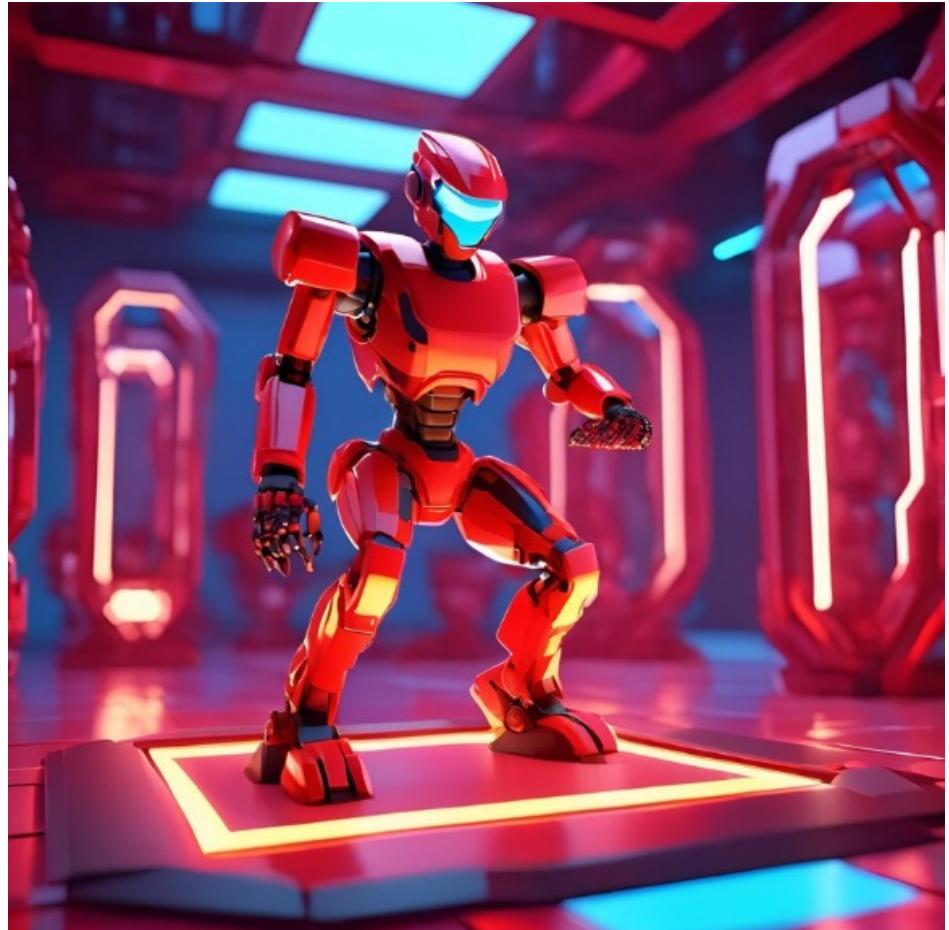


5 Supported Platforms
Azure ML, BigML, Vertex AI, MLFlow, SageMaker

MLOKit - History



Demos: Attack Scenarios



Obtaining Credentials

File Shares

Intranet Sites (e.g., internal wikis)

User Workstations

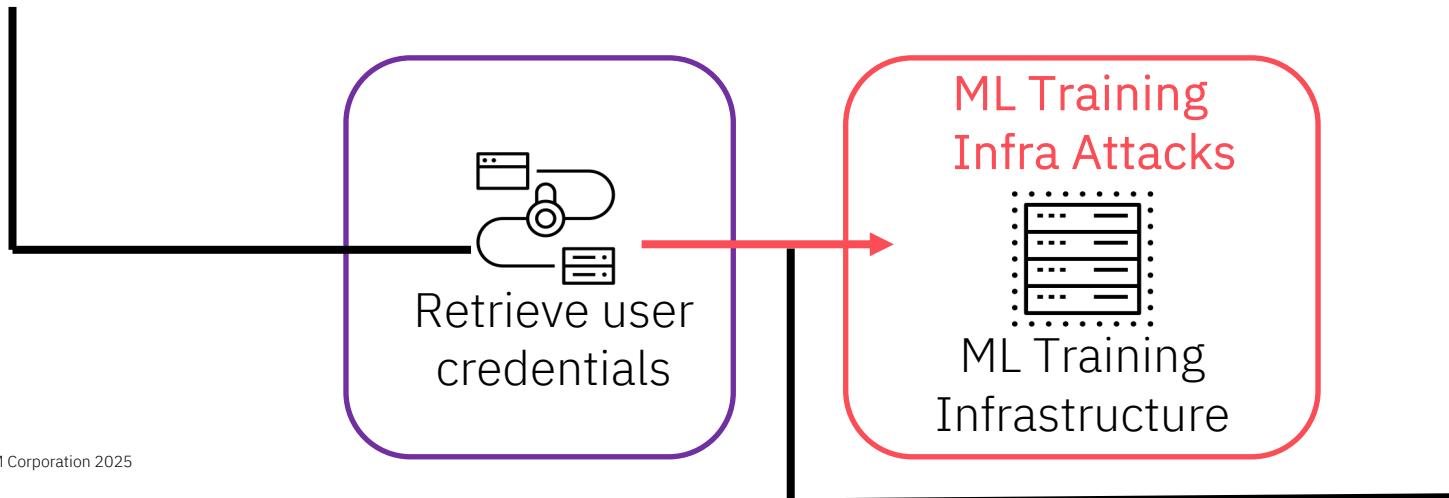
Social Engineering

Public Resources (e.g., Code Repos)

Unauthenticated Access

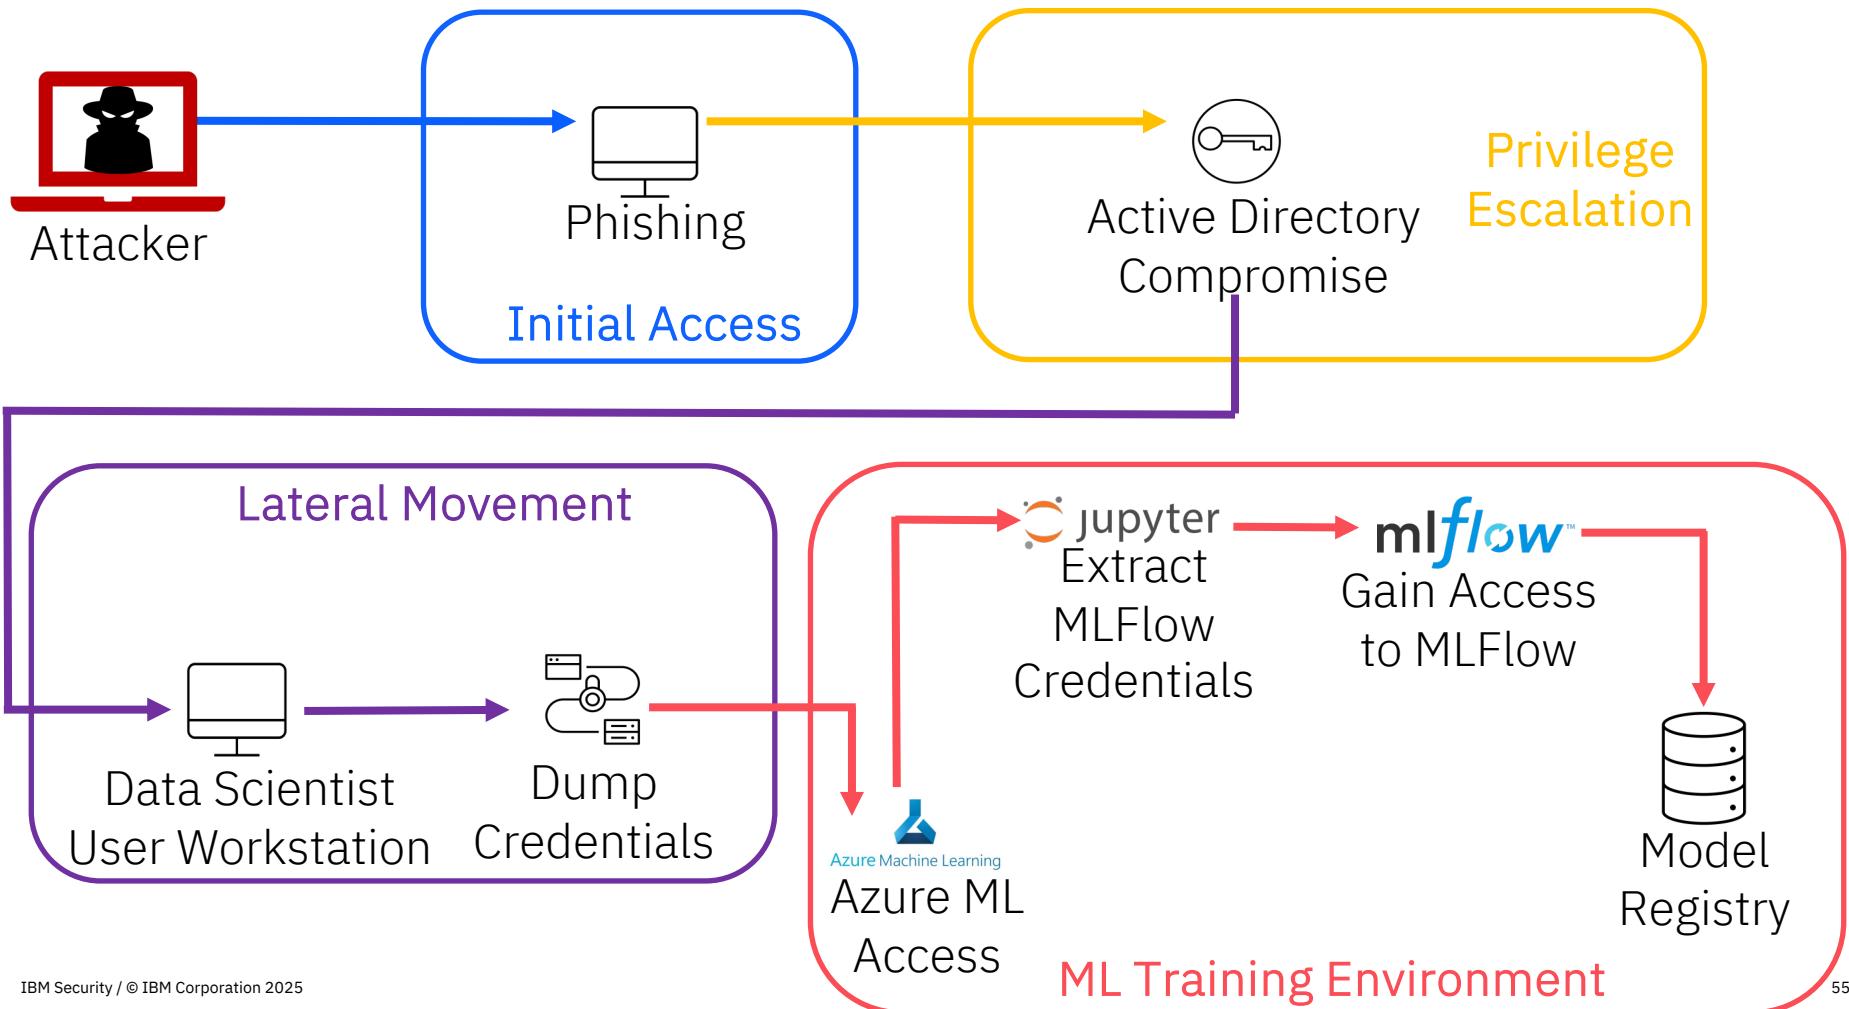
Public Data Breach Leaks

This research focuses
on attack paths
possible from here



Demos: Attack Scenarios

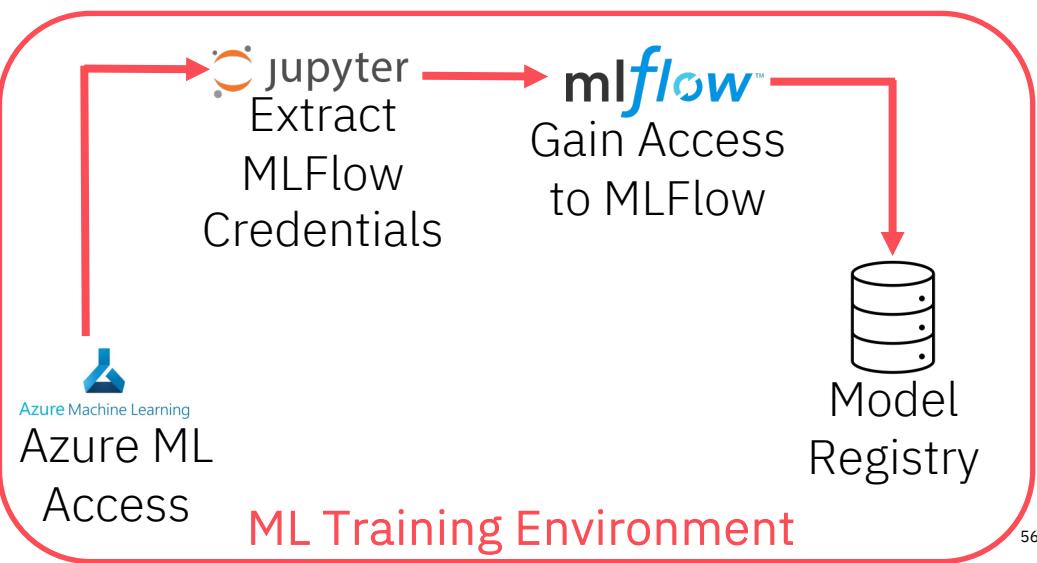
- #1: MLFlow - Initial Access and Model Theft from Model Registry
- #2: SageMaker - Lateral Movement from SCM System to Cloud Compute
- #3: SageMaker - Lateral Movement to Cloud Compute using Malicious Lifecycle Configuration
- #4: SageMaker - Model Theft from Model Registry
- #5: SageMaker - Model Poisoning to gain Code Execution
- #6: Azure ML - Model Poisoning to gain Code Execution





Attacker

Demo





C:\Demo>|



Demos: Attack Scenarios

#1: MLFlow - Initial Access and Model Theft from Model Registry

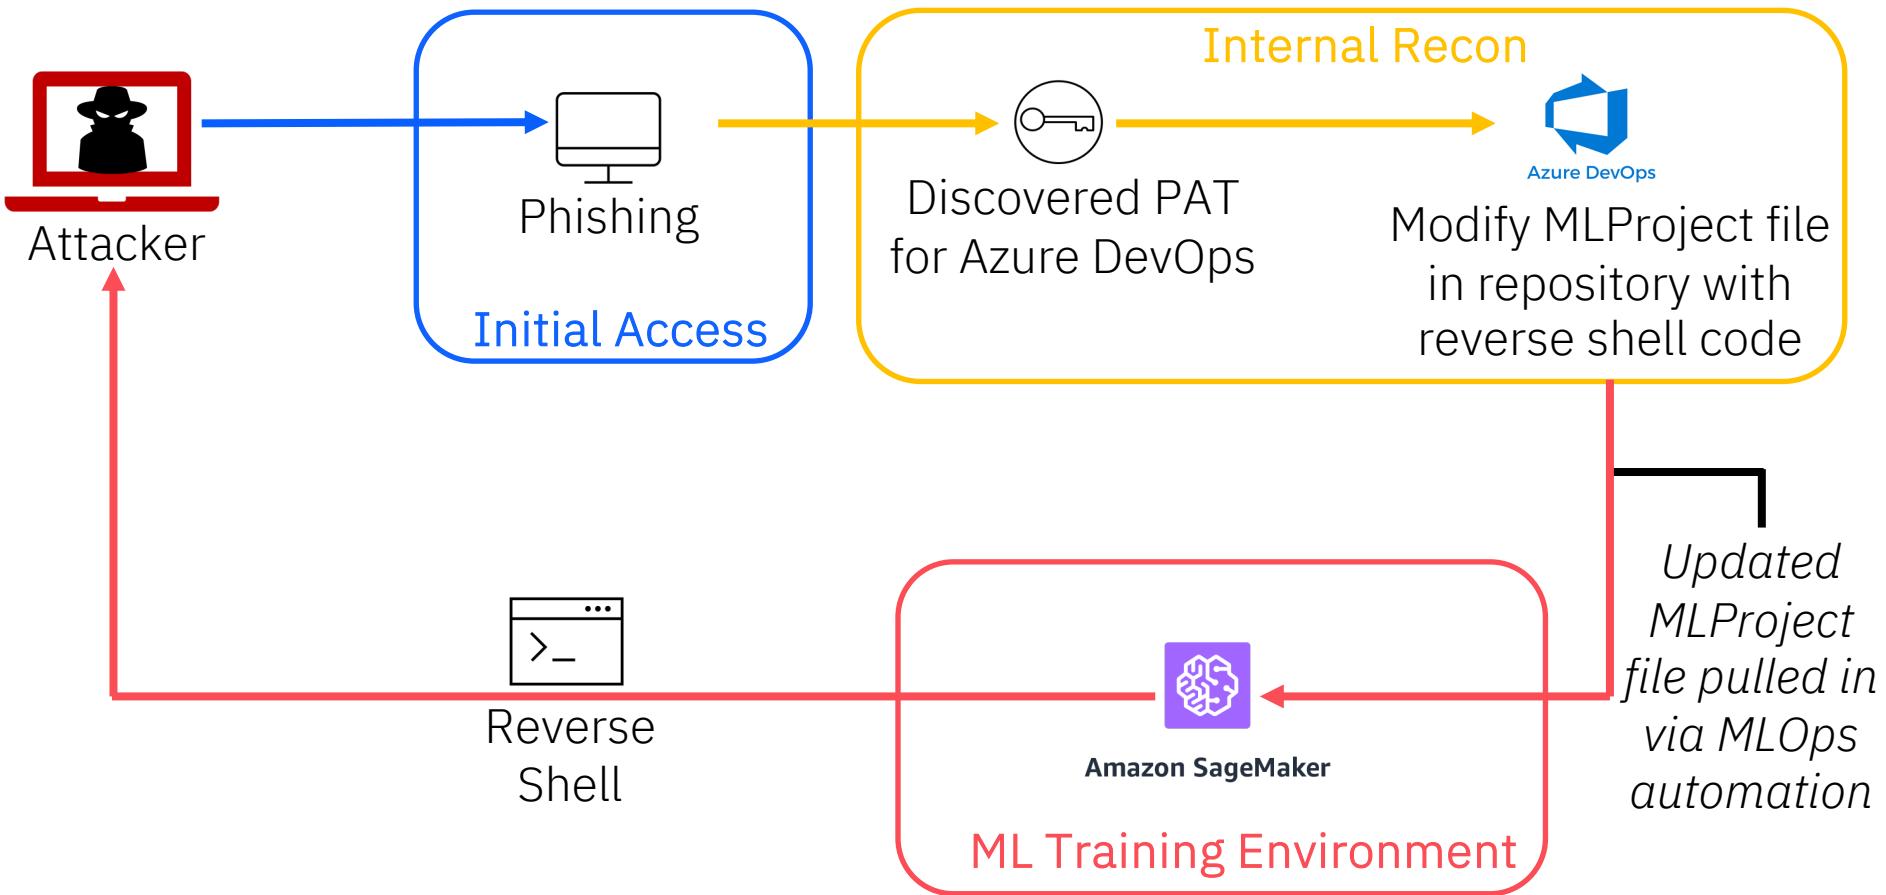
#2: SageMaker - Lateral Movement from SCM System to Cloud Compute

#3: SageMaker - Lateral Movement to Cloud Compute using Malicious Lifecycle Configuration

#4: SageMaker - Model Theft from Model Registry

#5: SageMaker - Model Poisoning to gain Code Execution

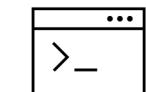
#6: Azure ML - Model Poisoning to gain Code Execution





Attacker

Demo



Reverse
Shell

Internal Recon



Modify MLProject file
in repository with
reverse shell code



Amazon SageMaker

ML Training Environment

*Updated
MLProject
file pulled in
via MLOps
automation*



Clone ML training code repo using stolen PAT

Demos: Attack Scenarios

#1: MLFlow - Initial Access and Model Theft from Model Registry

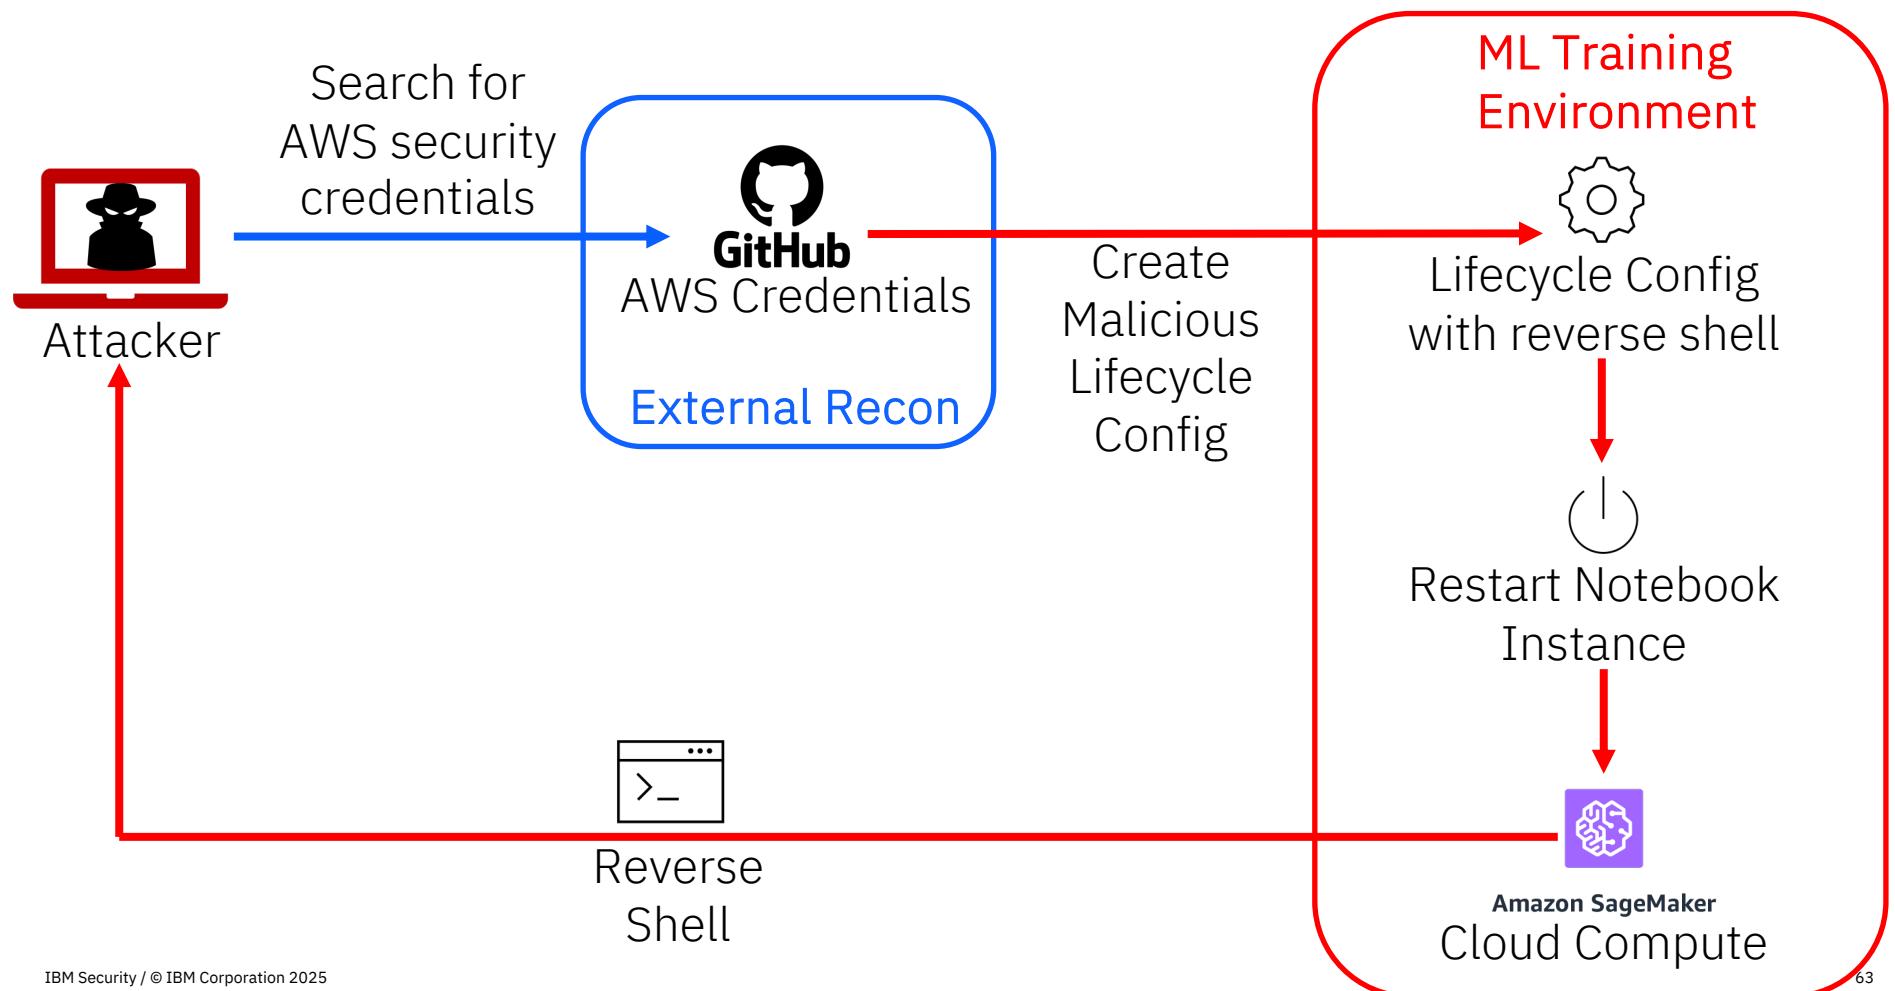
#2: SageMaker - Lateral Movement from SCM System to Cloud Compute

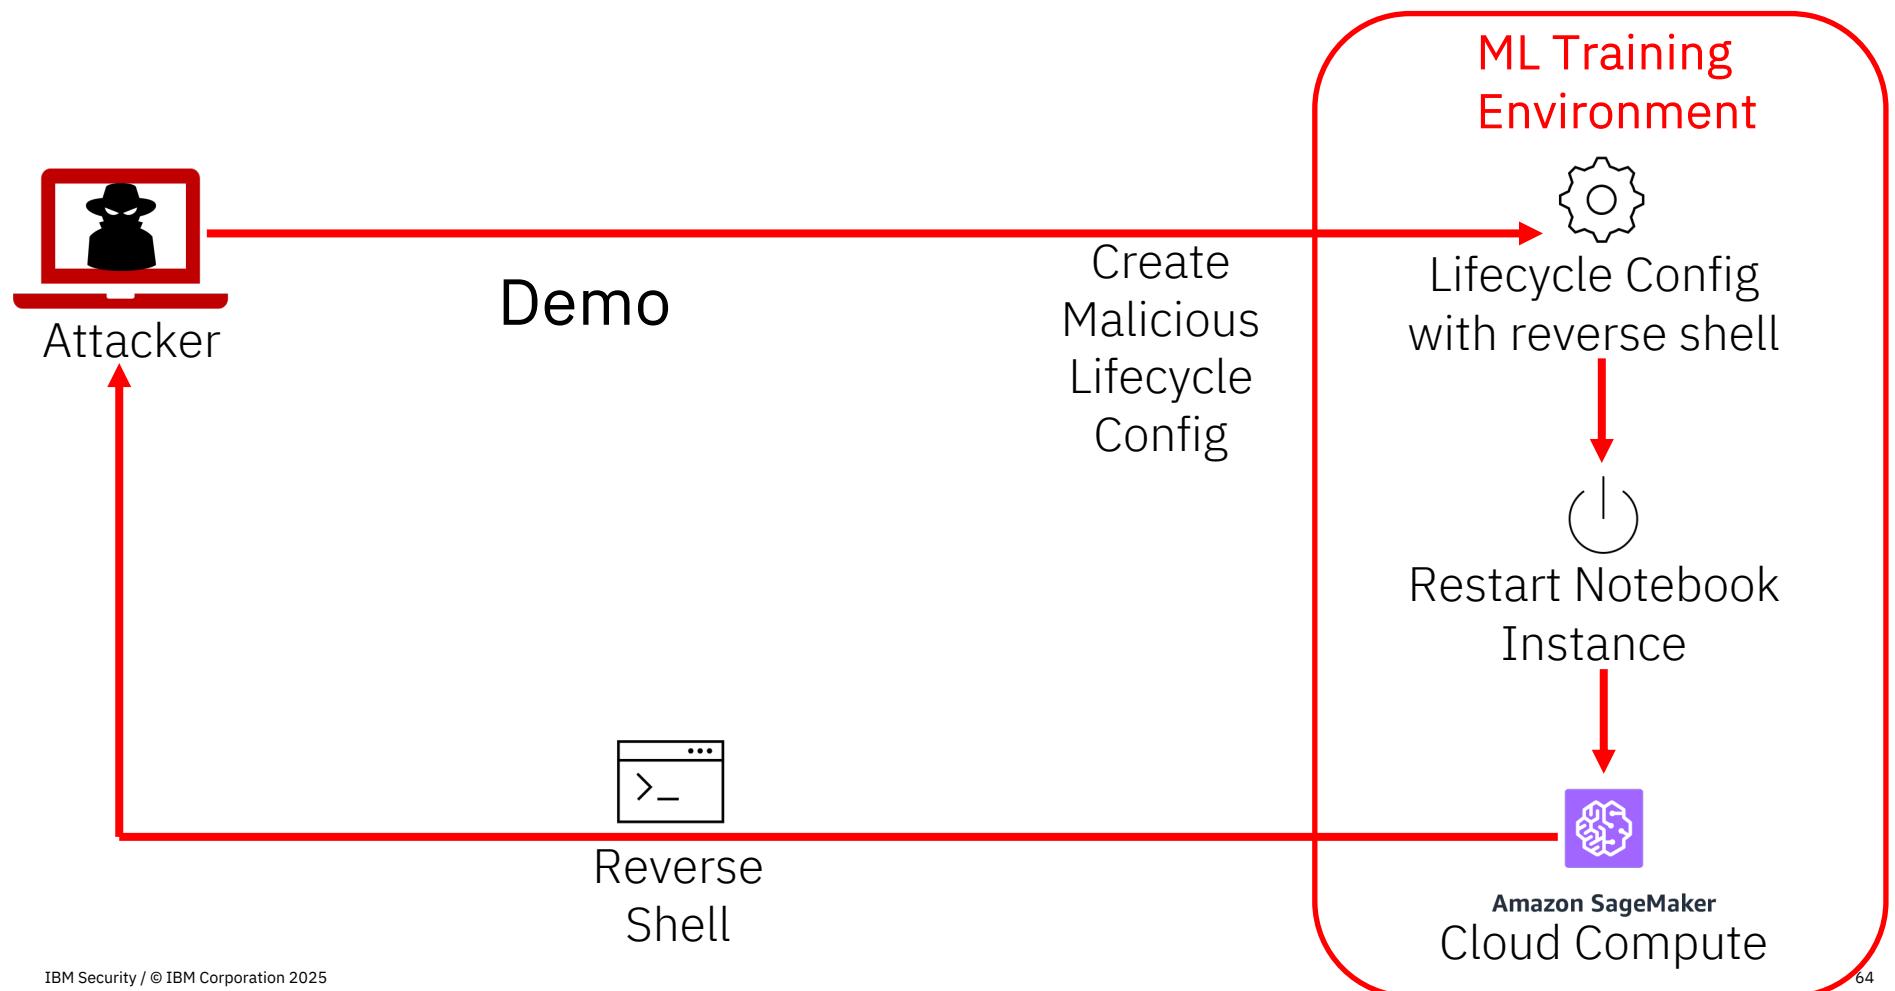
#3: SageMaker - Lateral Movement to Cloud Compute using Malicious Lifecycle Configuration

#4: SageMaker - Model Theft from Model Registry

#5: SageMaker - Model Poisoning to gain Code Execution

#6: Azure ML - Model Poisoning to gain Code Execution







C:\Demo>



Demos: Attack Scenarios

#1: MLFlow - Initial Access and Model Theft from Model Registry

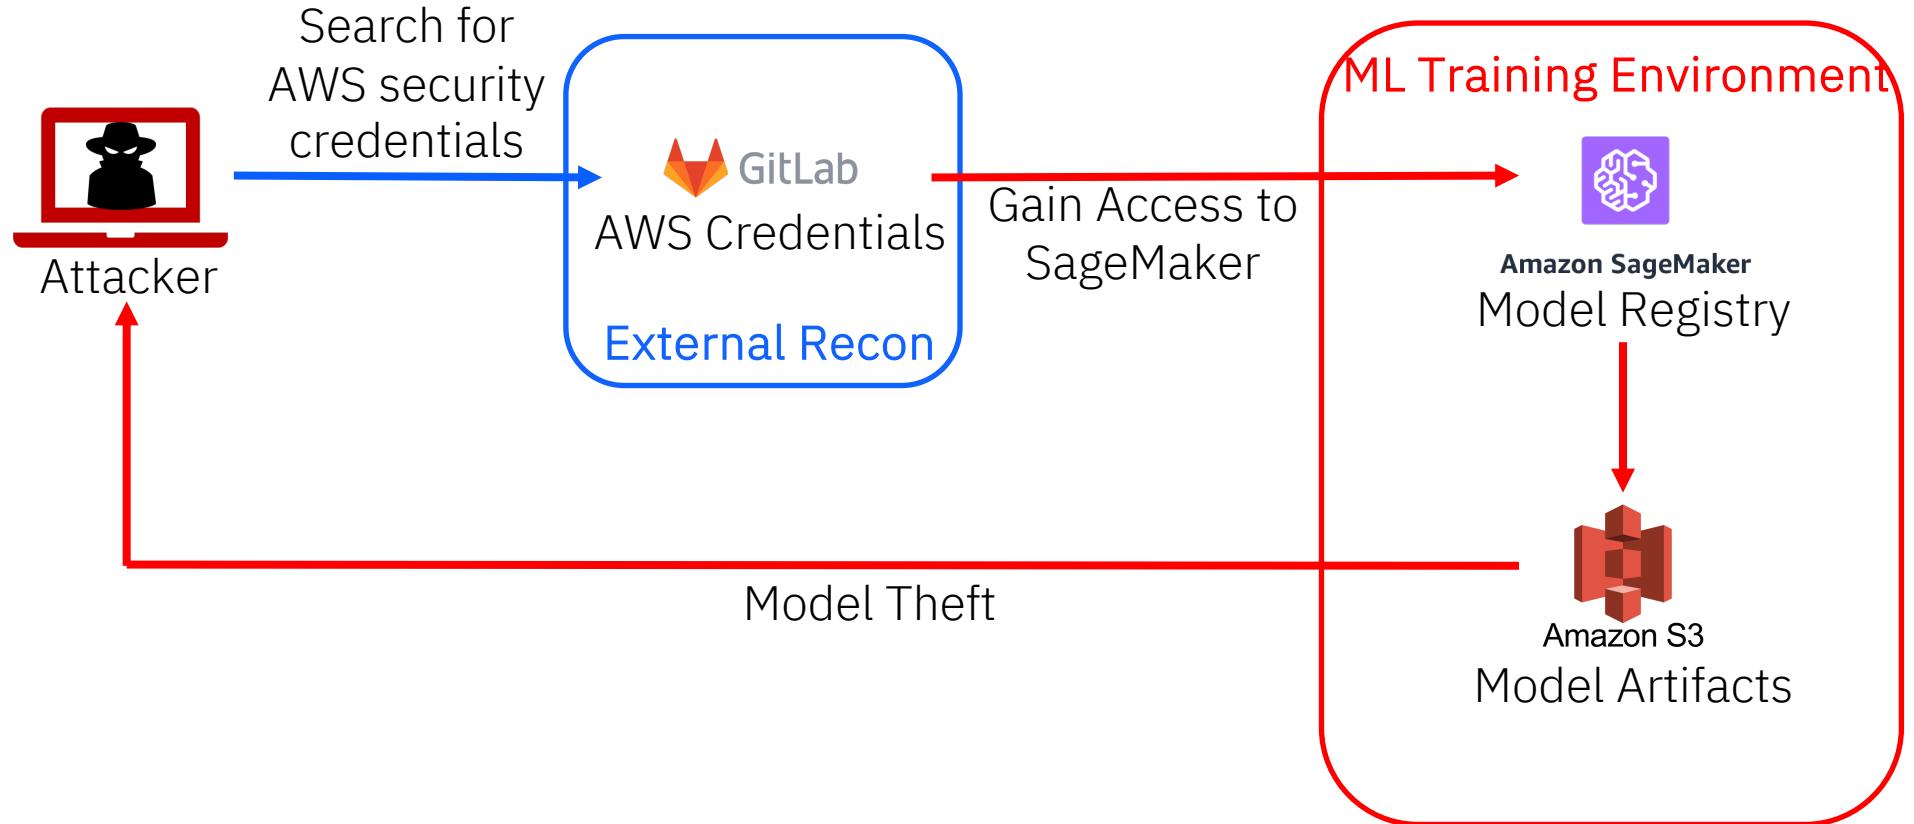
#2: SageMaker - Lateral Movement from SCM System to Cloud Compute

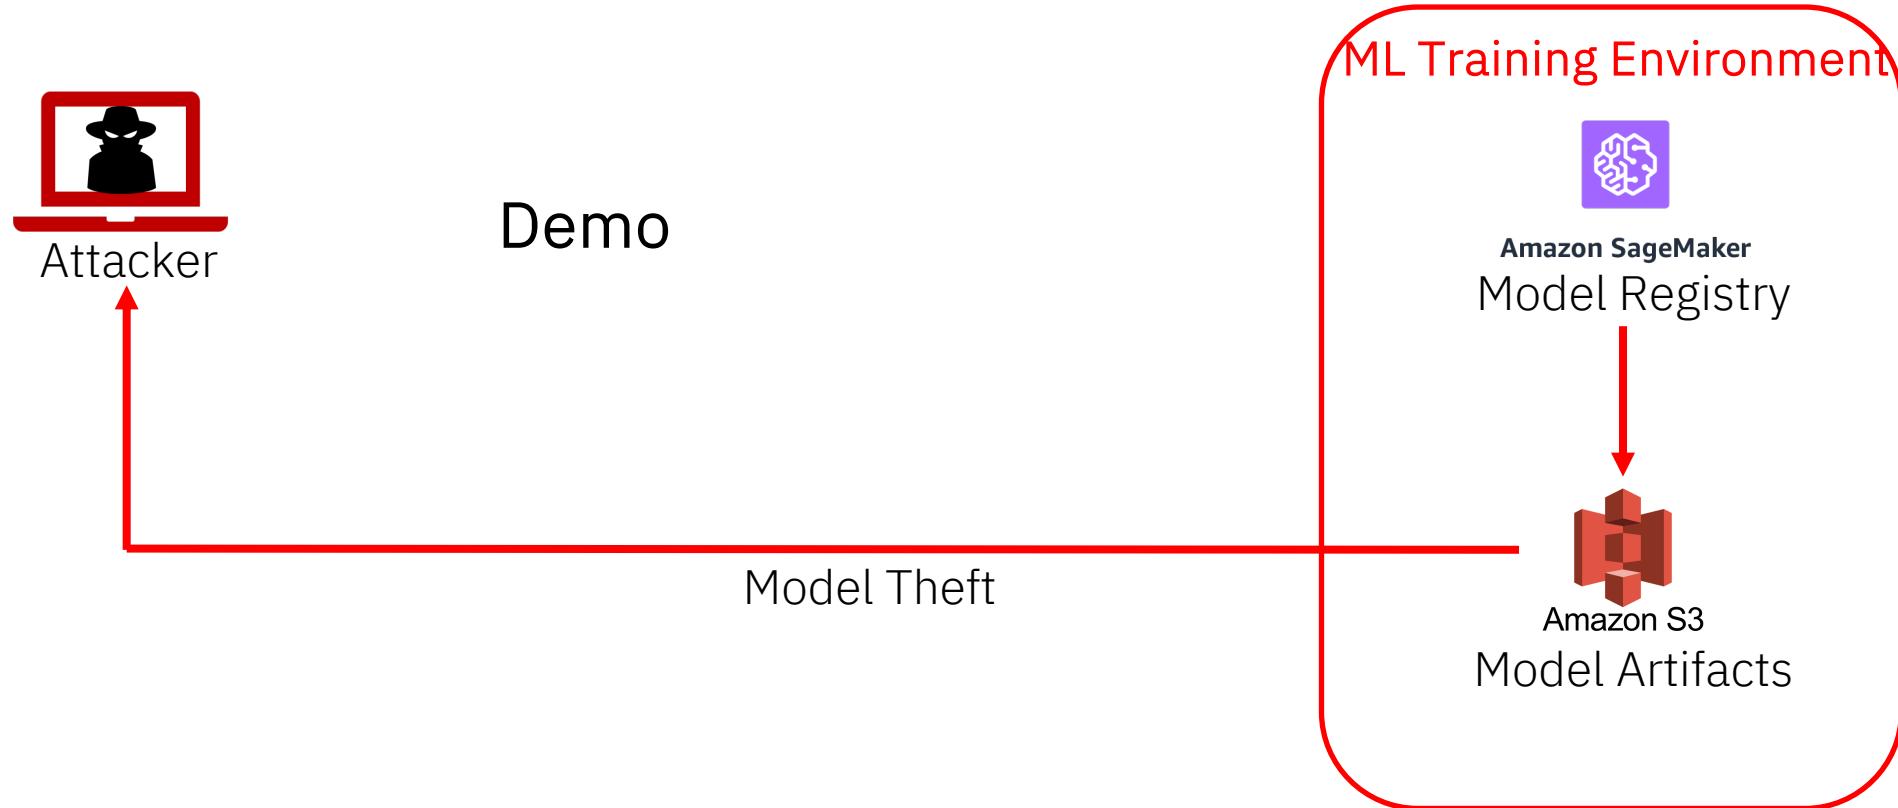
#3: SageMaker - Lateral Movement to Cloud Compute using Malicious Lifecycle Configuration

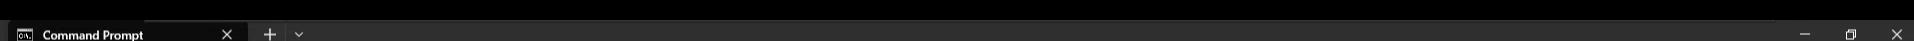
#4: SageMaker - Model Theft from Model Registry

#5: SageMaker - Model Poisoning to gain Code Execution

#6: Azure ML - Model Poisoning to gain Code Execution



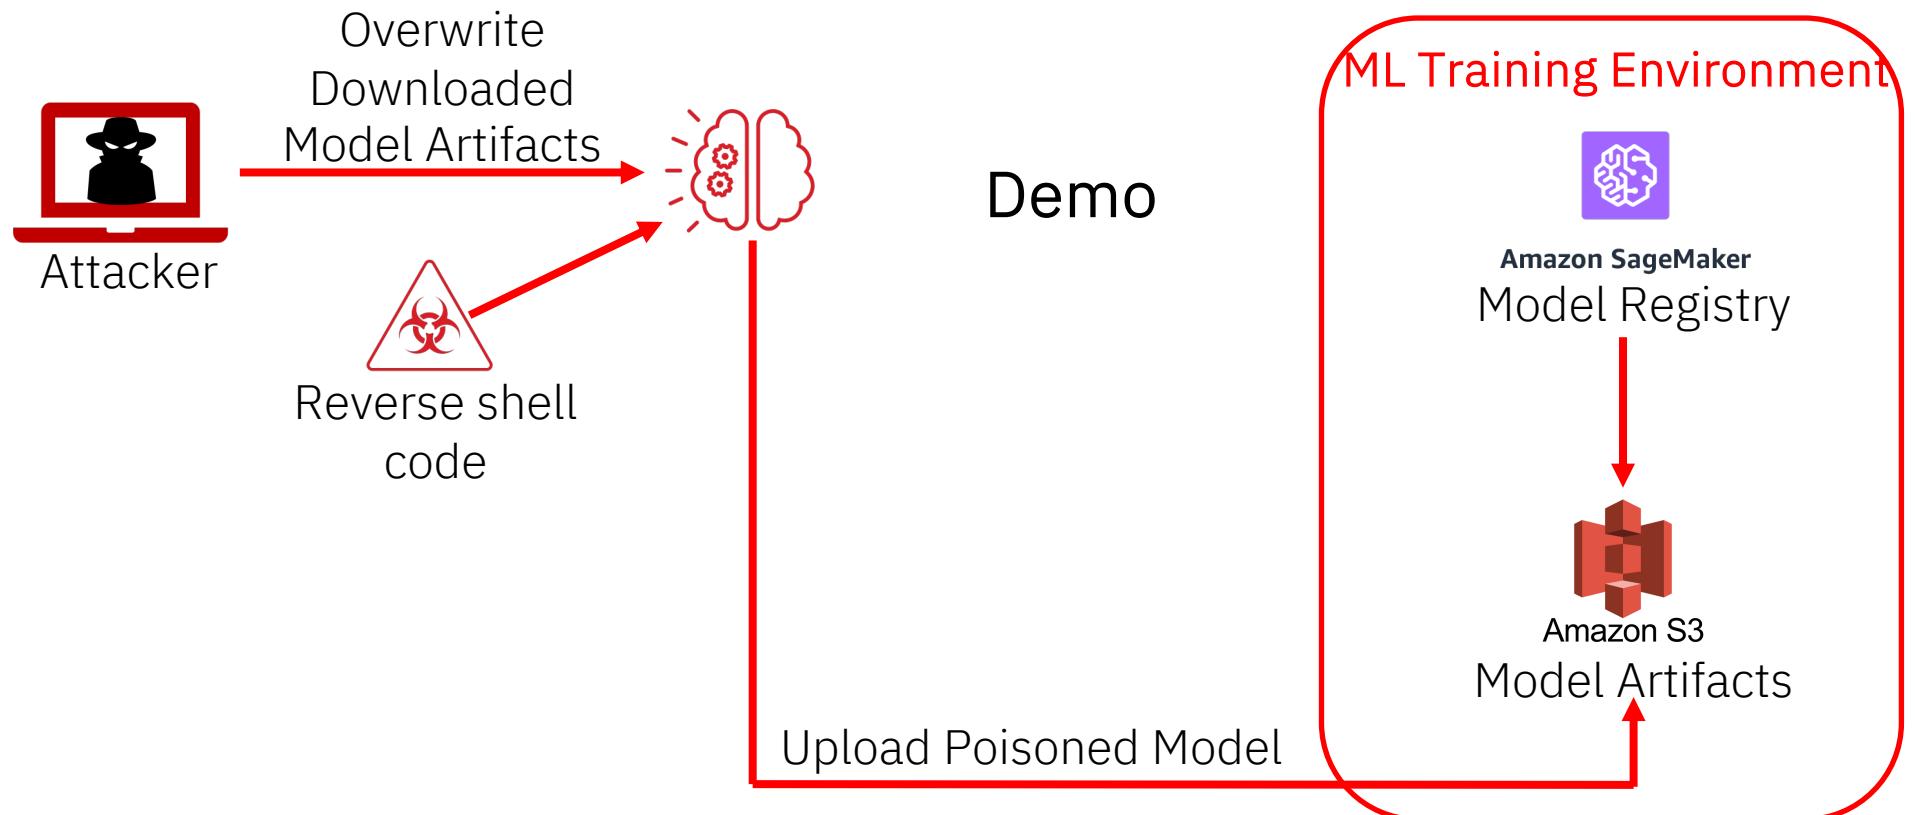




C:\Demo>

Demos: Attack Scenarios

- #1: MLFlow - Initial Access and Model Theft from Model Registry
- #2: SageMaker - Lateral Movement from SCM System to Cloud Compute
- #3: SageMaker - Lateral Movement to Cloud Compute using Malicious Lifecycle Configuration
- #4: SageMaker - Model Theft from Model Registry
- #5: SageMaker - Model Poisoning to gain Code Execution**
- #6: Azure ML - Model Poisoning to gain Code Execution





C:\Demo>



**Show ing model
artifacts of model
we previously
down loaded**

Demos: Attack Scenarios

#1: MLFlow - Initial Access and Model Theft from Model Registry

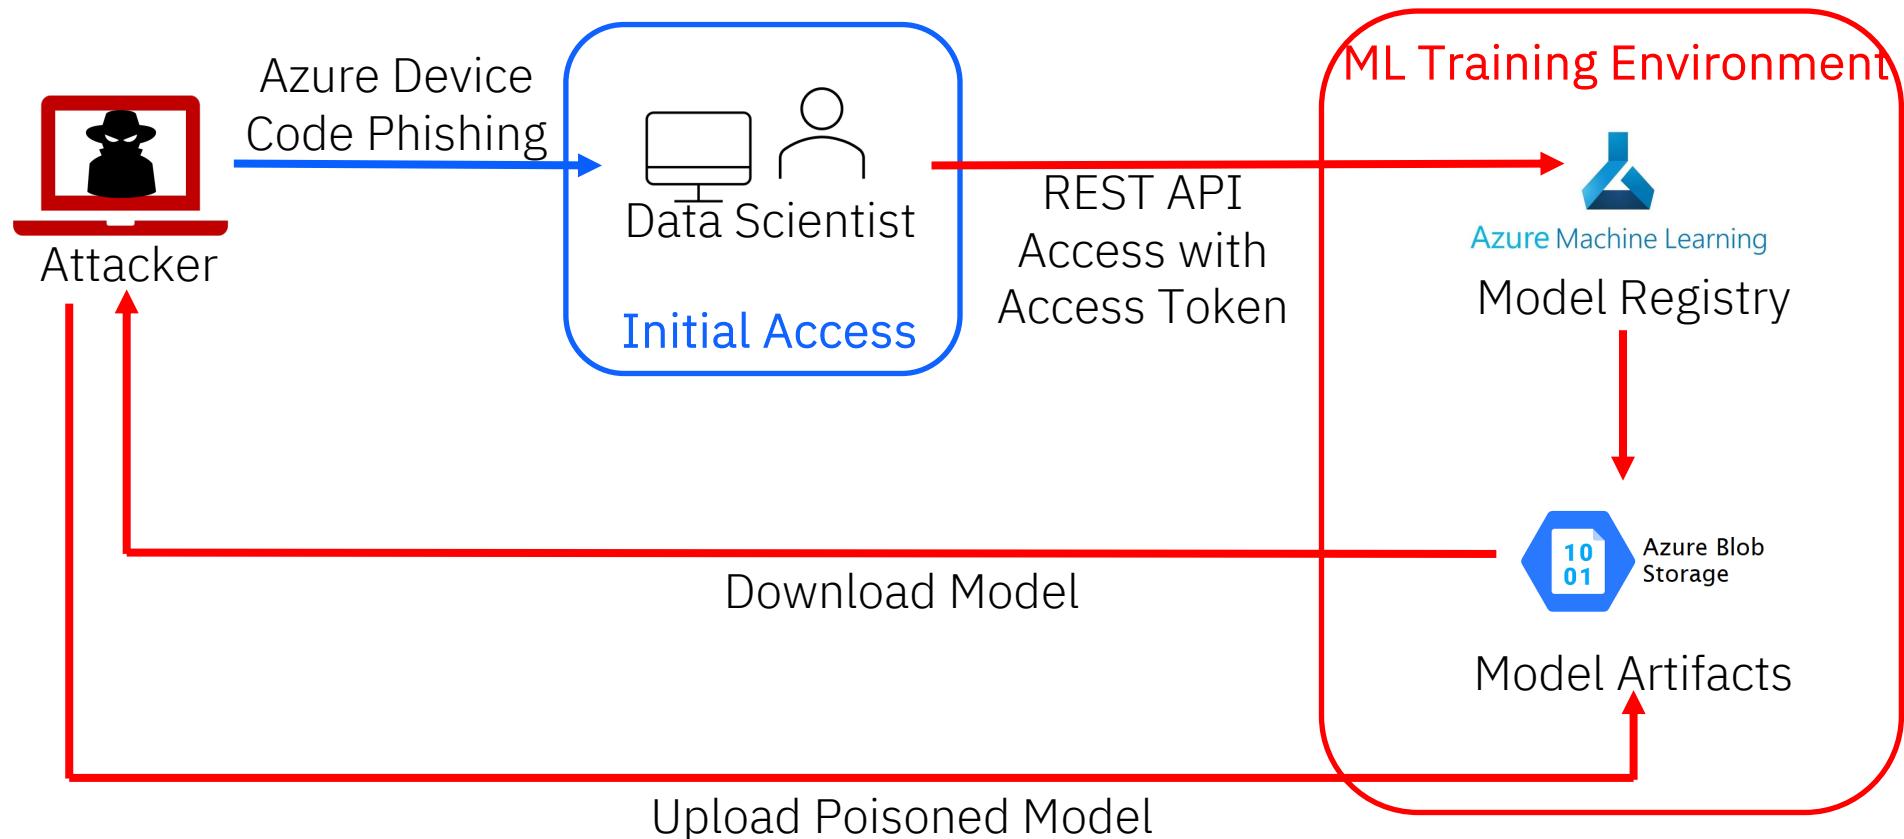
#2: SageMaker - Lateral Movement from SCM System to Cloud Compute

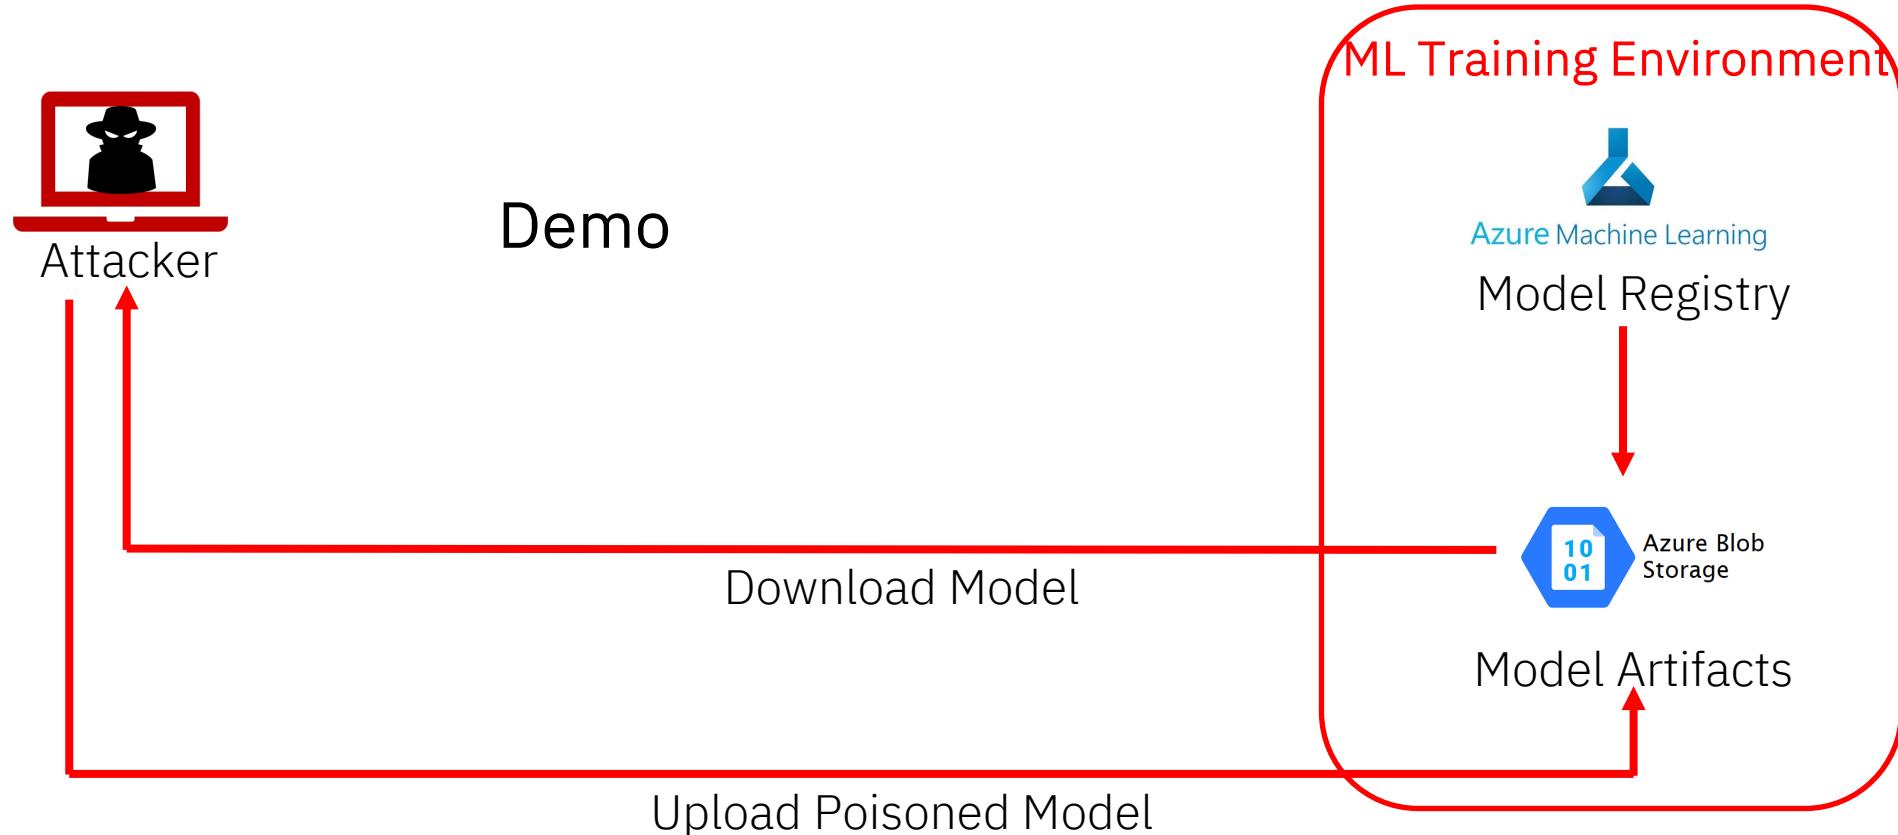
#3: SageMaker - Lateral Movement to Cloud Compute using Malicious Lifecycle Configuration

#4: SageMaker - Model Theft from Model Registry

#5: SageMaker - Model Poisoning to gain Code Execution

#6: Azure ML - Model Poisoning to gain Code Execution





C:\Users\hawk\Desktop\Demo.txt - Notepad++

File Edit Search View Encoding Language Settings Tools Macro Run Plugins Window ?

Demo.txt poison_model.py

```
1 ===Azure ML===
2
3 MLOKit.exe check /platform:azureml /credential:%CRED%
4
5 MLOKit.exe list-projects /platform:azureml /subscription-id:47c5aaab-dbda-44ca-802e-00
6
7 MLOKit.exe |list-models /platform:azureml /credential:%CRED% /subscription-id:47c5aaab-
8
9 MLOKit.exe download-model /platform:azureml /credential:%CRED% /subscription-id:47c5aaab
10
11 MLOKit.exe poison-model /platform:azureml /credential:%CRED% /subscription-id:47c5aaab
12
13
14
15
16
17
18
19
20
21
22
23
```

length : 886 lines : 28 Ln : 7 Col : 12 Pos : 206 Windows (CR LF) UTF-8 IN8

Protecting ML Training Environments



Users

Use password management system

Separate admin account

Additional monitoring controls

Security Awareness Training

MFA enabled and in use

PATs with expiration dates



Notebook Environments

Password protect notebook

IP address restrictions

Limits to kernel execution times

Use virtual environment

Run as non-root account

No cleartext credentials/secrets



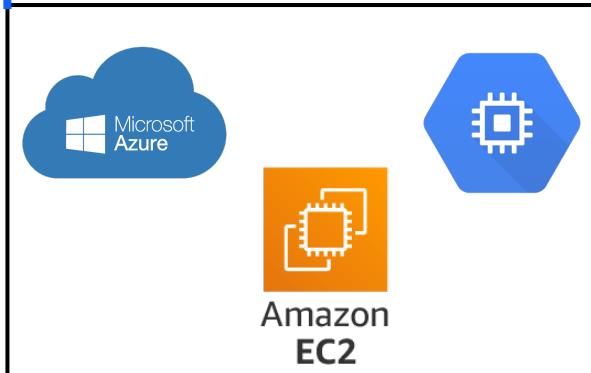
Cloud Compute

Enable auto-shutdown and auto-start schedule

Delete compute if no longer needed

Disable unneeded services

Configure role-based access



Model Artifact Storage and Registry

Cleanup/delete old model artifacts

Restrict access to backend storage

IP-address restrictions

Enable logging and apply detection rules

Implement model integrity verification



Azure Machine Learning



Amazon S3



Amazon SageMaker



Azure Blob
Storage

Detection Guidance – Summary

Azure ML Detections



<https://github.com/h4wkst3r/KQL-Queries>

Dataset Poisoning

Dataset Recon

Dataset Theft

Model Poisoning

Model Recon

Model Theft

SageMaker Detections



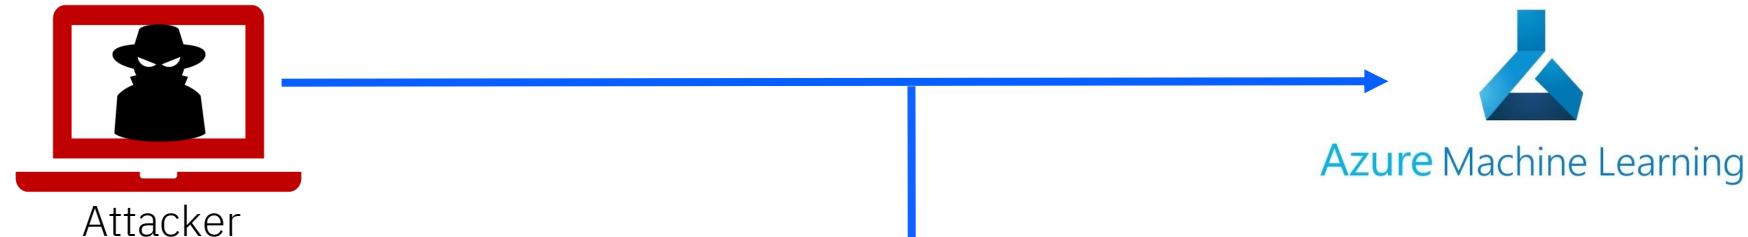
<https://github.com/h4wkst3r/CloudTrail-Queries>

Model Poisoning

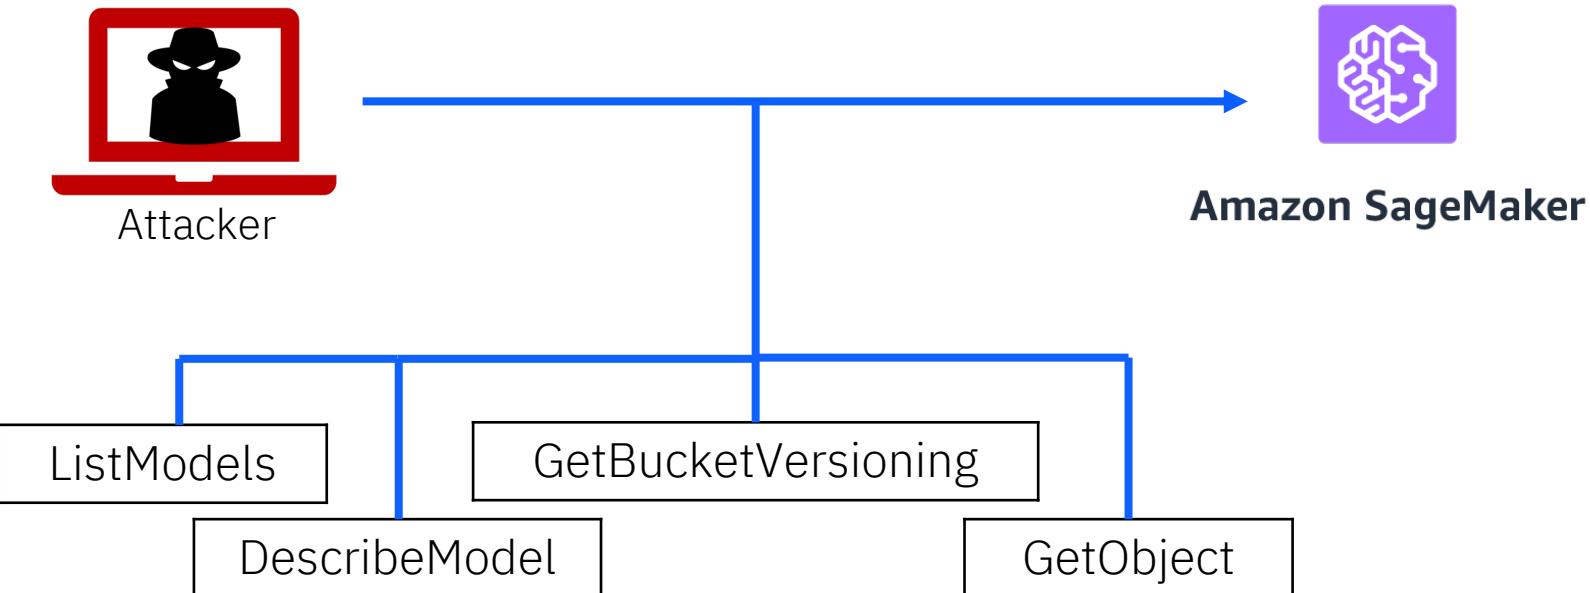
Model Theft

Malicious Lifecycle Config

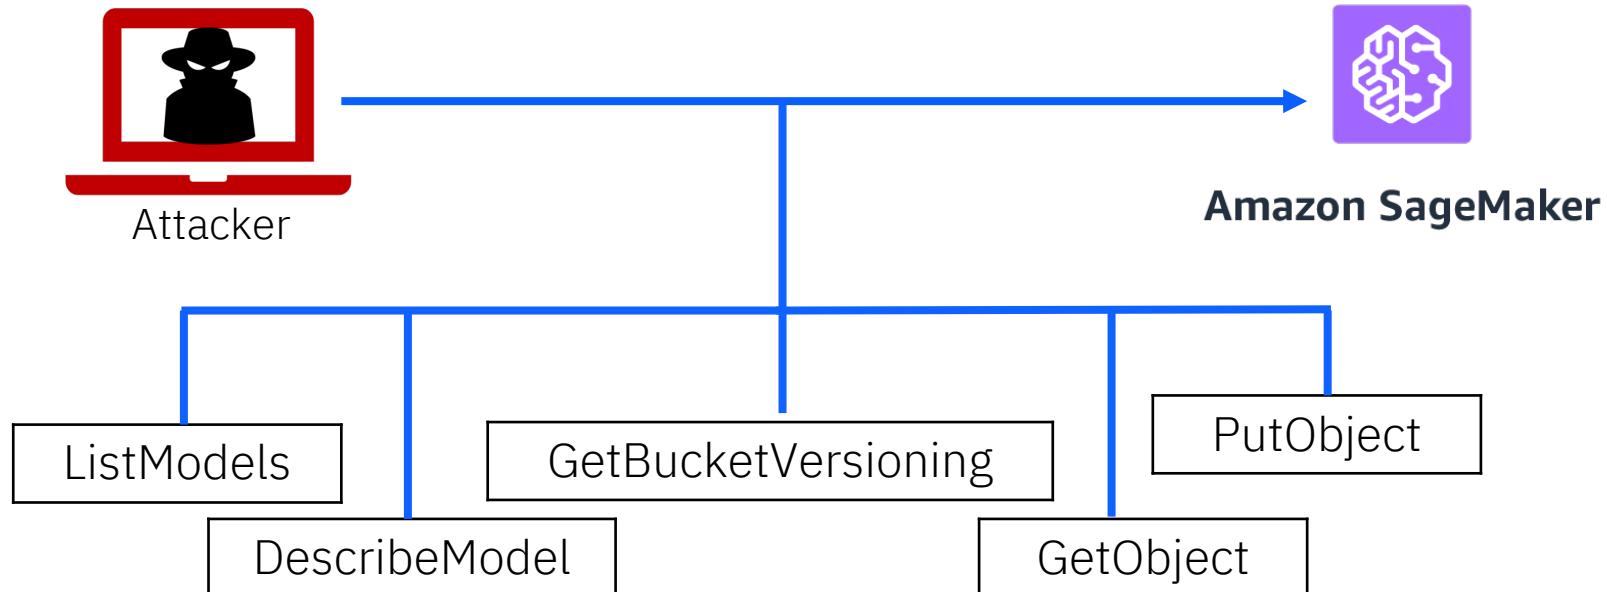
Detection Guidance – Azure ML Model Poisoning



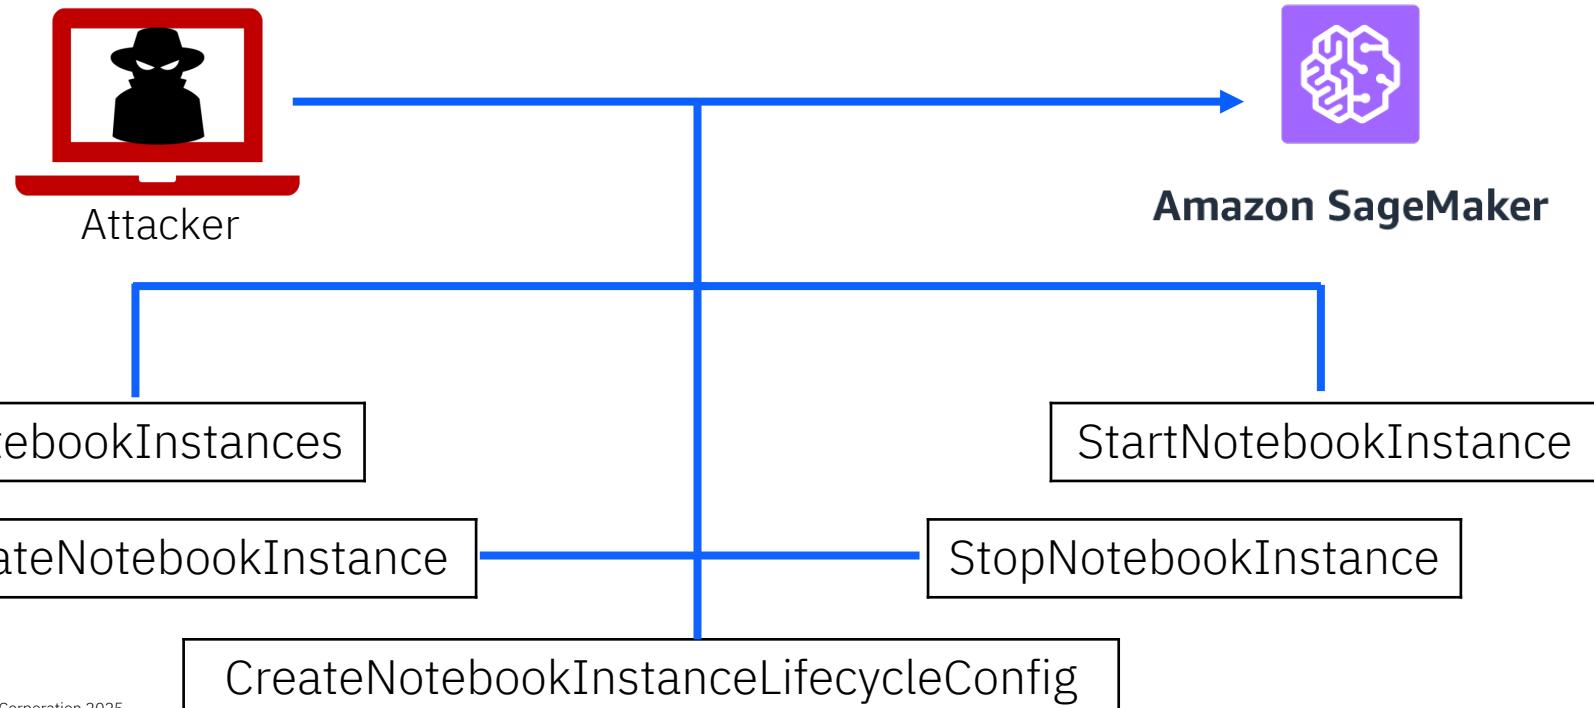
Detection Guidance – SageMaker Model Theft



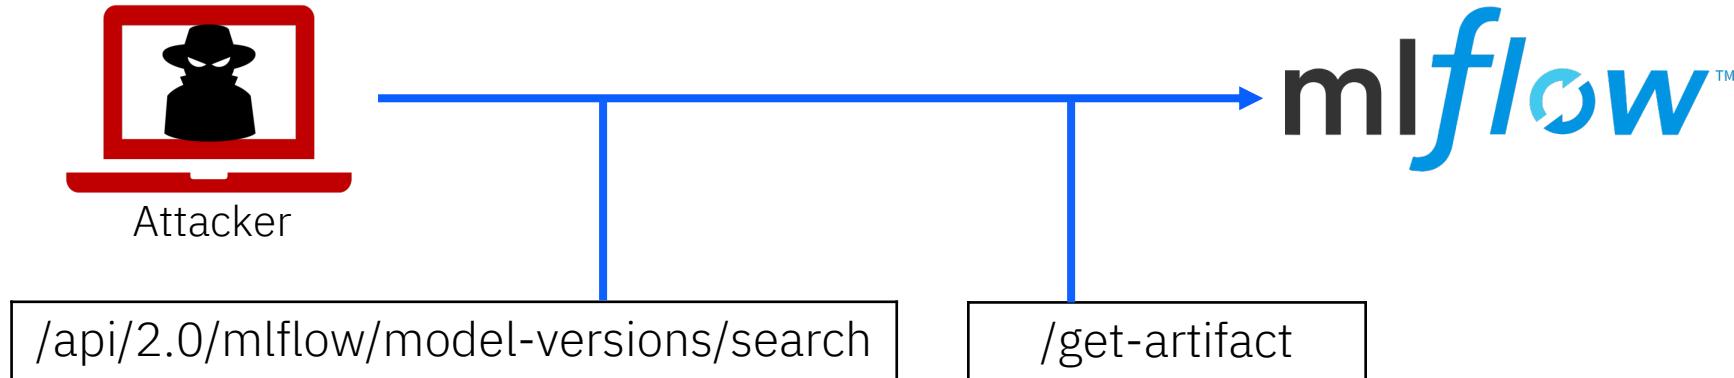
Detection Guidance – SageMaker Model Poisoning



Detection Guidance – SageMaker Malicious Lifecycle Configuration



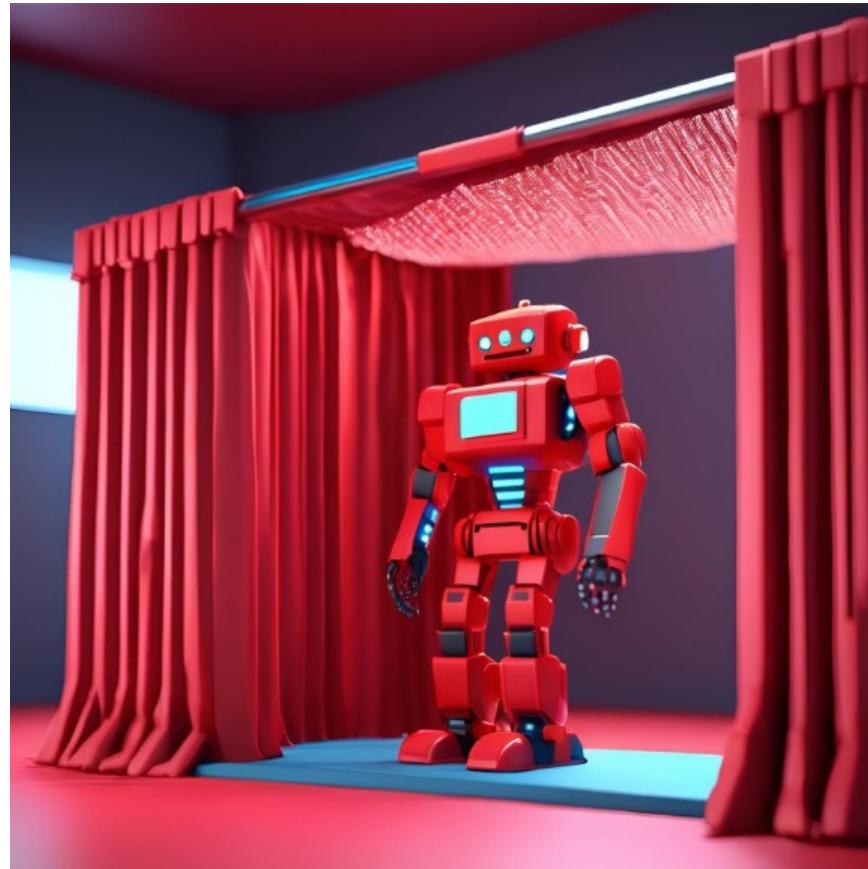
Detection Guidance – MLFlow



```
[root@el7 ~]# curl -X GET http://127.0.0.1:5000/get-artifact?path=ElasticNet/code/input_example.json&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6
grep -i /api/2.0/mlflow/model-versions/search access.log
[2025:20:59:42 -0500] "GET /api/2.0/mlflow/model-versions/search HTTP/1.1" 200 904 "-" "MLOKit-e977ac02118a3cb2
[2025:20:59:42 -0500] "GET /api/2.0/mlflow/model-versions/search HTTP/1.1" 200 904 "-" "MLOKit-e977ac02118a3cb2
[2025:21:05:59 -0500] "GET /get-artifact?path=ElasticNet/code/data.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:00 -0500] "GET /get-artifact?path=ElasticNet/code/params.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:01 -0500] "GET /get-artifact?path=ElasticNet/code/train.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:01 -0500] "GET /get-artifact?path=ElasticNet/code/utils.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:02 -0500] "GET /get-artifact?path=ElasticNet/conda.yaml&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:02 -0500] "GET /get-artifact?path=ElasticNet/input_example.json&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:03 -0500] "GET /get-artifact?path=ElasticNet/model.pkl&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:03 -0500] "GET /get-artifact?path=ElasticNet/python_code.yaml&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
```

```
grep -i /get-artifact access.log
[2025:21:05:59 -0500] "GET /get-artifact?path=ElasticNet/code/data.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:00 -0500] "GET /get-artifact?path=ElasticNet/code/params.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:01 -0500] "GET /get-artifact?path=ElasticNet/code/train.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:01 -0500] "GET /get-artifact?path=ElasticNet/code/utils.py&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:02 -0500] "GET /get-artifact?path=ElasticNet/conda.yaml&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:02 -0500] "GET /get-artifact?path=ElasticNet/input_example.json&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:03 -0500] "GET /get-artifact?path=ElasticNet/model.pkl&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
[2025:21:06:03 -0500] "GET /get-artifact?path=ElasticNet/python_code.yaml&run_id=f0b129a5e11c4bb6b1c0459a2f5ae7f6 HTTP/1.1" 200 0 "-" "MLOKit-e977ac02118a3cb2
```

Conclusion



Conclusion

01

ML training environments contain highly sensitive and business critical data

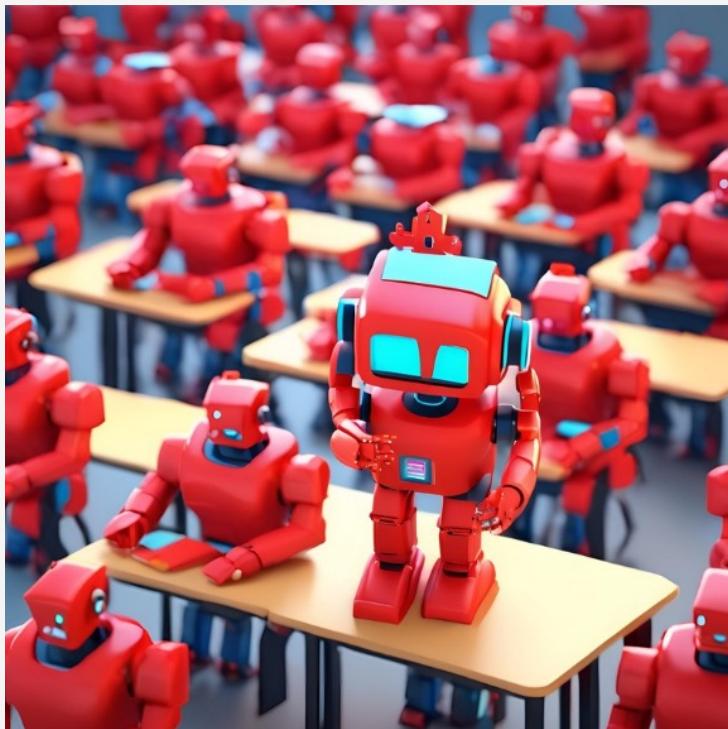
02

We need to understand these systems so we can protect them

03

Unauthorized access to these environments could be significant

Questions?



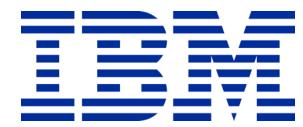
Brett Hawkins
[@h4wkst3r](https://twitter.com/h4wkst3r)  

Blog Post -



MLOKit Tool -





References

- <https://twitter.com/retBandit>
- <https://twitter.com/h4wkst3r>
- <https://www.ibm.com/downloads/documents/us-en/11630e2cbc302316>
- <https://github.com/xforceder/MLOKit>
- <https://twitter.com/azarzaror>
- <https://web.archive.org/web/20241214000321/https://www.panoptica.app/blog/protect-your-environment-when-working-with-amazon-sagemaker>
- <https://www.ibm.com/topics/mlops>
- <https://neptune.ai/blog/mlops-tools-platforms-landscape>
- <https://www.databricks.com/glossary/jupyter-notebook>
- <https://www.ibm.com/docs/en/z-devops-guide?topic=applications-source-code-management>
- <https://azure.microsoft.com/en-us/resources/cloud-computing-dictionary/what-is-cloud-computing>
- <https://mlflow.org/docs/latest/tracking/artifacts-stores.html>

References

- <https://wandb.ai/site/articles/what-is-an-ML-model-registry/>
- <https://learn.microsoft.com/en-us/azure/devops/user-guide/what-is-azure-devops>
- <https://about.gitlab.com/>
- <https://github.com/>
- <https://aws.amazon.com/sagemaker/>
- <https://learn.microsoft.com/en-us/azure/machine-learning/overview-what-is-azure-machine-learning>
- <https://mlflow.org/>
- <https://docs.aws.amazon.com/sagemaker/latest/dg/model-registry.html>
- <https://learn.microsoft.com/en-us/azure/machine-learning/how-to-manage-models?view=azureml-api-2&tabs=cli>
- <https://aws.amazon.com/s3/>
- <https://learn.microsoft.com/en-us/azure/storage/blobs/storage-blobs-introduction>
- <https://www.ibm.com/services/adversary-simulation>

References

- <https://mlflow.org/docs/latest/rest-api.html>
- <https://learn.microsoft.com/en-us/azure/devops/organizations/accounts/use-personal-access-tokens-to-authenticate?view=azure-devops&tabs=Windows>
- <https://mlflow.org/docs/latest/projects.html>
- <https://github.com/xforceder/ADOKit>
- <https://aws.amazon.com/blogs/machine-learning/build-an-end-to-end-mlops-pipeline-using-amazon-sagemaker-pipelines-github-and-github-actions/>
- <https://docs.aws.amazon.com/IAM/latest/UserGuide/security-creds.html>
- <https://docs.aws.amazon.com/sagemaker/latest/dg/security-iam-awsmanpol.html>
- <https://docs.aws.amazon.com/sagemaker/latest/dg/nbi.html>
- <https://docs.aws.amazon.com/sagemaker/latest/dg/notebook-lifecycle-config.html>
- <https://docs.aws.amazon.com/AmazonS3/latest/userguide/security-iam-awsmanpol.html>
- <https://jfrog.com/blog/from-mlops-to-mloops-exposing-the-attack-surface-of-machine-learning-platforms/>

References

- <https://docs.python.org/3/library/pickle.html>
- <https://gist.github.com/h4wkst3r/2c30a3d39e20b7cd8606211ba3132d85/raw/e4e253b052f5b916134409fc b61b91c49b91d912/CreatePickle.py>
- <https://github.com/coldwaterq/MaliciousPickles>
- <https://github.com/moohax/Charcuterie>
- <https://github.com/trailofbits/fickling>
- <https://github.com/hiddenlayerai/HiddenPickle>
- <https://docs.aws.amazon.com/sagemaker/latest/dg/how-it-works-deployment.html>
- <https://www.microsoft.com/en-us/security/blog/2025/02/13/storm-2372-conducts-device-code-phishing-campaign/>
- <https://learn.microsoft.com/en-us/entra/identity-platform/access-tokens>
- <https://learn.microsoft.com/en-us/rest/api/azureml/>
- <https://learn.microsoft.com/en-us/azure/machine-learning/concept-workspace?view=azureml-api-2>
- <https://docs.gunicorn.org/en/stable/settings.html>

References

- <https://learn.microsoft.com/en-us/microsoft-365/enterprise/subscriptions-licenses-accounts-and-tenants-for-microsoft-cloud-offerings?view=o365-worldwide>
- <https://learn.microsoft.com/en-us/azure/machine-learning/concept-data?view=azureml-api-2>
- <https://learn.microsoft.com/en-us/azure/machine-learning/concept-endpoints?view=azureml-api-2>
- <https://blog.gitguardian.com/how-to-handle-secrets-in-jupyter-notebooks/>
- <https://medium.com/@techlatest.net/security-best-practices-for-ai-ml-in-jupyter-notebooks-a-blog-post-on-the-security-best-practices-c0e0659cfccb>
- <https://github.com/h4wkst3r/KQL-Queries>
- <https://github.com/h4wkst3r/CloudTrail-Queries>
- <http://www.ibm.com/downloads/documents/us-en/11630e2cbc302316>
- <https://github.com/aws-samples/aws-customer-playbook-framework/blob/main/docs/Responding%20to%20SageMaker.md>
- <https://aws.amazon.com/blogs/machine-learning/securing-mlflow-in-aws-fine-grained-access-control-with-aws-native-services/>

References

- <https://learn.microsoft.com/en-us/azure/storage/blobs/monitor-blob-storage?tabs=azure-portal>
- <https://learn.microsoft.com/en-us/azure/machine-learning/monitor-azure-machine-learning?view=azureml-api-2>
- <https://github.com/h4wkst3r/KQL-Queries/blob/main/AzureML/AzureMLModelPoisoning.kql>
- <https://github.com/h4wkst3r/CloudTrail-Queries/blob/main/SageMakerModelTheft.sql>
- <https://github.com/h4wkst3r/CloudTrail-Queries/blob/main/SageMakerModelPoisoning.sql>
- <https://github.com/h4wkst3r/CloudTrail-Queries/blob/main/SageMakerMaliciousLifecycleConfig.sql>
- <https://www.ibm.com/think/x-force/becoming-the-trainer-attacking-ml-training-infrastructure>